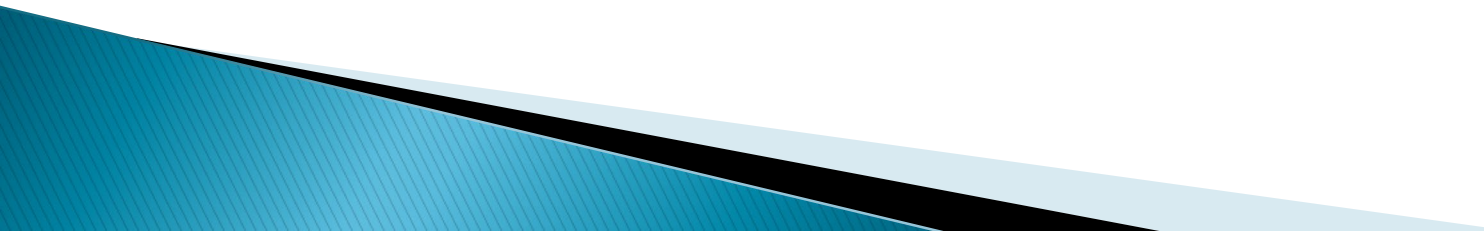


Additional Rendering Considerations



Topics

- ▶ Forward rendering
 - ▶ Object culling
 - ▶ Light culling
 - ▶ Dynamic light loops
- 

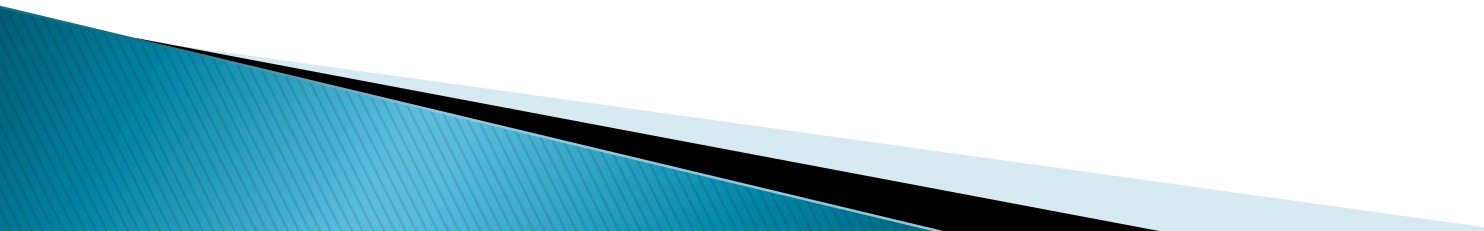
Forward Rendering



Forward rendering



- ▶ The rendering technique we've been using
- ▶ Calculating radiance* for a pixel as geometry is being drawn
 - *Radiance: emitted or reflected light
- ▶ Required information:
 - Geometry: position, normal, etc.
 - Material: shininess, color, etc.
 - Light: color, direction, etc.

Forward rendering: Pros

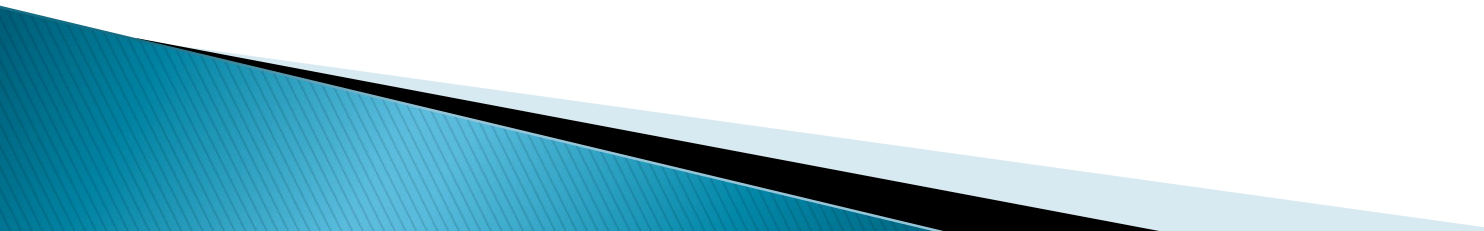
- ▶ Easy to implement
 - ▶ Makes a lot of sense
 - Gotta run a pixel shader – might as well do lighting there!
 - ▶ Works with opaque & transparent objects
 - As long as transparent objects are rendered after all opaque objects
 - And transparent objects are sorted back to front
 - ▶ Works with hardware anti-aliasing (MSAA)
- 

Forward rendering: Cons

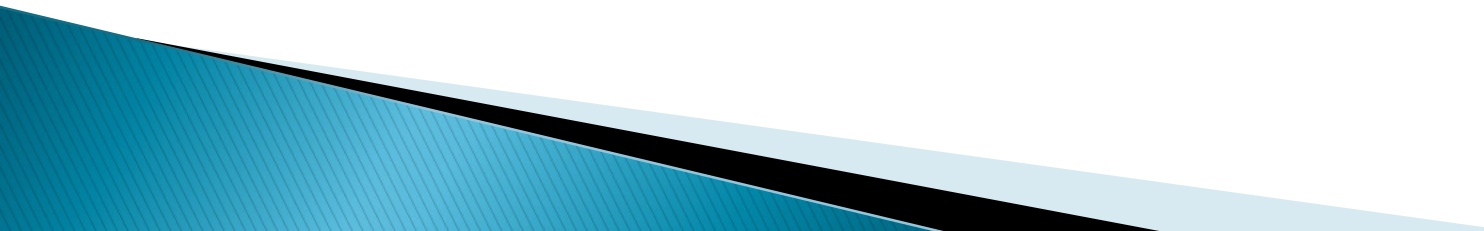
- ▶ Overdraw* is expensive
 - *Drawing a pixel, then covering it
 - Later within the same frame
- ▶ Lighting calculated for entire objects
 - Even when lights have very small ranges
- ▶ Lighting complexity is
 - $O(\text{number of objects} * \text{number of lights})$



Alternatives to forward rendering

- ▶ Deferred rendering
 - Defers lighting until after all geometry is rendered
 - Requires multiple render targets
 - ▶ Forward+ rendering
 - Divides screen into very small tiles (8x8 pixels)
 - Process lights whose bounds overlap a tile
 - Usually involves compute shaders
 - ▶ Advanced topics requiring different C++ render architecture
- 

Improvements to forward rendering

- ▶ $O(\text{objects} * \text{lights}) \rightarrow$ Reduce objects or reduce lights!
 - ▶ Cull objects
 - Check camera bounds (frustum) in C++
 - Only draw if object is visible
 - A bounding volume hierarchy (octree or similar) is useful here
 - ▶ Cull lights
 - Determine which lights might actually affect an object
 - Only pass those lights to the shader
- 


Culling Objects

Culling objects

- ▶ Skipping objects we can't see – *before* we render them!
- ▶ If it won't be seen this frame...
 - Don't even attempt to render it
 - Saves us from even having to transform the vertices
- ▶ Multiple ways of handling this
 - Per-object checking
 - Grid-based checking
 - Bounding volume hierarchy (BVH)

Culling & the camera

- ▶ The camera frustum
 - The bounds of the camera
 - Defined by the projection, specifically
- ▶ Anything inside the frustum can be seen
- ▶ Anything outside the frustum cannot
- ▶ If we can detect objects inside/outside in C++
 - We can decide if we even need to render




Culling requires entity bounds

- ▶ Need to know an entity's world-space bounds
- ▶ Mesh could store simple AABB bounding box
 - But we need more than this
 - This only tells us the size of the mesh before transformations
- ▶ Needs to take into account entity's transform!
 - Transform bounds by world matrix → Oriented Bounding Box (OBB)
 - Or build new AABB from transformed AABB corners

Culling: Per-object

- ▶ Compare each object's bounds to the frustum
 - Box/Frustum intersection test




- ▶ Easy but somewhat slow: $O(n)$

Culling: Grid-based


- ▶ Divide scene into grid (3D array)
 - Grid holds all objects
 - Test grid cell against frustum
 - Draw any objects in a “valid” cell

- ▶ Usually faster if cells hold > 1 object
 - Can result in false positives!

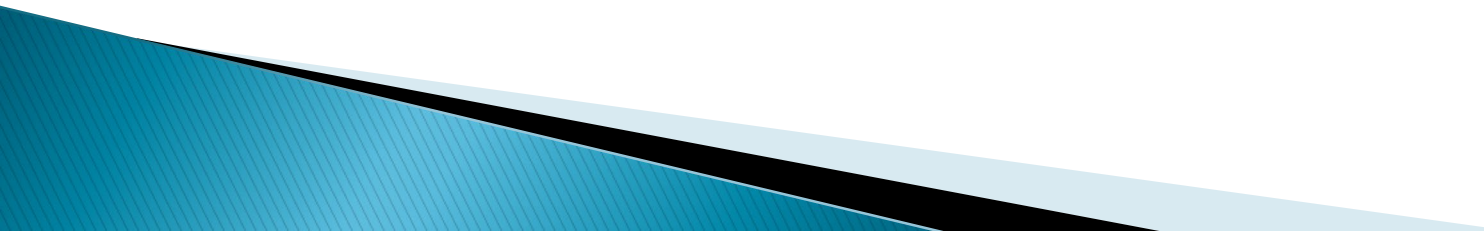


Culling: Bounding Volume Hierarchy (BVH)

- ▶ A hierarchy (tree) of bounding volumes (chunks of space)
 - Quadtree (2D)
 - Octree (3D)
 - Binary space partitioning tree (3D)
- ▶ Similar to grid-based, but with a tree
- ▶ Generally fastest: $O(\log n)$
 - Also trickiest to implement



When do we cull?

- ▶ Beginning of Draw()
 - Camera & all entities need to update first
 - Hopefully all bounding boxes are updated at this point
 - ▶ Render loop can perform intersection tests
 - Only render entities that “pass the test”
 - ▶ Alternatively, loop through everything first
 - Collect list of objects that need to be rendered
 - Can perform other useful operations (sorting?) on this list
- 

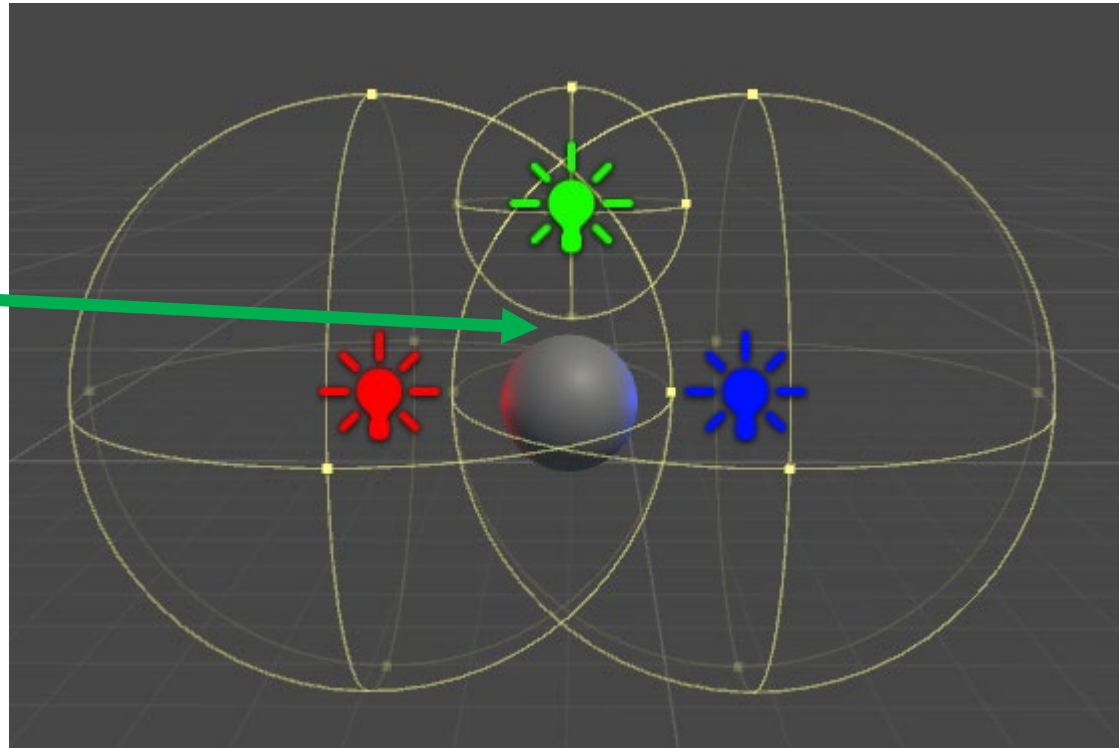
Culling Lights

Light culling

- ▶ Don't process lights that won't affect objects
 - No sense in doing unnecessary work

Green light doesn't intersect sphere.

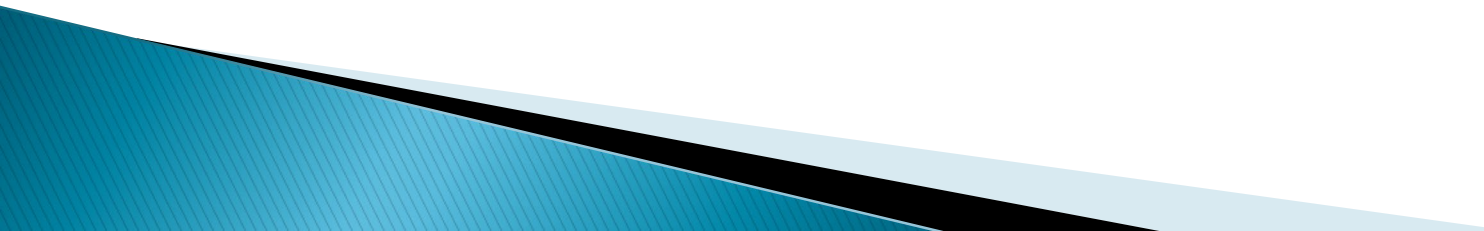
No need to even send to the shader.



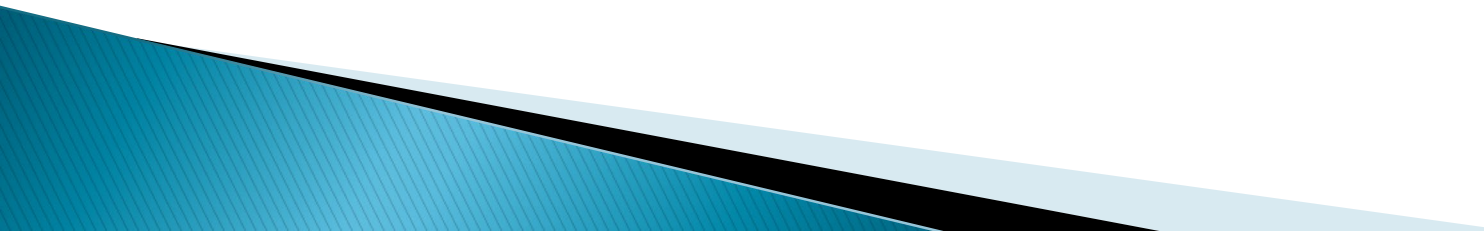
Light culling

- ▶ Culling occurs in C++
 - Requires range-based point/spot attenuation
- ▶ Determine which lights affect objects:
 - **Point:** Based on light's range
 - **Spot:** Based on range, cone size and direction
 - **Directional:** These hit *everything*
- ▶ Requires each entity to know its bounds
 - Mesh could store simple AABB bounding box
 - Needs to take into account entity's transform!

Light culling – What to send to GPU?

- ▶ Before rendering each object:
 - Loop through lights
 - Check each point & spot light against object bounds
 - Add overlapping lights to *intersecting lights* list
 - ▶ Only send *intersecting lights* to shader
 - **Light array** only holds lights that might actually hit object
 - **Light count** is number of lights in *intersecting lights* list
 - ▶ Requires a light loop in HLSL that iterates based on CPU data
- 

Is light culling worth it?

- ▶ Is this faster than just letting the GPU chug through lights?
 - Maybe? Maybe not? Depends on your implementation!
 - ▶ What is the cost of each light in HLSL?
 - More expensive BRDFs take more time
 - ▶ What is the cost of light culling?
 - Brute force is slower than using a good bounding volume hierarchy
 - ▶ Test and find out for your specific application
- 

Dynamic Light Loop

In a shader



Dynamic light loop

- ▶ Process a dynamic # of lights in a shader
 - Loop through an array of lights
 - Add the results of each light calculation
- ▶ This is similar to what we've discussed already
- ▶ Except with a *dynamic* number of lights
 - Don't hardcode light1, light2, etc.
 - Don't assume there are always N lights

Dynamic light count – Light array

```
// Define the max number of lights
```

```
#define MAX_LIGHTS 128
```

```
cbuffer externalData : register(b1)
```

```
{
```

```
    // An array of lights (const size)
```

```
    LightStruct lights[MAX_LIGHTS];
```

```
    // Amount of lights to process THIS FRAME
```

```
    int lightCount;
```

```
}
```



Example light struct setup

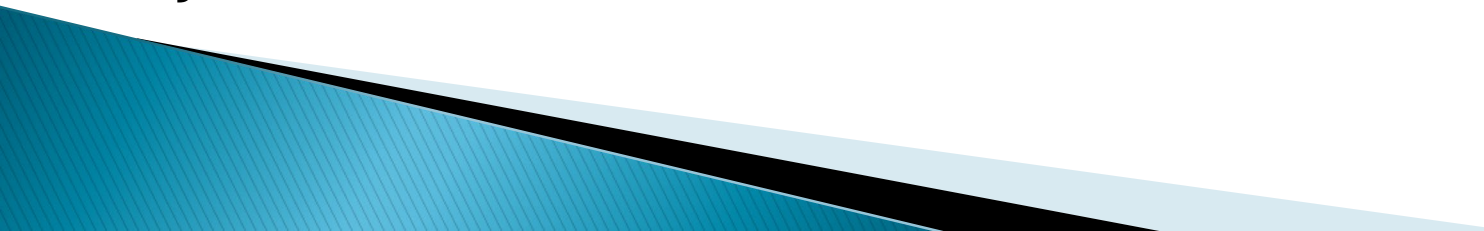
```
#define LIGHT_TYPE_DIR      0
#define LIGHT_TYPE_POINT   1
#define LIGHT_TYPE_SPOT    2

// Make a matching struct on CPU (C++) side
struct LightStruct
{
    int      Type;
    float3 Direction;      // 16 bytes
    float  Range;
    float3 Position;       // 32 bytes
    float  Intensity;
    float3 Color;           // 48 bytes
    float  SpotFalloff;
    float3 Padding;        // 64 bytes
};
```

Light loop pseudo-code

```
float3 c = float3(0,0,0); // Total color

// Loop through all lights THIS FRAME
// (Assumes lightCount < MAX_LIGHTS)
for (int i = 0; i < lightCount; i++)
{
    switch (lights[i].Type)
    {
        case LIGHT_TYPE_DIR:    c += DirLight(...); break;
        case LIGHT_TYPE_POINT:  c += PointLight(...); break;
        case LIGHT_TYPE_SPOT:   c += SpotLight(...); break;
    }
}
```



Dynamic light loop: Pros & cons

- ▶ This is the method I usually use
- ▶ Pros:
 - Relatively simple to implement
 - Flexible amount of lights (up to the max)
- ▶ Cons:
 - Branches (luckily each pixel takes the same branch – best case)
 - One fairly expensive shader – overdraw is *rough*
 - Lighting complexity is still $O(\text{objects} * \text{lights})$