

DirectX Math

Quick overview & examples

DirectX Math

- ▶ SIMD-friendly C++ data types and functions
 - Common linear algebra
 - Common graphics functions
- ▶ Vectors
- ▶ Matrices (creation and math)
- ▶ Quaternions
- ▶ Colors, planes, triangle-ray intersections and more!
- ▶ [DirectXMath Programming Guide](#) (overview)
- ▶ [DirectXMath Programming Reference](#) (function list)

DirectX Math

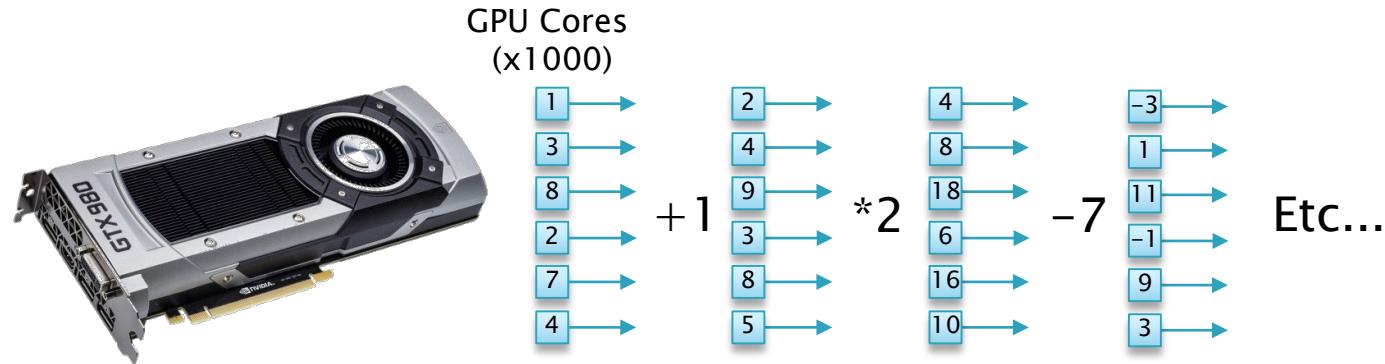
- ▶ This library requires a very specific usage pattern
- ▶ Highly optimized *for specific use cases*
- ▶ SIMD-friendly
 - “Single Instruction, Multiple Data”
 - Uses SSE intrinsics (Streaming SIMD Extension)
 - Allows parallel processing of data components

SIMD

Single Instruction, Multiple Data

SIMD?

- ▶ Remember our discussion of SIMD GPU cores?



- ▶ How GPUs can achieve massive parallelism:
 - Give each core a **program** to follow (shaders)
 - Give each program **different data** (vertex data)
 - Allow them to access **shared data** (constant buffers)
 - Get different **results** (pixel positions & colors)

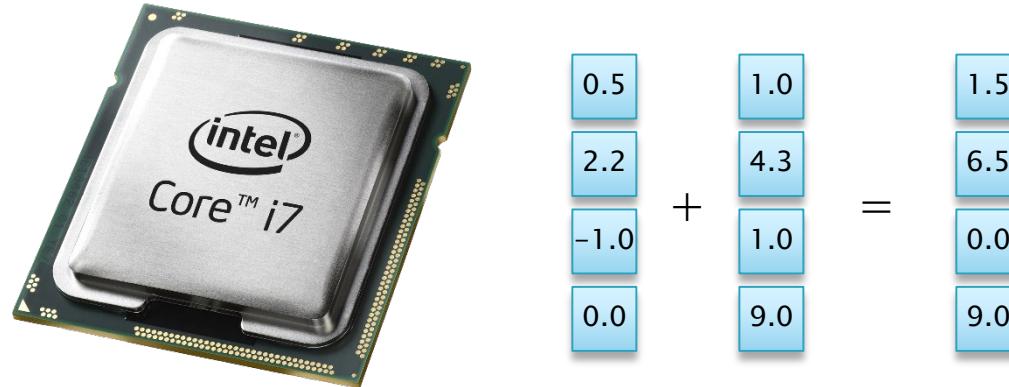
SIMD everywhere?

- ▶ SIMD
 - Great for number crunching!
- ▶ Can we apply this everywhere? C++ code?
 - SIMD isn't great w/ branching
 - If statements, switches, etc.
- ▶ Doesn't work well for general purpose code
 - GPUs custom-built to work this way
 - But...



CPU SIMD for vector math!

- ▶ But we DO perform lots of *vector math!*
 - Each vector has up to 4 numbers
 - Could we apply SIMD to vectors *in C++?*



- ▶ Yes!
 - Modern CPUs can speed up C++ vector math
 - Have special registers for 4-component SIMD math
 - CPU-side SIMD-specific math library required

Using DirectX Math

Using DirectX Math

- ▶ Basic header file is DirectXMath.h
- ▶ “DirectX” namespace
 - Contains all data types and functions
 - Also contains operator overloads for some data types
 - Don’t forget the using statement (in .cpp files)!

Categories of data types

- ▶ **Storage data types**

- Many of these
- No alignment issues
- Perfect for use as class/struct member variables

- ▶ **SIMD Math data types**

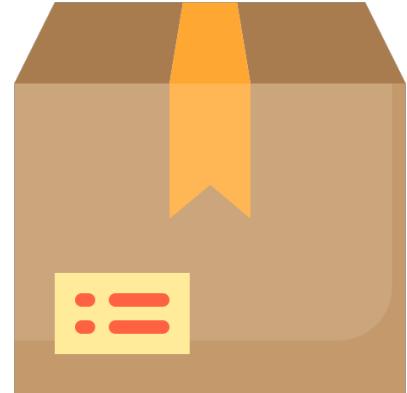
- Exactly two of these
- Required by all math functions in DirectXMath
- Special alignment rules in memory
- Should **not** be member variables!

Storage Data Types

XMFLOAT2, XMFLOAT3, XMFLOAT4, XMFLOAT4X4, etc.

Storage data types

- ▶ Begin with “XMFLOAT”
 - Vectors: `XMFLOAT2`, `XMFLOAT3`, `XMFLOAT4`
 - Matrix: `XMFLOAT3X3`, `XMFLOAT3X4`, `XMFLOAT4X3`, `XMFLOAT4X4`
- ▶ Structs of float values
 - Can be interpreted as arrays of floats
 - Has members for each component (.x for example)
- ▶ Safe to use as class/struct members



Creating XMFLOAT vectors

- ▶ Different data types for 2, 3, and 4–component vectors

```
// Simple initialization of stack objects
XMFLOAT2 pos2D(2.5f, 3.0f);
XMFLOAT3 direction3D(0.0f, 0.0f, 1.0f);
XMFLOAT4 color(1.0f, 0.0f, 1.0f, 0.5f);
```

```
// Same as above but technically require a copy (less efficient)
XMFLOAT2 pos2D          = XMFLOAT2(2.5f, 3.0f);
XMFLOAT3 direction3D    = XMFLOAT3(0.0f, 0.0f, 1.0f);
XMFLOAT4 color          = XMFLOAT4(1.0f, 0.0f, 1.0f, 0.5f);
```

Creating XMFLOAT matrices

- ▶ You can make matrices this way too
 - But we'll see better ways to initialize common transform matrices

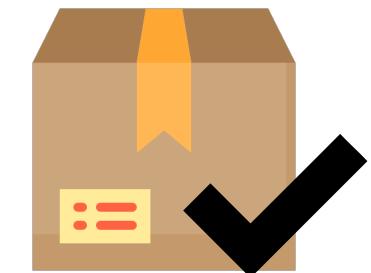
```
XMFLOAT4X4 matrix(1,0,0,0,0,1,0,0,0,0,1,0,0,0,0,1);
```

```
// The same, but a bit easier to read
```

```
XMFLOAT4X4 matrix(  
    1, 0, 0, 0,  
    0, 1, 0, 0,  
    0, 0, 1, 0,  
    0, 0, 0, 1);
```

Using XMFLOATs

- ▶ Okay, we can create and store data
- ▶ How do we do math with XMFLOATs?
- ▶ We *don't*
- ▶ They're for storage only!
 - No math functions accept or produce XMFLOATs
 - No overloaded operators
 - Explicitly designed for *not doing math*

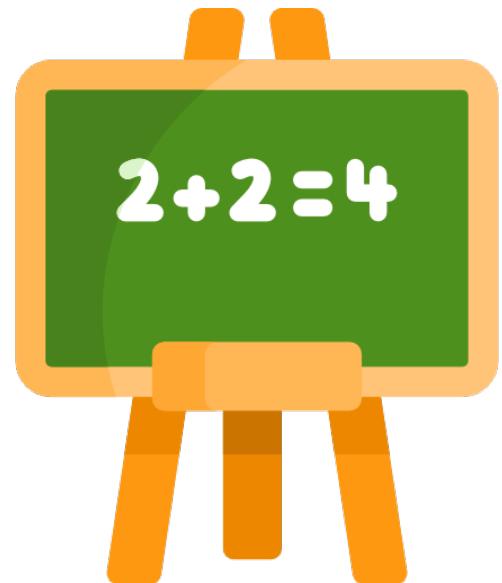


SIMD Math Data Types

XMVECTOR and XMMATRIX

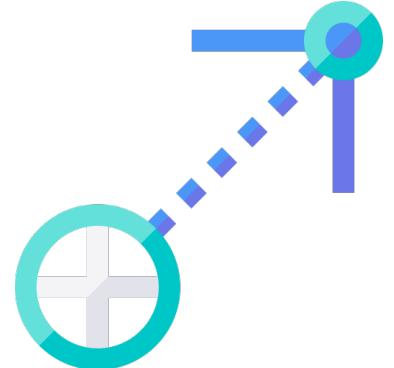
SIMD Math data types

- ▶ Exactly 2 of these:
 - **XMFLOAT3** – Always has room for 3 floats
 - **XMMATRIX** – Always has room for 16 floats
- ▶ Designed to be used as **local (stack) variables**
 - Mapped to hardware vector registers
 - Must be aligned optimally in memory
- ▶ These types are the “work horses” of DXMath
 - Every function *uses* or *produces* these



XMVECTOR

- ▶ Always stores 4 components
 - X, y, z, w
- ▶ Interpreted differently depending on usage
 - XMVector2Length() Uses first 2 components
 - XMVector3Dot() Uses first 3 components
 - XMVector4Transform() Uses all 4 components



Creating XMVECTORs from 4 values

- ▶ Use XMVectorSet()
 - Creates an XMVECTOR from 4 distinct values
 - Always requires all 4 values (even if *you* don't need them)

```
// Assuming we're "using namespace DirectX;" here
XMVECTOR v1 = XMVectorSet(1.0f, 0.0f, 2.5f, 0.0f);
XMVECTOR v2 = XMVectorSet(0.0f, 1.0f, -1.0f, 0.0f);
```

Creating XMVECTORs from XMFLOATs

- ▶ More commonly, you'll convert from *storage* to *math* types
- ▶ Known as *loading* – loading into special CPU registers

```
// Created during initialization perhaps
XMFLOAT2 pos2D(2.5f, 3.0f);
XMFLOAT3 direction3D(0.0f, 0.0f, 1.0f);
XMFLOAT4 color(1.0f, 0.0f, 1.0f, 0.5f);

// Later in your code, when you want to do some math
XMVECTOR pos2DVec = XMLoadFloat2(&pos2D);
XMVECTOR dir3DVec = XMLoadFloat3(&direction3D);
XMVECTOR colorVec = XMLoadFloat4(&color);
```

XMMATRIX

- ▶ Always stores 16 components
- ▶ We'll exclusively use them for 4x4 matrices
 - 3D transformations
 - Camera view matrices
 - Camera projection matrices

$$\left\{ \begin{array}{cccc} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{array} \right\}$$

Creating XMMATRIX from XMFLOATS

- ▶ Same idea as vectors
 - Load from an XMFLOAT4X4
 - Into an XMMATRIX

```
// Storage type - created during init
XMFLOAT4X4 identity(
    1, 0, 0, 0,
    0, 1, 0, 0,
    0, 0, 1, 0,
    0, 0, 0, 1);

// Load a stored matrix to do some math
XMMATRIX identMat = XMLoadFloat4x4(&identity);
```

Easily creating useful matrices

- ▶ Probably don't want to type out 16 numbers for each matrix
 - Lots of math to remember, too
- ▶ DirectXMath has many matrix transformation functions

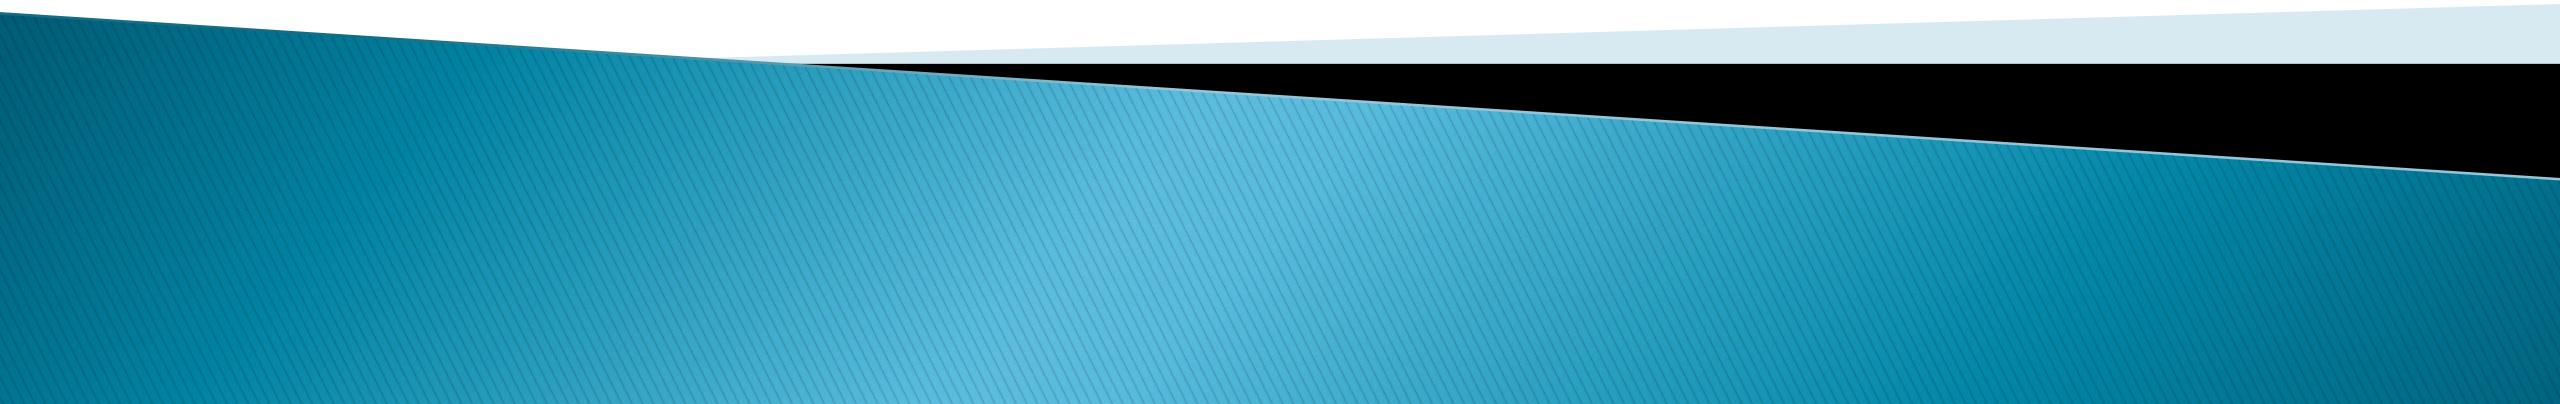
```
XMMATRIX ident      = XMMatrixIdentity();
XMMATRIX translate = XMMatrixTranslation(x, y, z);
XMMATRIX rotate    = XMMatrixRotationRollPitchYaw(pitch, yaw, roll);
XMMATRIX scale     = XMMatrixScaling(sx, sy, sz);
```

Creating XMATRIX from XVECTORS

- ▶ Already have your data in an XVECTOR?
 - No need to separate out that data
 - Call the “FromVector” variants

```
XMMATRIX translate = XMMatrixTranslationFromVector(positionVec);  
XMMATRIX rotate    = XMMatrixRotationRollPitchYawFromVector(pyrVec);  
XMMATRIX scale     = XMMatrixScalingFromVector(scaleVec);  
  
XMMATRIX rotQuat   = XMMatrixRotationQuaternion(quaternion); // XVECTOR param
```

Doing Some Math



Does DirectXMath do...

- ▶ Yes, it probably does!
- ▶ Look up functions
 - All start with “XM”
 - Use the programming reference [function list](#)
 - Or Intellisense while coding
- ▶ Need a left-handed orthographic project matrix? Yup
- ▶ Need a right-handed look-at view matrix? Yup
- ▶ Need to rotate a 3D vector by a quaternion? Yup
- ▶ Need to decompose a matrix? Yup

Math examples – Vectors

- ▶ Use various DirectX Math functions
 - Once you have your data in math types

```
XMVECTOR dot = XMVector3Dot(dirVec1, dirVec2);
```

```
XMVECTOR rotatedVec = XMVector3Rotate(startVec, quaternion);
```

```
XMVECTOR crossVec = XMVector3Cross(startVec, rotatedVec);
```

Math examples – Matrices

```
// Create individual transformations & combine  
XMMATRIX trans = XMMatrixTranslation(10, 0, 0);  
XMMATRIX rot = XMMatrixRotationRollPitchYaw(0, 0.1f, 0);  
XMMATRIX scale = XMMatrixScaling(0.5f, 0.5f, 0.5f);  
  
XMMATRIX worldMatrix =  
    XMMatrixMultiply( XMMatrixMultiply(scale, rot), trans);  
  
// Alternatively (slower than XMMatrixMultiply)  
XMMATRIX worldMatrix = scale * rot * trans;
```

Storing Your Results

Storing your results

- ▶ We did some math
 - Now what?
- ▶ *Store* your results back into a storage data type

```
XMFLOAT3 pos3D;  
XMStoreFloat3(&pos3D, posVec); // posVec is XMVECTOR
```

- ▶ The opposite of *loading*

Store functions

- ▶ Store functions exist for all storage types

```
XMStoreFloat();    // Just the first component  
XMStoreFloat2();  
XMStoreFloat3();  
XMStoreFloat4();  
XMStoreFloat4x4();
```

- ▶ Each has similar parameters:
 - Address of a corresponding XMFLOAT variable
 - An XMFLOAT4 or XMMATRIX with data to store

Loading and storing example

```
// Example storage type variables in memory
XMFLOAT3 dir1(1.0f, 1.0f, 0.0f);
XMFLOAT3 dir2(0.0f, 1.0f, 1.0f);

// Load into aligned XMVECTORs
XMVECTOR vecDir1 = XMLoadFloat3(&dir1);
XMVECTOR vecDir2 = XMLoadFloat3(&dir2);

// Do any calculations necessary
// (Often more than one - do all the work here!)
XMVECTOR dot = XMVectorDot(vecDir1, vecDir2);

// Store in "regular" variable when completely done
float dotProduct;
XMStoreFloat(&dotProduct, dot);
```

Simplified load/store example

```
// Example “regular” variables in memory
XMFLOAT3 dir1(1.0f, 1.0f, 0.0f);
XMFLOAT3 dir2(0.0f, 1.0f, 1.0f);

// Load and calculate in one statement
XMVECTOR dot = XMVectorDot(
    XMLoadFloat3(&dir1),
    XMLoadFloat3(&dir2));

// Store when done
// If you have other calculations, don't store yet!
float dotProduct;
XMStoreFloat(&dotProduct, dot);
```

Complete example – Relative movement

```
// Moves relative to current rotation
// Ex: Values (0,0,1) will move along local forward axis
void MoveRelative(float x, float y, float z)
{
    // Rotate desired movement by our rotation
    XMVECTOR relativeDirection = XMVector3Rotate(
        XMVectorSet(x, y, z, 0),
        XMLoadFloat4(&currentRot)); // <-- Quaternion

    // Add to position and store
    XMStoreFloat3(
        &currentPosition,
        XMLoadFloat3(&currentPosition) + relativeDirection);
}
```

Shortcuts?

- ▶ Necessary to load/store for every operation?
 - Maybe not for very simple ones

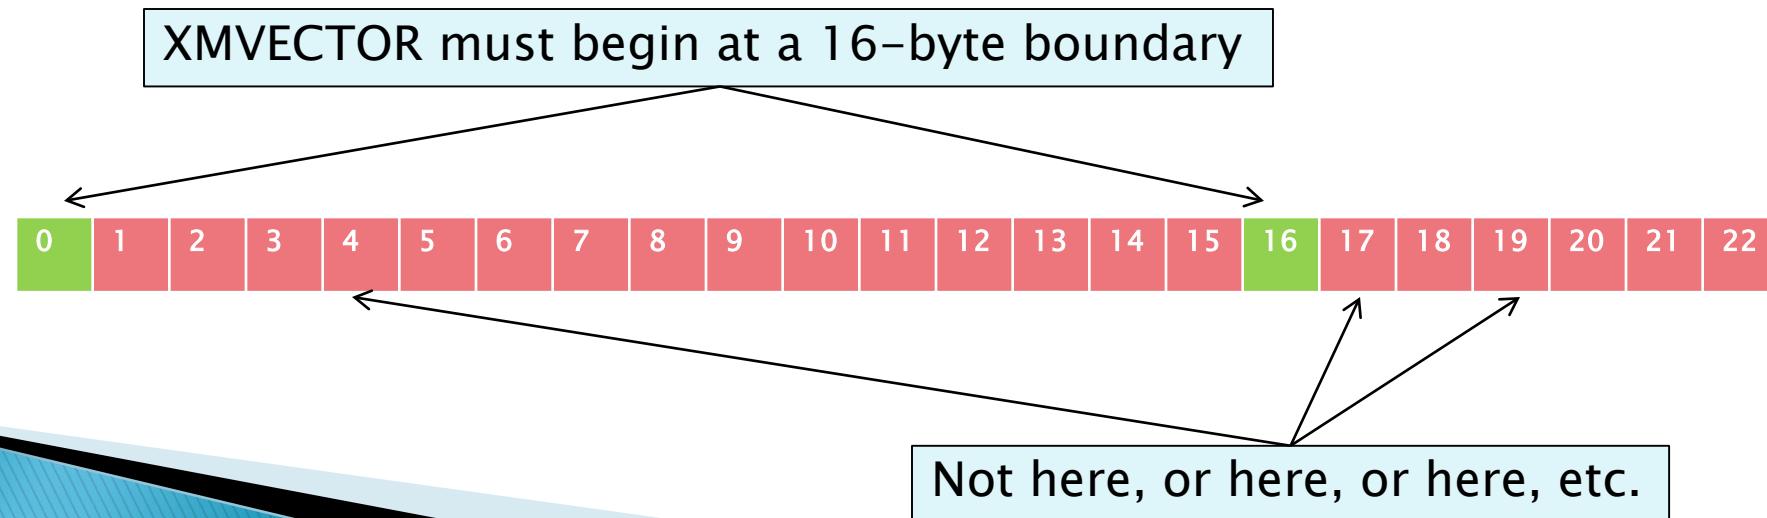
```
void MoveAbsolute(float x, float y, float z)
{
    // currentPosition is an XMFLOAT3
    // Just 3 adds? Probably fast enough to do manually
    currentPosition.x += x;
    currentPosition.y += y;
    currentPosition.z += z;
}
```

But...why?

This seems like a lot

Math data types – Issues

- ▶ These types are optimized for SIMD registers
 - Up to 4 of the same calculation at a time on the CPU
- ▶ Must be aligned in memory
 - Must not cross 16-byte boundaries

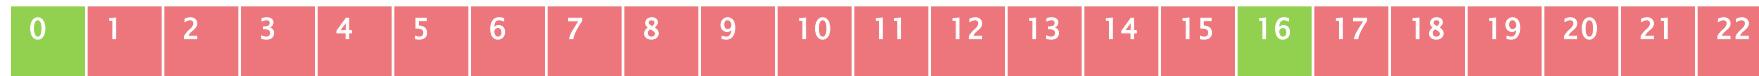


Member variables are NOT always aligned!

- ▶ Don't use Math Types as member variables!
 - Unless you can ensure they're aligned in memory
 - Might need a custom heap allocator
- ▶ Compiler does align *local variables* correctly
 - Variables on the stack
 - Core types are fine as *temporary local variables*
 - This is their intended use!

Alignment problem – Example

```
struct MyData  
{  
    char otherData;    // 1 byte  
    XMVECTOR position;  
}
```



If your struct starts here...

Its XMVECTOR could start further in memory.
Problem!

That seems like more work!

- ▶ If everything uses XMVECTOR or XMMATRIX...
- ▶ Why even bother with XMFLOAT types?
 - XMVECTOR/XMMATRIX alignment is **required**
 - Otherwise the math literally doesn't work right
 - If you can't ensure that, you **WILL** have problems
- ▶ This *is* slightly more work
 - For the *programmer*
 - The trade-off is performance

General work flow

- ▶ Define member variables using XMFLOATs
- ▶ “Load” into XMVECTOR/XMMATRIX before a calculation
- ▶ Perform all necessary operations
 - Do *as much as you can* at this step
 - For this particular calculation
- ▶ “Store” result(s) back into XMFLOATs when done

Class/struct members - Incorrect

```
class GameObject
{
private:
    float otherData;
    XMVECTOR position;
    XMMATRIX worldMat;
}
```

```
struct MyVertex
{
    float otherData;
    XMVECTOR position;
    XMVECTOR color;
    XMVECTOR uv;
}
```

DO NOT DO THIS!

DO NOT DO THIS EITHER!

Correct!

```
class GameObject
{
private:
    float otherData;
    XMFLOAT3 position;
    XMFLOAT4X4 worldMat;
}
```

```
struct MyVertex
{
    float otherData;
    XMFLOAT3 position;
    XMFLOAT4 color;
    XMFLOAT2 uv;
}
```

Much better!

Also much better!

This seems complicated

- ▶ It IS more complicated than other libraries
 - Gain performance
 - Lose simplicity
- ▶ SIMD-optimized math
 - Fast! Up to 4x speed-up of your math!
 - But requires extra work on your part
- ▶ Try it out
 - Not bad once you get the hang of it

What about 64-bit allocations?

- ▶ Generally already aligned to 16-byte boundary
 - So theoretically, the problem goes away, right?
 - We can just use XMVECTOR/XMMATRIX everywhere?
- ▶ Sure...
 - Unless you're using STL
 - Or another library that allocates memory for you
 - So you need to write your own memory allocator
- ▶ (Or just use *storage* types as described)

DirectX Math (CPU) vs. HLSL (GPU)

Different layouts for different architectures

Note on matrix storage

- ▶ DirectX Math vs. HLSL (shaders)
- ▶ Store their matrices differently in memory
- ▶ Due to SIMD optimizations on CPU
- ▶ DirectX Math: **Row-major** storage
- ▶ HLSL Shaders: **Column-major** storage

Storage: Row vs. column

- ▶ Row-major storage (CPU)
 - Layout in linear memory:
 - **abcdefghijklmnop**

a	b	c	d
e	f	g	h
i	j	k	l
m	n	o	p

- ▶ Column-major storage (GPU)
 - Layout in linear memory:
 - **aeimbjfncgkodhlp**

a	e	i	m
b	f	j	n
c	g	k	o
d	h	l	p

How does storage order affect us?

- ▶ CPU writes data to GPU memory
- ▶ GPU (shaders) use data expecting particular layout

CPU writes this:

abcdefghijklmnp

GPU expects this:

aeimbjfncgkodhlp

Uh-oh

a~~b~~i~~c~~m~~d~~f~~e~~g~~f~~h~~g~~i~~g~~j~~k~~n~~g~~c~~h~~o~~d~~k~~h~~l~~p~~

Options for row vs. column?

- ▶ **Transpose** your all of your matrices
 - Call XMMatrixTranspose() before storing them
- ▶ **Force** shader matrices to be row major
 - Define each matrix with “row_major” keyword
- ▶ **Force** entire shaders to be row major
 - `#pragma pack_matrix(row_major)`
 - Put at the top of each shader
- ▶ **Rearrange** your matrix math in your shader

Which is best?

- ▶ I suggest **rearranging matrix math** in shaders
- ▶ Simple operation
 - Less error-prone than other options
- ▶ Rerranged shader matrix example:

```
// This code will eventually appear in your Vertex Shader  
// We want World * View * Proj, but must multiply “backwards”  
matrix wvp = mul(proj, mul(view, world));
```