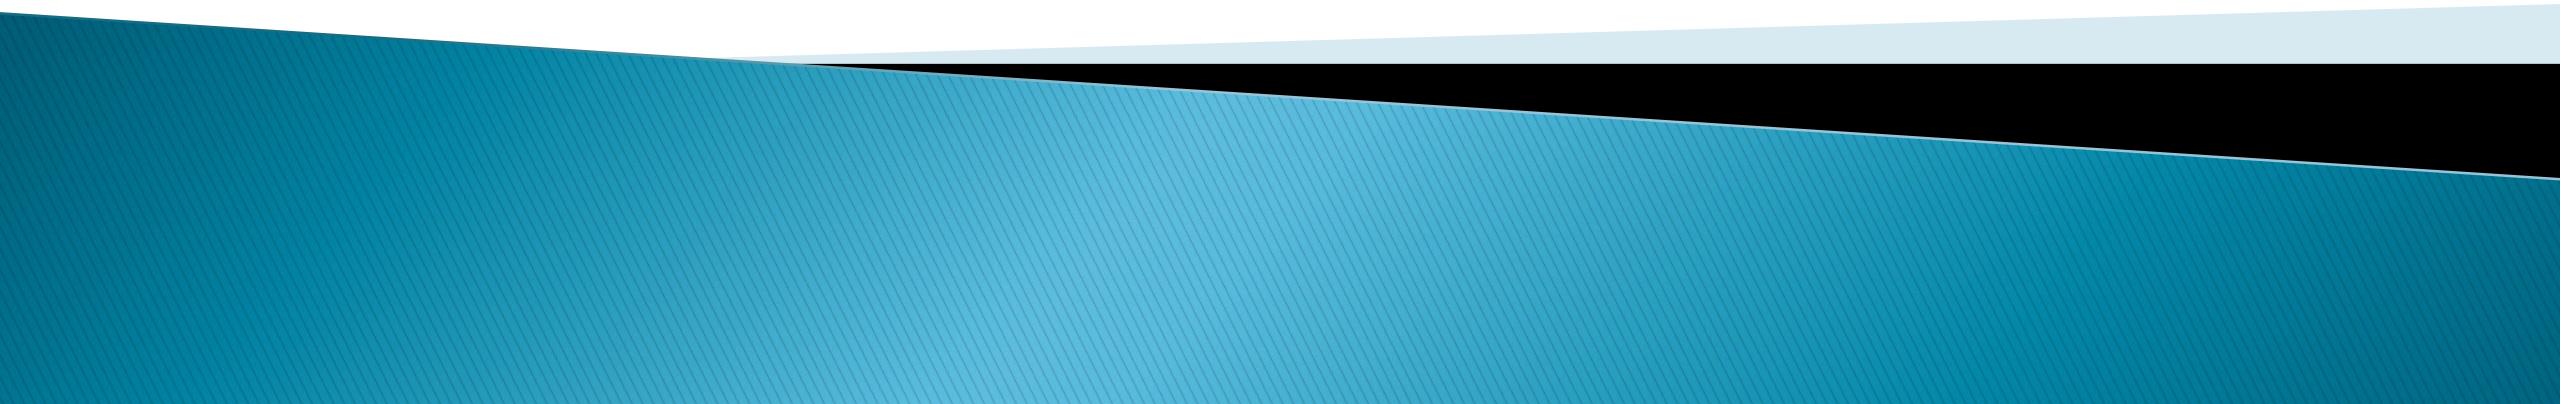


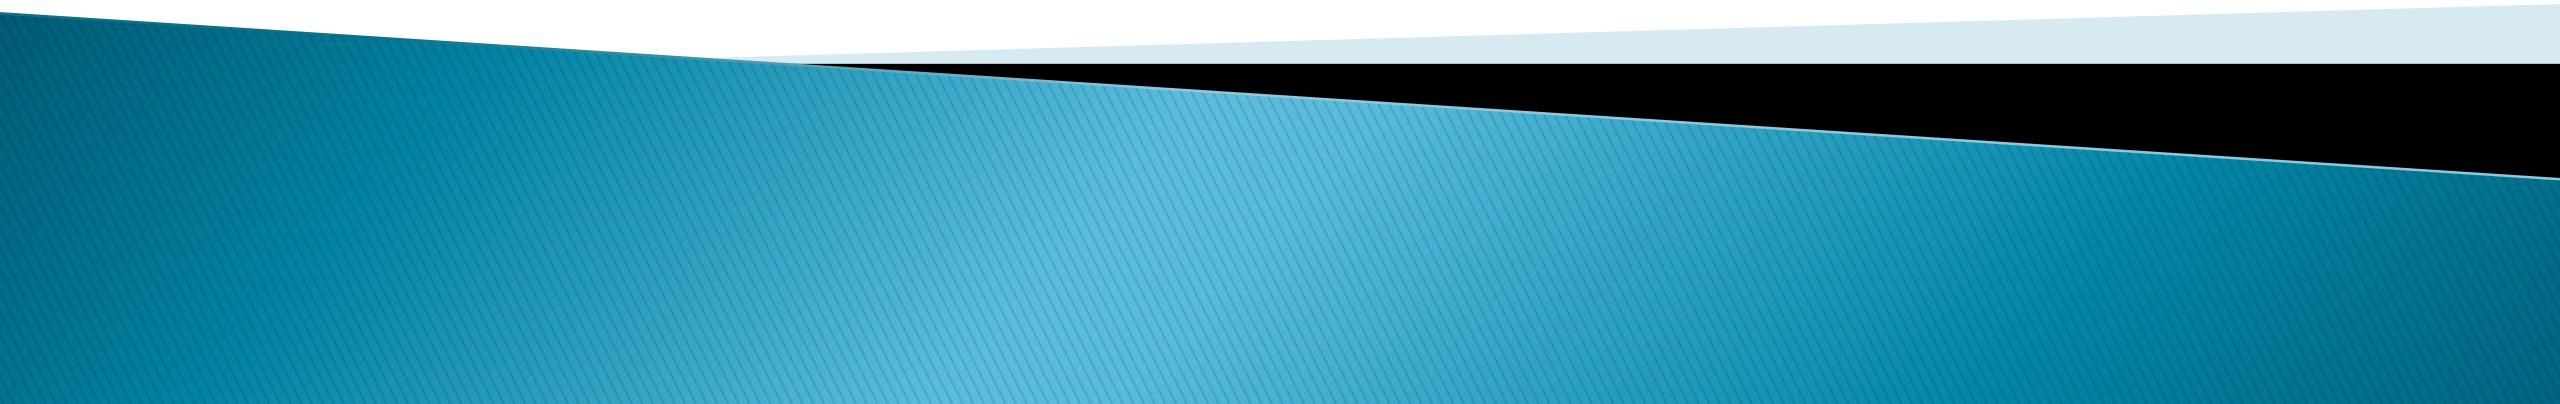
# HLSL In-Depth



# Topics

- ▶ Data types & operators
- ▶ Registers & semantics
- ▶ Flow control
- ▶ Intrinsic functions
- ▶ Custom functions
- ▶ Examples

# Data Types & Operators



# Major Data Types in HLSL

- ▶ **Scalars:** One-component
- ▶ **Vectors:** 1 to 4-component
- ▶ **Matrices:** 1 to 16-component (1x1 up to 4x4)
- ▶ **Other types:**
  - Textures
  - Samplers
  - User-defined Structs

# Major Scalar Types

- ▶ **bool**: true/false
- ▶ **int**: 32-bit signed integer
- ▶ **uint**: 32-bit unsigned integer
- ▶ **float**: 32-bit floating point value
  
- ▶ Other advanced types exist
  - Above types are good enough in most cases

[Scalar Types – MSDN](#)

[Scalar Literal Grammer – MSDN](#)

# Vector Types

- ▶ Contains between 1 and 4 scalar components
  - Yes, HSL supports 1-component vectors
  - Yes, they're basically scalars
- ▶ All components must be of the same type
- ▶ Can be declared in multiple ways
  - Built-in type declaration
  - Template (generic) declaration

# Vector Type Declaration

- ▶ Built-in type declaration:
  - [scalar type][number of components] varName;
  - `float3 direction;`
  - `int4 otherData;`
- ▶ Template (generic) declaration:
  - `vector<type, number of components> varName;`
  - `vector<float, 3> direction;`
  - `vector<int, 4> otherData;`

# Vector Initialization Examples

```
// These are all valid
```

```
float4 color = float4(1.0f, 0.5f, 0.0f, 1.0f);
```

```
float3 direction = { 1.0f, 0.0f, 0.0f };
```

```
int1 dumbVector = 6;
```

```
vector<int, 2> twoInts = { 5, 10 };
```

# Vector Components

- ▶ A vector's components may be accessed with:
  - r, g, b, a
  - x, y, z, w
- ▶ Accessing data in colors (or ANY vector):

```
float4 color      = float4(1,0,0,1);
float redAmount   = color.r;    // or color.x
float greenAmount = color.y;   // or color.g
```

# Swizzling

- ▶ Syntactical shortcut for rearranging vector components

```
float4 color      = float4(1.0f, 0.5f, 0.1f, 0.0f);
float4 foSwizzle = color.xxyx; // 1, 1, 0.5, 1
float4 worksFine = color.agbb; // 0, 0.5, 0.1, 0.1
float4 error      = color.xbya; // Can't mix notations
```

# Swizzling

- ▶ If fewer components than necessary are specified, the last one is repeated
- ▶ When a float4 is expected:
  - .xy really means .xyyy
  - .z really means .zzzz
  - .xxx really means .xxxx
- ▶ Code examples:
  - `float4 color = float4(1, 0.5f, 0.25f, 0);`
  - `float4 allZ = color.z; // 0.25, 0.25, 0.25, 0.25`
  - `float4 someY = color.xy; // 1, 0.5, 0.5, 0.5`

# Matrix Types

- ▶ Contain between 1 and 16 scalar components
  - Yes, HLSL supports  $1 \times 1$  matrices
  - Yes, they're basically scalars
- ▶ All components must be of the same type
- ▶ Can be declared in multiple ways
  - Built-in type declaration
  - Template (generic) declaration

# Matrix Type Declaration

- ▶ Built-in type declaration:
  - [scalar type][rows]x[columns] varName;
  - `float4x4 worldMatrix;`
  - `int3x1 otherData;`
- ▶ Template (generic) declaration:
  - `matrix<type, rows, columns> varName;`
  - `matrix<float, 4, 4> worldMatrix;`
  - `matrix<int, 3, 1> otherData;`
  - `matrix viewMatrix; // Shortcut for float4x4`

# Matrix Initialization Examples

```
// These are all valid
```

```
float2x2 mat = float2x2(1.0f, 0.0f, 1.0f, 0.0f);
```

```
float2x2 mat2 = { 1.0f, 0.5f, 1.0f, 2.0f };
```

```
matrix<float, 2, 2> mat3 = { 1.0f, 0.0f,
                               0.0f, 1.0f };
```

# Matrix Casting

- ▶ Sometimes you want a portion of a matrix
  - You can cast to the portion you want

```
// Declare a 4x4 matrix (and probably
// give it some data)
float4x4 world;

// Get the top left 3x3 portion
float3x3 topLeftPart = (float3x3)world;
```

# User-Defined Structures

- ▶ Made up of one or more member variables
  - Similar to structs in other languages
- ▶ Useful for passing multiple pieces of data between shader stages
  - Sending data from Vertex Shader to Pixel Shader
  - Returning multiple colors from Pixel Shader
- ▶ Useful for defining shader input data too

# Struct Examples

```
struct DirectionalLight
{
    float4 Color;
    float3 Direction;
};

struct Material
{
    float4 Color;
    float2 TextureOffset;
    float2 TextureScale;
};
```

# HLSL Operators

- ▶ Basic Math: +, -, \*, /, %
- ▶ Assignment: =, +=, -=, \*=, /=, %=
- ▶ Boolean: &&, ||, ?:
- ▶ Comparison: <, >, ==, !=, <=, >=
- ▶ Pre/postfix: ++, --
- ▶ Unary: !, -, +

[HLSL Operator MSDN Reference](#)

# More HLSL Operators

- ▶ Bitwise:            `~, <<, >>, &, |, ^`  
                        `<<=, >>=, &=, |=, ^=`
- ▶ Array:             `[index]`
- ▶ Cast:              `(type)`
- ▶ Structure  
member:             `.`

# Per-Component Operators

- ▶ Many operators are per-component
  - Operator occurs independently for each component in a multi-component variable (vector or matrix)
- ▶ Examples:
  - `int4 data = int4(1, 2, 3, 4);`
  - `int4 result = data * 5; // -> (5, 10, 15, 20)`
  - `float4 red = float4(1, 0, 0, 1);`
  - `float4 darkRed = red * 0.5f;`

# Optimizing Per-Component Math

- ▶ The following line can be optimized:

```
float4 answer = float4(1,1,1,1) * 5 * 10;
```

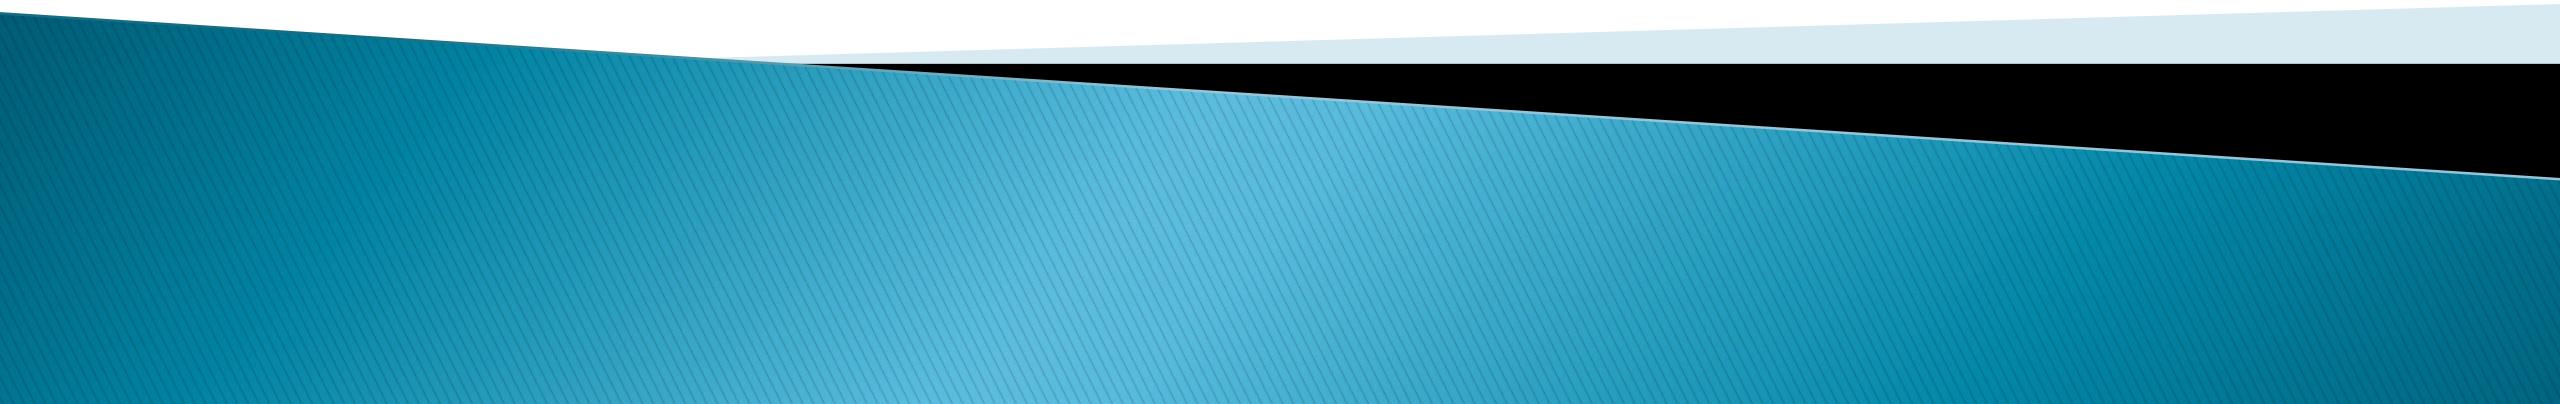
- Results in 8 multiplications

- ▶ Optimized version:

```
float4 answer = float4(1,1,1,1) * (5 * 10);
```

- Results in 5 multiplications

# Registers & Semantics



# Registers

- ▶ Where data is stored during the execution of a given shader
  - Similar to registers in a CPU
- ▶ Individual items (constant buffers, samplers, etc) are *always* bound to registers
- ▶ Can optionally be bound to specific registers
  - Using the *register* keyword
  - Basically an index

# Why Use Registers?

- ▶ A shader could use multiple textures
- ▶ CPU must bind those textures to the shader
  - For example: PSSetShaderResources()
  - Requires an index (called StartSlot)
- ▶ How does the CPU connect a texture to a specific texture register in a shader?
  - *Register* keyword allows you to specify the index

# Which Items Have Registers?

- ▶ Major register prefix & corresponding type:
  - b – constant buffer registers
  - t – texture registers
  - s – sampler registers
- ▶ Items must always use corresponding type
  - Or shader won't compile
  - Example: Textures must always use “t” prefix
- ▶ Register keyword allows you to specify the index of each item defined in the shader

# Register Example

```
// Registers require a prefix and an index
```

```
// Samplers - Must use “s” prefix
sampler LinearWrap : register(s0);
sampler PointClamp : register(s1);
```

```
// Textures - Must use “t” prefix
```

```
Texture2D BaseColor : register(t0);
Texture2D NormalMap : register(t1);
Texture2D SpecMap   : register(t2);
```

# Shader Parameter Semantics

- ▶ Semantics are strings attached to shader input or output parameter variables
- ▶ Specifies the intended usage of a parameter
- ▶ All variables passed between shader stages are required to have semantics
  - Including variables in input or output structs
- ▶ [Semantics MSDN Reference](#)

# Vertex/Pixel Shader Semantics

- ▶ Not all semantics are recognized everywhere
- ▶ Different shaders support different semantics
  - Different semantics for input & output too
  - See link on previous slide for specifics
- ▶ Semantic names can have an optional index
  - TEXCOORD – A texture coordinate (uv coord)
  - TEXCOORD0 – Same as TEXCOORD
  - TEXCOORD1 – Another texture coordinate

# Semantic Examples

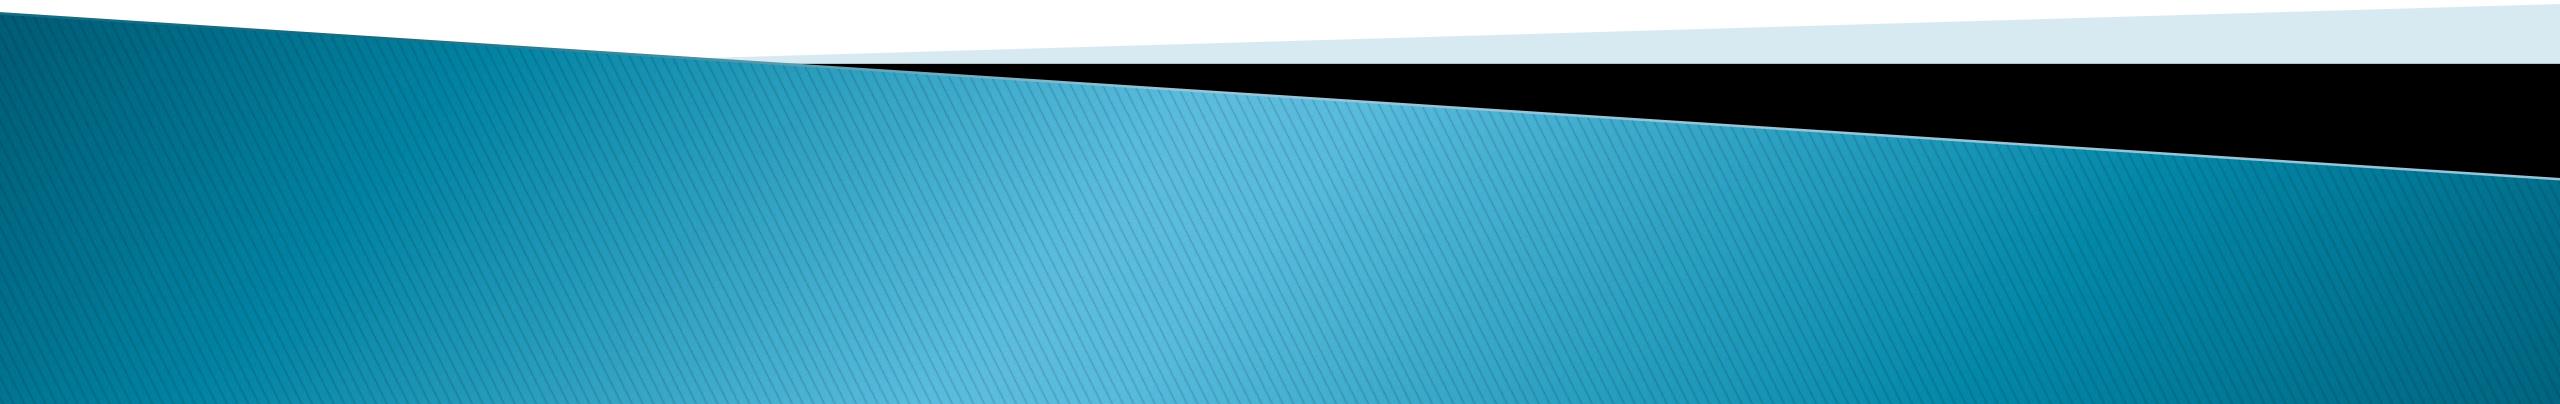
- ▶ Many semantics have existed since DirectX 9
  - POSITION
  - COLOR
  - TEXCOORD
  - NORMAL
- ▶ Some semantics are DX 10 and newer only
  - Example: SV\_Position
  - Called “System Value” semantics
  - Begin with SV\_ (so there are no naming conflicts)

# Semantic Example

```
struct VSInput
{
    float3 pos : POSITION;
    float2 uv : TEXCOORD;
};

struct VSOutput
{
    float4 pos : SV_POSITION; // DX10 and above
    float2 uv : TEXCOORD0;
    float other : TEXCOORD1;
};
```

# Flow Control



# Flow Control

- ▶ Determines which block of statements to execute next
- ▶ HLSL supports standard conditional and loop control statements
  - if / else if / else
  - switch / case / default
  - for / while / do-while
  - break / continue
- ▶ [Flow control MSDN Reference](#)

# Branching

- ▶ Sequential statements are executed in order
- ▶ Unless a *branch* statement is encountered
  - If statements
  - Loop conditions
- ▶ Branching causes different statements to potentially execute based on a condition

# Branching Performance

- ▶ Potential performance issues!
- ▶ Normal code execution:
  - Future instructions are preemptively fetched
  - Basically, upcoming code is being “prepared” early
- ▶ When a branch is encountered:
  - Which statement comes next?
  - Not known until branch is evaluated!
  - Processor must wait until current code is done before “preparing” next set of instructions

# CPU vs. GPU Branching

- ▶ CPU's have branch predictors
  - “Guesses” which branch might be taken
  - Begins fetch and speculative execution
  - If guess was correct – speed up!
- ▶ GPU's don't have branch predictors
  - Branching can potentially cause performance issues
  - Unless branch cost is cheaper than code execution

# When To Branch?

- ▶ Branch when:
  - Cost of branching is less than cost of code block
  - Branch potentially skips most/all of shader
  - Shader performance isn't a high priority
  - Learning HLSL
- ▶ Don't branch when:
  - Code can be rewritten without a branch
  - Branching can be replaced by built-in functions
- ▶ [Shader Optimization MSDN Reference](#)

# Branching & Loops

- ▶ Loops contain conditions
  - Determine when the loop ends
  - This is a branch!
  - Potential performance issue again!
  
- ▶ HLSL compiler can *unroll* some loops
  - Removes part or all of the loop itself
  - Replaces it with copies of the code block

# Unrolling Loop Example

- ▶ Loop version:

```
for( int i = 0; i < 3; i++ )  
    DoSomething(i);
```

- ▶ Unrolled version:

```
DoSomething(0);  
DoSomething(1);  
DoSomething(2);
```

# HLSL Loop Attributes

- ▶ Compiler can unroll some loops automatically
- ▶ Unrolling isn't always desirable
  - Unrolling results in more instructions
  - Can potentially make shaders “too long” to compile
- ▶ Optional loop attributes control compilation:
  - [unroll] – Unrolls loop (if possible)
  - [loop] – Ensures loop actually loops

# Loop Attribute Examples

[unroll]

```
for( int i = 0; i < 10; i++ )  
    DoSomething(i);
```

[loop]

```
for( int i = 0; i < 10; i++ )  
    DoSomething(i);
```

- ▶ [For Loop MSDN Reference](#)

# Intrinsic Functions

# Intrinsic Functions

- ▶ HLSL has many built-in (intrinsic) functions
  - Mostly math-related
  - Some only work in certain shader types
- ▶ Some functions are per-component
  - They can be used with scalars, vectors or matrices
  - Will perform their function on the whole input
  - Output type will match input type
- ▶ [List of HLSL Intrinsic Functions](#)

# Major Math Functions

- ▶ Basic math:

- abs, floor, ceil, frac, sign
- round, trunc, sqrt, pow

- ▶ Trig functions:

- sin, cos, tan, sincos
- asin, acos, atan, atan2

# Vector & Matrix Functions

- ▶ Vector functions:

- dot, cross
- length, distance
- normalize

- ▶ Matrix functions:

- mul

# Clamping Values

- ▶ `clamp( data, min, max )`
  - Clamps *data* between *min* and *max*
  - If *data* < *min*, it becomes *min*
  - If *data* > *max*, it becomes *max*
  - Otherwise it doesn't change
  
- ▶ `saturate( data )`
  - Clamps *data* between 0.0 and 1.0
  - Same as `clamp( data, 0.0, 1.0 )`

# Interpolation

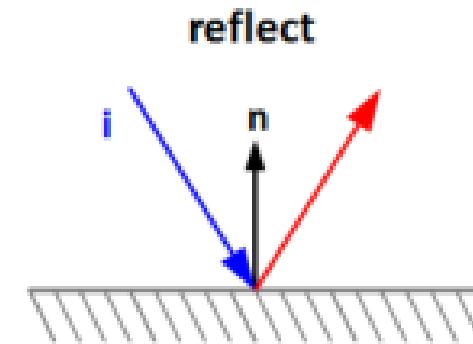
- ▶ `lerp( v1, v2, percent )`
  - Interpolates between v1 and v2 based on percent
  - Basically this function:  $v1 + percent * (v2 - v1)$
  - A percent of 0.0 gives v1
  - A percent of 1.0 gives v2
  
- ▶ `smoothstep( min, max, data )`
  - A smooth Hermite interpolation between 0 & 1
  - If data < min, returns 0
  - If data > max, returns 1
  - Otherwise returns a value between 0 and 1

# Choices

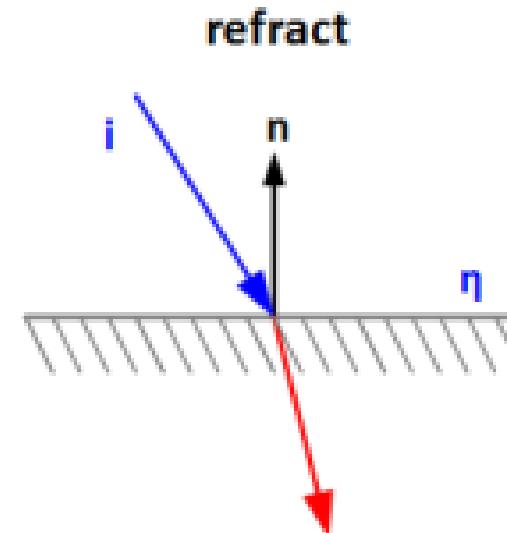
- ▶ `step( v1, v2 )`
  - Returns either 0 or 1
  - Returns 1 if  $v1 \geq v2$
  - Returns 0 otherwise
- ▶ `max( v1, v2 )`
  - Returns either  $v1$  or  $v2$ , whichever is greater
- ▶ `min( v1, v2 )`
  - Returns either  $v1$  or  $v2$ , whichever is smaller

# Vector Helpers

- ▶ `reflect( incident, normal )`
  - **Incident** is the incoming vector
  - **Normal** is the surface normal



- ▶ `refract( incident, normal, refractionIndex )`
  - **Indicent** is the incoming vector
  - **Normal** is surface normal
  - **RefractionIndex** defines how much the vector “bends”



# Per-Component Comparisons

- ▶ HLSL has 2 functions that look for non-zero values in a multi-component variable
- ▶ `all( data )`
  - Returns true if *a//*components of *data* are non-zero
- ▶ `any( data )`
  - Returns true if *any* components of *data* are non-zero

# Per-Component Comparisons

- ▶ Comparison operators can return multiple booleans
  - If both sides are vectors, for example

```
float4 big    = float4(5,5,5,0);
```

```
float4 small = float4(1,1,1,9);
```

```
bool4 result = big > small;
```

```
// Result is (true, true, true, false)
```

# Per-Component Comparisons

- ▶ But if statements only work with single booleans

```
float4 big    = float4(5,5,5,0);
float4 small = float4(1,1,1,9);
if( big > small ) // ERROR: Won't compile!
```

```
// This works (but there is an easier way)
if( big.x > small.x && big.y > small.y &&
    big.z > small.z && big.w > small.w )
```

```
// Cleaner - All components must be true
if( all( big > small ) )
```

# Custom Functions

# Custom Functions

- ▶ Individual shaders are simply functions
  - Vertex, pixel, etc.
- ▶ Can write custom helper functions too
- ▶ Often helper methods are put in external files
  - Can be #include'd into other .hlsl files
- ▶ [Function MSDN Reference](#)

# Function Parameters

- ▶ Parameters are “in” params by default
  - Represent data being passed in to the function
  - Standard parameter behavior
- ▶ Parameters can also be “out” parameters
  - Represent data being passed back out
  - Allows a function to “return” more than one value
- ▶ Also supports “inout” parameters
  - Data coming in, which can change and is passed back out at the end of the function

# Function Performance

- ▶ All custom functions in HLSL are “inline”
  - Basically the function body is pasted into the shader
  - Rather than actually “calling” the function
- ▶ No performance hit when using functions!
  - No function call overhead since they’re inline
- ▶ Still adds to the instruction count
  - Could possibly generate a shader that is “too long”

# Example

A simple example of the pieces necessary  
for a basic directional light

# Example - Vertex Shader

# Example – VS Constant Buffer

```
// The name of the cbuffer is largely
// irrelevant. What matters is the
// register and the data inside!
cbuffer DataFromCPU : register(b0)
{
    matrix World;
    matrix WorldInverseTranspose;
    matrix View;
    matrix Projection;
}
```

# Example – VS Input Struct

```
// This struct represents the data of a
// single vertex in our vertex buffer.
// This must also match our Input Layout
struct VS_Input
{
    float3 localPosition : POSITION;
    float3 normal        : NORMAL;
    float2 uv             : TEXCOORD;
};
```

# Example – VS Output Struct

```
// Represents the data coming out of the
// Vertex Shader, which must also match
// the input of the Pixel Shader
struct VS_Output
{
    float4 screenPosition : SV_POSITION;
    float3 normal         : NORMAL;
    float2 uv              : TEXCOORD;
};
```

# Example – Vertex Shader Signature

```
// Beginning of the vertex shader itself.  
// Its return type is a struct.  
// Its parameter is also a struct.  
VS_Output main( VS_Input input )  
{  
    // Create eventual output variable  
VS_Output output;  
  
    // Continues on next slide...
```

# Example - VS - Screen Position

```
// Calculate World View Projection matrix
// (Backwards because of DXMath / HLSL transpose issue)
matrix wvp = mul(Projection, mul(View, World));

// Use final matrix to calculate the
// vertex's screen position
output.screenPosition =
    mul(wvp, float4(input.localPosition, 1.0f));

// Continues on next slide...
```

# Example - VS - Normal & UV

```
// Vertex normal must be transformed too
output.normal =
    mul((float3x3)WorldInverseTranspose, input.normal);

// Pass the uv through, since it comes
// from the vertex but isn't used here
output.uv = input.uv;

// End of vertex shader - return data
return output;
}
```

# Example – Pixel Shader

# Example – PS Input Struct

```
// Represents the input data of the Pixel Shader, which
// must match the output struct of the Vertex Shader
struct PS_Input
{
    float4 screenPosition : SV_POSITION;
    float3 normal         : NORMAL;
    float2 uv              : TEXCOORD;
};
```

# Example – PS Helper Struct

```
// Represents a single directional light
// with ambient and diffuse color
struct DirectionalLight
{
    float4 AmbientColor;
    float4 DiffuseColor;
    float3 Direction;
};

// Will be part of the constant buffer
```

# Example – PS Constant Buffer

```
// We need to get the light data from
// the CPU side, which means passing
// it in through a constant buffer
cbuffer PSDataFromCPU : register( b0 )
{
    DirectionalLight Light;
};
```

# Example - Other PS Variables

```
// The texture for the model
Texture2D DiffuseTexture : register(t0);

// Sampler which describes how to pull
// pixels from a texture
SamplerState TrilinearSampler : register(s0);
```

# Example – Pixel Shader Signature

```
// Beginning of Pixel Shader itself
// Return type is float4 (a color)
// Parameter is PS_Input struct
float4 main(PS_Input input) : SV_TARGET
{
    // Continues on next slide...
```

# Example – PS – Basic Lighting

```
// Normalize the input normal
float3 normal = normalize(input.normal);

// Invert the light's direction
float3 lightDir = -Light.Direction;

// Calculate the diffuse light factor
float lightAmount = dot(normal, lightDir);

// Clamp light factor between 0 - 1
lightAmount = saturate(lightAmount);
```

# Example – PS – Texture

```
// Sample the texture using the sampler  
// and the interpolated uv coords  
float4 textureColor =  
    DiffuseTexture.Sample(TrilinearSampler, input.uv);
```

# Example – PS – Final Color

```
// Put it all together! First figure
// out final diffuse light color
float4 lightColor = lightAmount * Light.DiffuseColor;

// Add in ambient light
lightColor += Light.AmbientColor;

// Return texture modulated by light
return textureColor * lightColor;
}
```