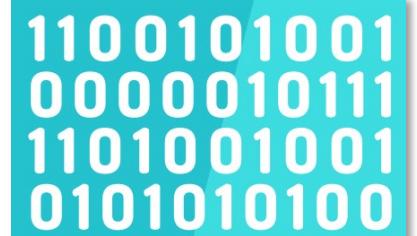


Graphics APIs, Direct3D & The Rendering Pipeline



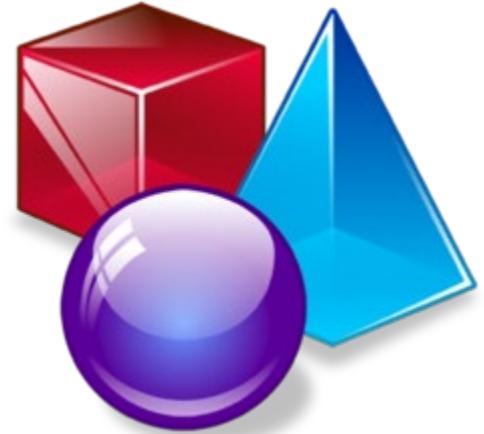
Terminology



1100101001
0000010111
1101001001
0101010100

- ▶ **Buffer** – A place to store a “chunk” of data
 - Generally holds sets of numbers
- ▶ **CPU-Side**
 - The code running on your CPU *or* the data in system RAM
 - Your C++ code and what it directly interacts with
- ▶ **GPU-Side**
 - Video card hardware, video RAM or shader code
 - Anything that happens on the video card
 - Many Direct3D constructs (buffers, textures, etc) are GPU-side

Terminology – Geometry



- ▶ **Vertex** – A single point
 - Often part of a larger geometric primitive
 - Can contain data in addition to a position in space
- ▶ **Geometric Primitive** – Simple geometric shapes
 - Points, lines & triangles
 - This is what the GPU draws
- ▶ **Mesh** – A unique set of geometry
 - Vertices and potentially indices
 - No transformation information
 - Can be shared among render-able entities

Terminology - Output buffers

- ▶ **Texture** – An image on the GPU
 - A “buffer of pixels”
- ▶ **Render Target**
 - A texture on the GPU that can hold rendering output
 - Any number of these can exist
- ▶ **Back Buffer** – Specific render target used for screen output
- ▶ **Depth Buffer** – Texture holding depth of each pixel rendered
 - Another “buffer of pixels”
 - These pixels hold a single depth value (not a color)
 - Allows us to render 3D objects in any order



Terminology – Rendering

▶ **Rendering Pipeline**

- Steps to render triangles
- The actual rendering process
- Performed by Direct3D/drivers/GPU



▶ **Fixed-Function Pipeline** – All stages are hardwired

- You choose pre-defined options
- Can't execute arbitrary code (no shaders)

▶ **Programmable Pipeline** – Supports executing arbitrary code

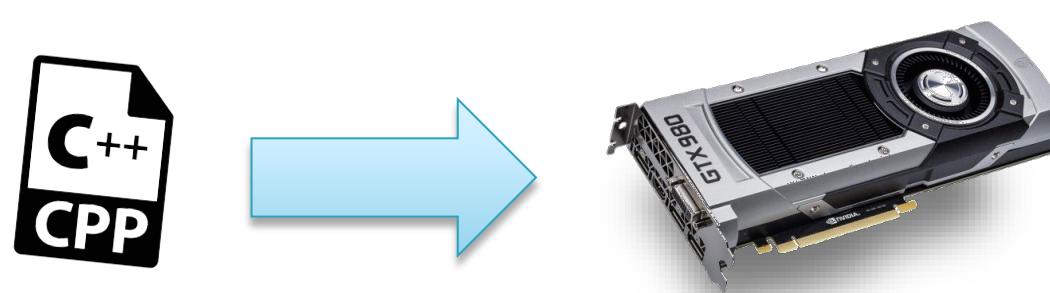
- Certain steps require custom code (shaders)
- Allows for far more variety and power

Graphics APIs



What's a Graphics API?

- ▶ C++ code library
- ▶ Allows an application to utilize built-in features of our graphics hardware (GPUs) to greatly speed up rendering



- ▶ Technically...
 - C++ Code → Graphics API → Graphics Drivers → GPU Hardware

Examples of Graphics APIs



Vulkan®

OpenGL®



Direct3D 11

- ▶ The Graphics API we'll be using this semester
 - Powerful
 - Higher-level
 - Many examples exist
 - Integrates well with our tools
- ▶ Created by Microsoft
- ▶ Used on PCs and Xbox
- ▶ One of several APIs in “DirectX”



Direct3D



DirectX vs D3D

- ▶ DirectX: Collection of APIs allowing “direct” hardware access
- ▶ Direct3D: Specific API for rendering with a GPU
- ▶ The two terms are often used interchangeably

Direct3D is...

- ▶ A library to interface with a GPU for rendering
- ▶ Not explicitly a *game engine* or a *game framework*
 - Unity is a game engine
 - MonoGame is a game framework
 - Both can use Direct3D under the hood
- ▶ Anything else we want, we must build ourselves!
 - Game loop? Window?
 - We need to code those!

Official helper libraries

- ▶ DirectX Math
 - CPU (C++) side, SIMD-friendly math library
 - Has structures/functions for 2D & 3D math
 - Comes with DirectX / Windows SDK
- ▶ DirectX Toolkit
 - Collection of D3D helper classes for common tasks
 - Created by Microsoft
 - Available on [GitHub](#) & NuGet
- ▶ More on these soon!

Components of Direct3D

And most other graphics APIs

D3D is object-oriented

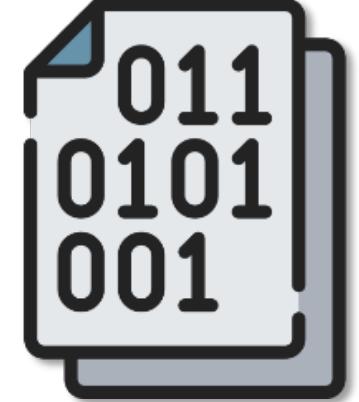
- ▶ All functions are accessed through objects
 - Very few global functions
 - OpenGL/Vulkan differ here!
- ▶ Main Direct3D C++ objects:
 - **Device**
 - **Device Context** (“Context” for short)
 - **Swap Chain**

Major components (of any graphics API)

- ▶ **Resources** – Data used during rendering
- ▶ **Graphics state** – Options for rendering
- ▶ **Issuing commands** – Change state, update resources, draw!

Resources

- ▶ Data used during the rendering process
 - Input buffers (vertices)
 - Shader variable data
 - Textures
- ▶ Resources live in GPU memory
- ▶ Graphics API creates and manages resources
 - D3D's **Device** object creates resources
- ▶ Our job is to *bind* (activate) resources prior to each draw call



Binding Resources

- ▶ We could have 100's or 1000's of resources in GPU memory
 - 250 vertex buffers
 - 122 index buffers
 - 7 shader variable buffers
 - 46 textures
- ▶ Which one(s) should be used for the next render?
- ▶ We specify the set of “active” resources by *binding* them
 - Call a function to make a resource active
 - Unbinds previously active resource

Graphics state

- ▶ The current set of all graphics options and bound resources
 - Used by the API/Drivers/GPU when drawing
 - D3D's **Context** object tracks graphics state
- ▶ Examples:
 - Which vertex buffer is active?
 - Which shaders are active?
 - What are the rasterizer options?
 - Where should rendering output be stored?
- ▶ All state persists until we change it



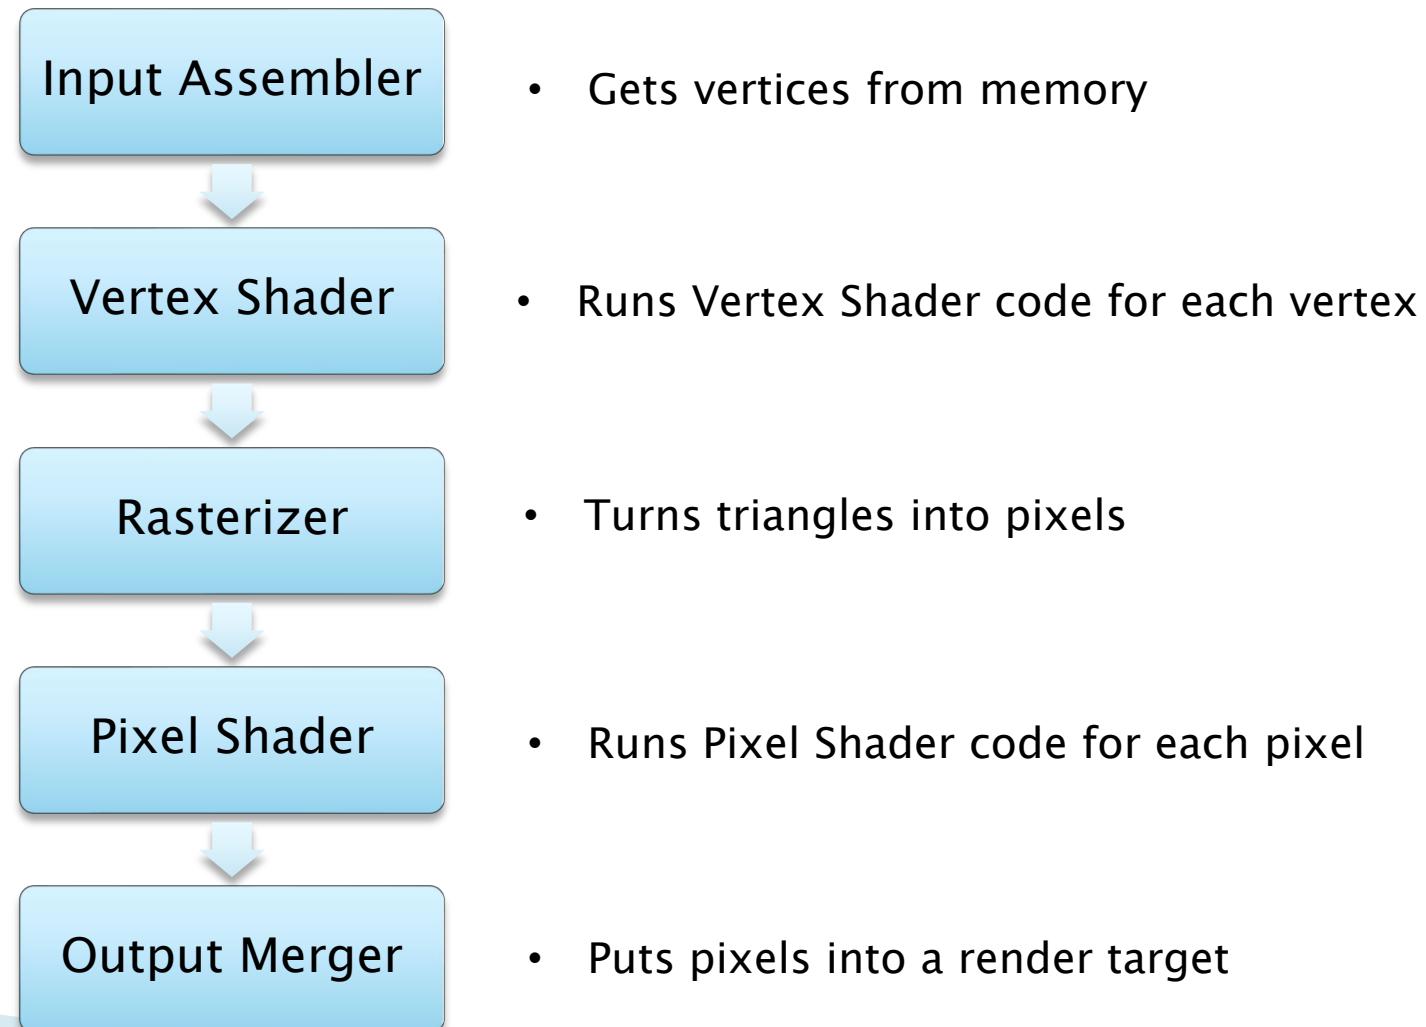
Issuing commands

- ▶ D3D's **Context** object has two broad categories of functions
- ▶ Functions to “prepare for rendering”
 - Change state
 - Update resources
- ▶ Functions to “launch GPU work”
 - Clear output buffers
 - Begin rendering process (“draw calls”)

Rendering Pipeline



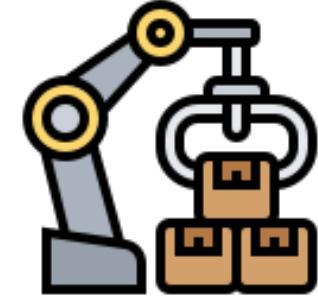
Rendering pipeline stages



Pipeline I/O

- ▶ **Input**
 - Vertices (required) & indices (optional)
 - Vertices & indices must be on GPU already
- ▶ **Output**
 - Pixel colors in a render target (an image)
 - Default target is the “back buffer” (screen)
- ▶ **Multi-step process**
 - Once started, all steps performed by Direct3D, drivers & GPU
 - Once started, we **do not** get results back in C++!

Stage 1: Input Assembler

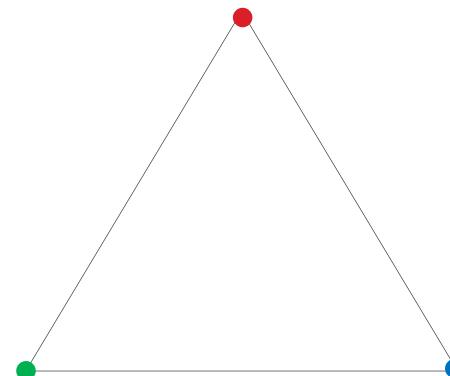


- ▶ Several tasks
 - Gathers information from the buffers
 - Assembles data into geometric primitives
 - Feeds primitives into the next stage
- ▶ Essentially: launches vert shader once per vertex
 - Each of which is processed in parallel (at the same time)
- ▶ Options for this stage include:
 - **Input layout**
 - **Primitive topology**

Data after Input Assembler

- Vertices assembled into primitives (triangles)

pos: 0, 0.5, 0
color: 1,0,0,1

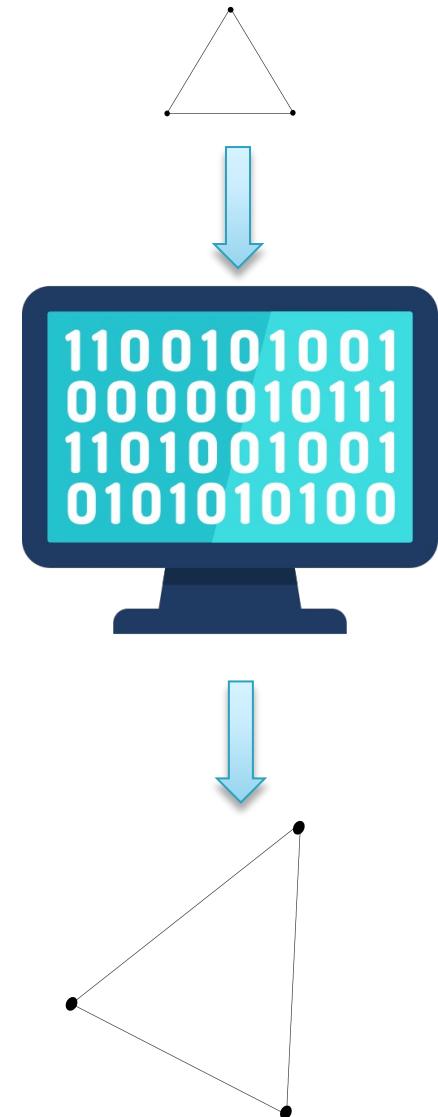


pos: -0.5, -0.5, 0
color: 0,1,0,1

pos: 0.5, -0.5, 0
color: 0,0,1,1

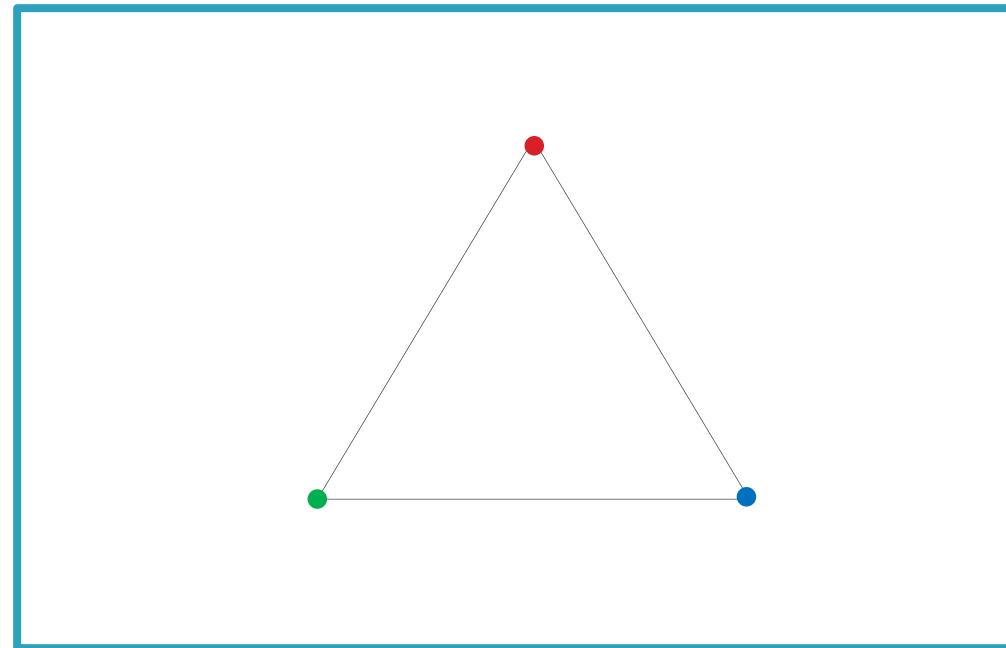
Stage 2: Vertex Shader

- ▶ Programmable and required
- ▶ Performs operations on a single vertex
 - Transformations, lighting, skinning (animation), etc.
 - Every vertex goes through this stage
- ▶ Main goal is to transform vertex positions:
 - Local to world space
 - World to view space
 - View to screen space



Data after Vertex Shader

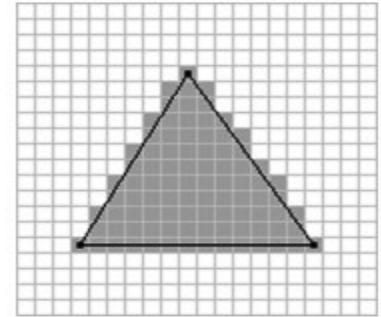
- ▶ Triangle is now in screen space
 - Exact location depends on transformation
 - And potentially camera matrices, too



Stage 3: Rasterizer

- ▶ Converts primitives to pixels
 - Feeds those pixels into the next stage
 - Also clips primitives that are not visible

- ▶ Interpolates vertex data into per-pixel data
 - UV coords, colors, normals, etc.
 - Takes data at the vertices and calculates an adjusted value for each pixel of the triangle

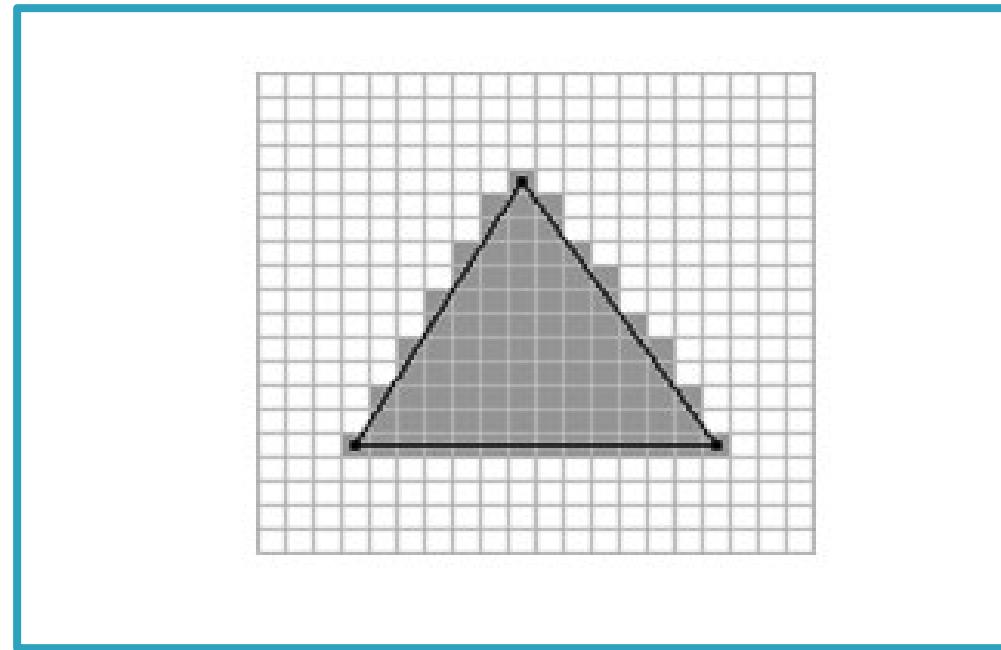


Rasterizer

- ▶ Automatically culls back-facing triangles
 - Triangles you “can’t see” are skipped
 - Optimization – assumes most models are “shells”
 - Can be customized
- ▶ Options for this stage include:
 - **Viewport** for rendering/clipping
 - **Fill mode** – wireframe or solid
 - **Cull mode** – triangle front, back or neither

Data after Rasterizer

- ▶ Triangle now rasterized



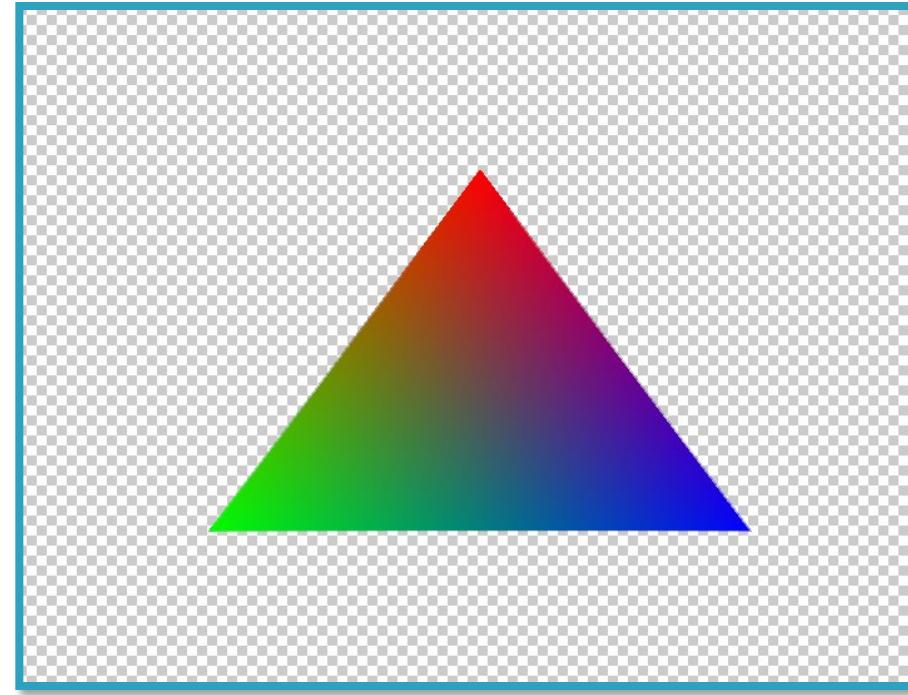
Stage 4: Pixel Shader

- ▶ Programmable and required
- ▶ Calculates final pixel color
- ▶ Where most special effects are calculated
 - Lighting, normal mapping, textures, etc.
- ▶ Return a color (4 numbers): R, G, B, A



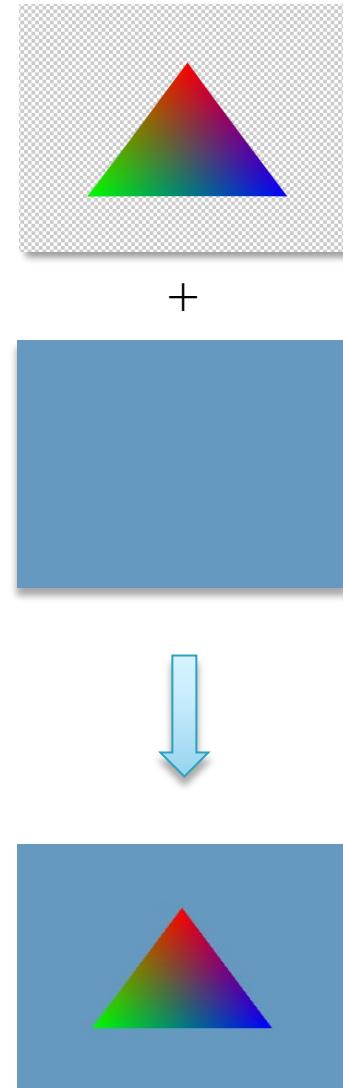
Data after Pixel Shader

- ▶ Just this triangle's pixels have been “shaded”



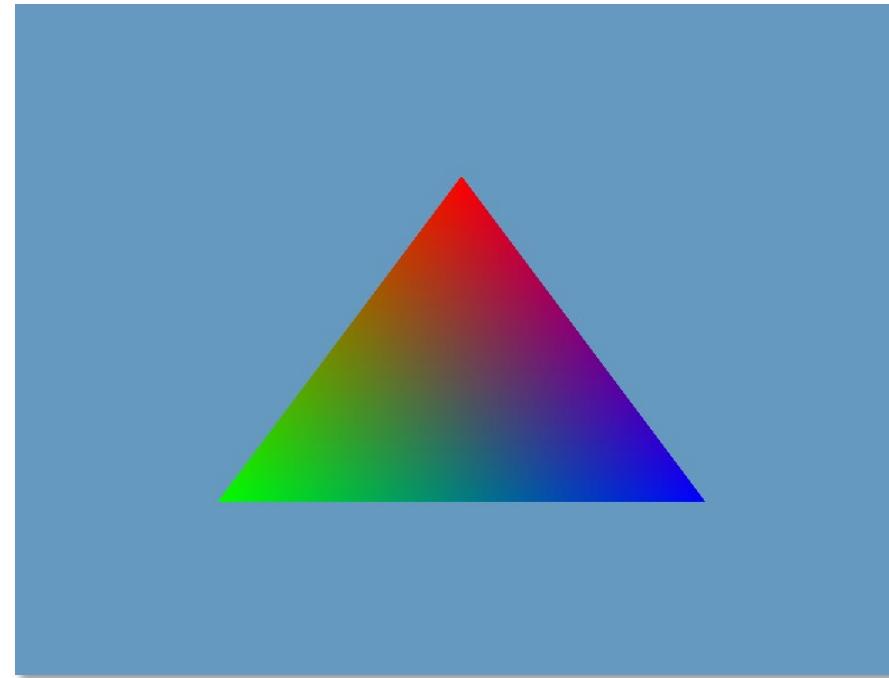
Stage 5: Output Merger

- ▶ Merges output data (pixels)
 - Handles alpha blending if necessary
 - Combines colors together based on settings
- ▶ Writes to depth buffer
 - Records “depth” (distance) of each pixel
 - Can discard pixels that are “behind” others
- ▶ Options for this stage include:
 - **Blending options** – How to combine colors
 - **Render targets** – Where to put pixels
 - **Depth buffer** – Can turn read/write on and off

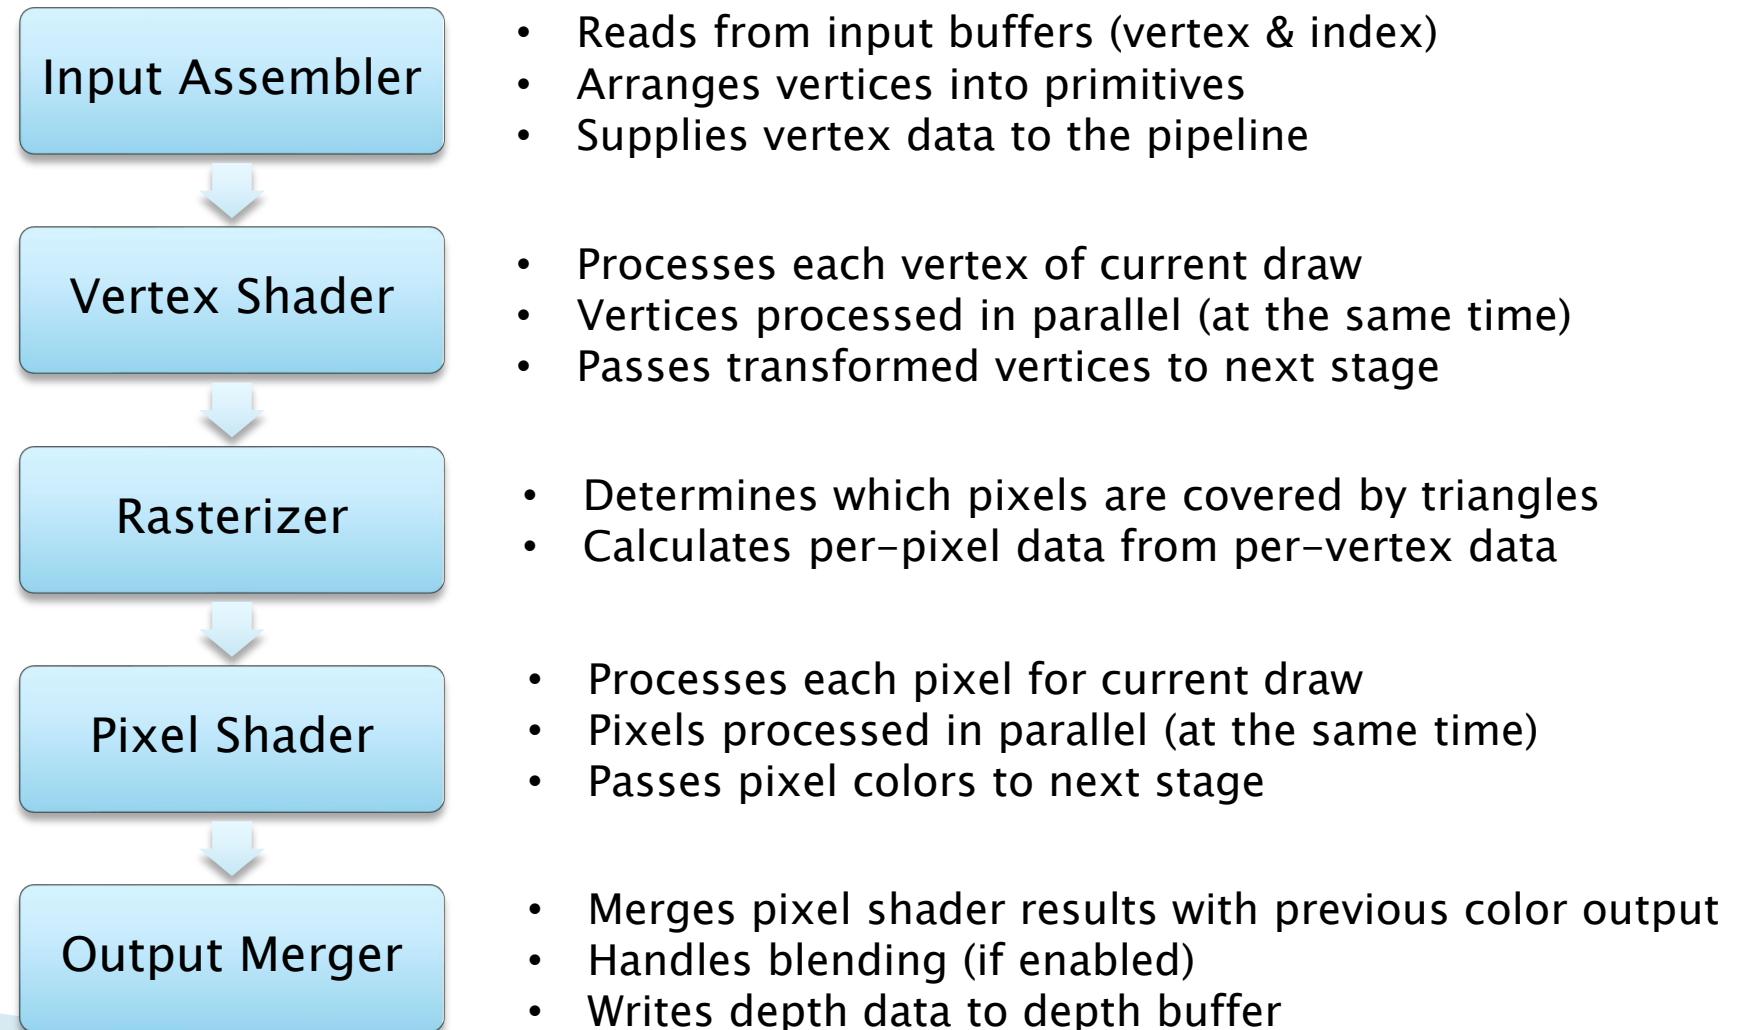


Data after Output Merger

- ▶ Pixels merged with existing background



Pipeline stages – Recap



Other Pipeline Details



Rendering Pipeline is asynchronous

- ▶ Draw calls are not blocking!
 - C++ does *not* wait for a draw call to complete
 - Pipeline starts
 - Takes time to complete
 - Our code continues queueing up commands for Direct3D
- ▶ Common to call Draw() many times before the first is finished
 - Rendering is usually $\frac{1}{2}$ – 1 frame behind game loop
- ▶ Part of Direct3D's job is to handle that synchronization for us

CPU-Side vs. GPU-Side

Our C++

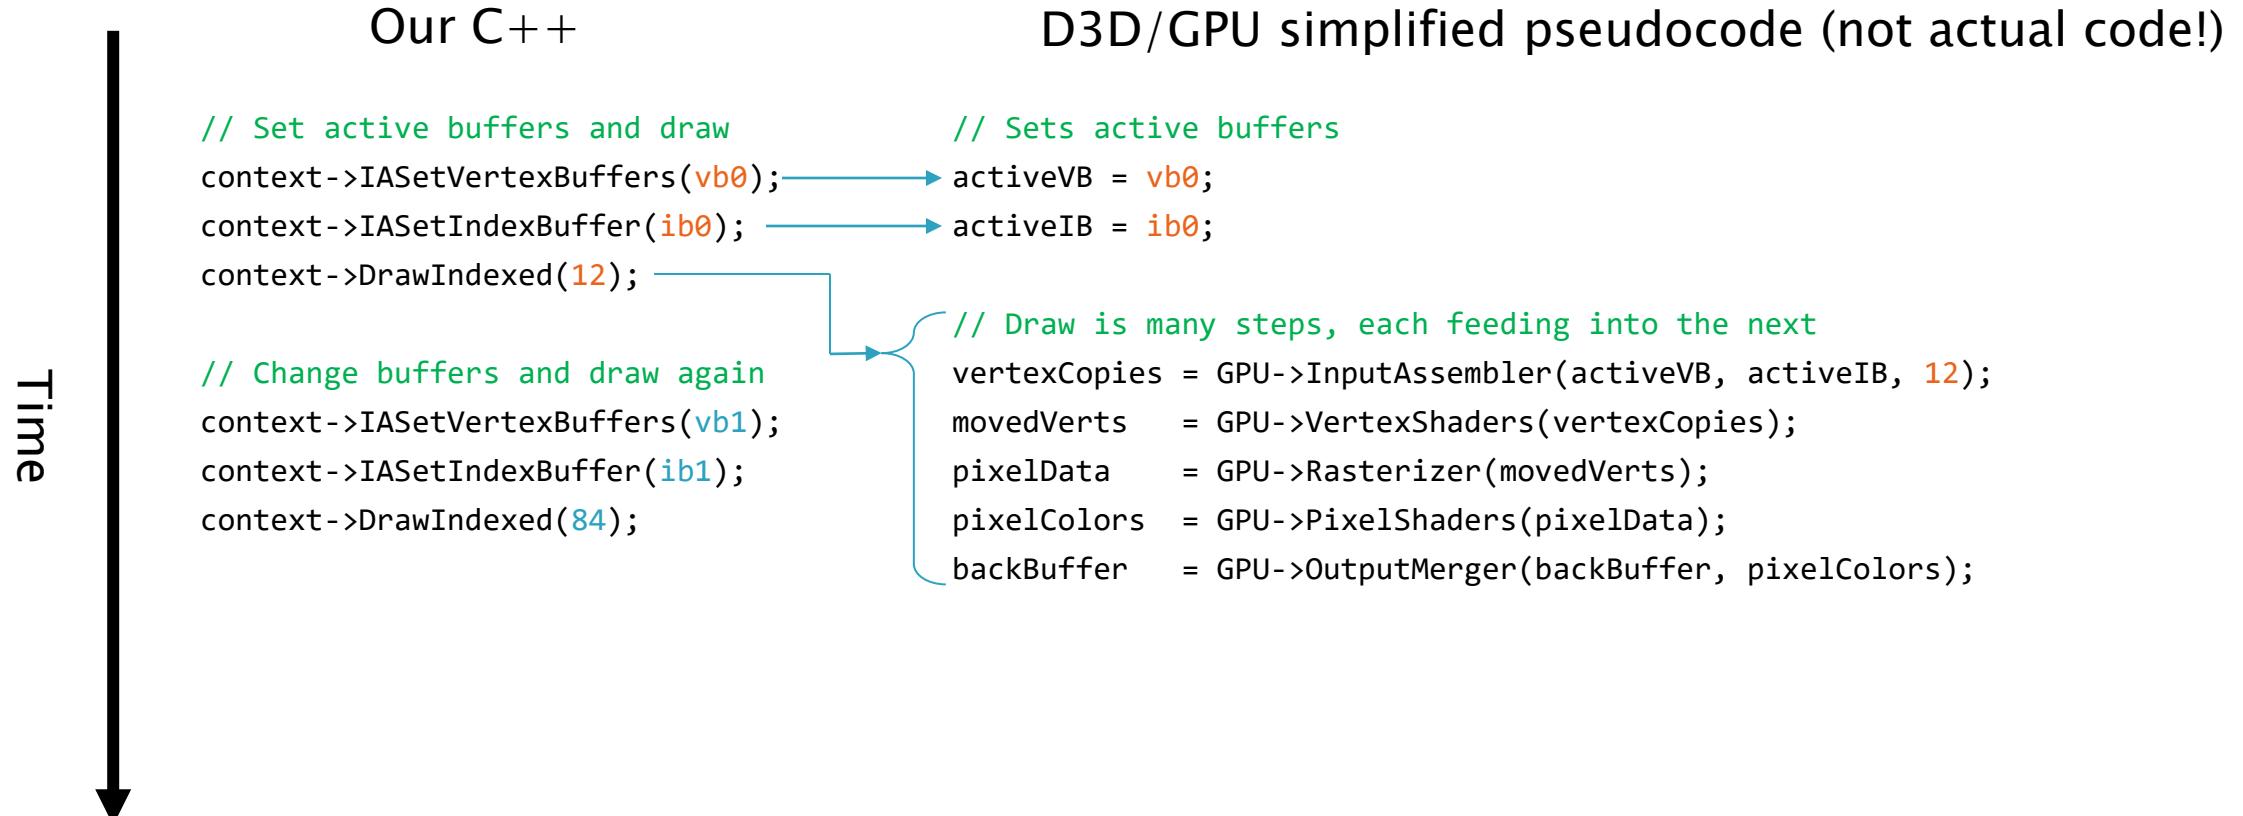
```
// Set active buffers and draw
context->IASetVertexBuffers(vb0);
context->IASetIndexBuffer(ib0);
context->DrawIndexed(12);
```

```
// Change buffers and draw again
context->IASetVertexBuffers(vb1);
context->IASetIndexBuffer(ib1);
context->DrawIndexed(84);
```

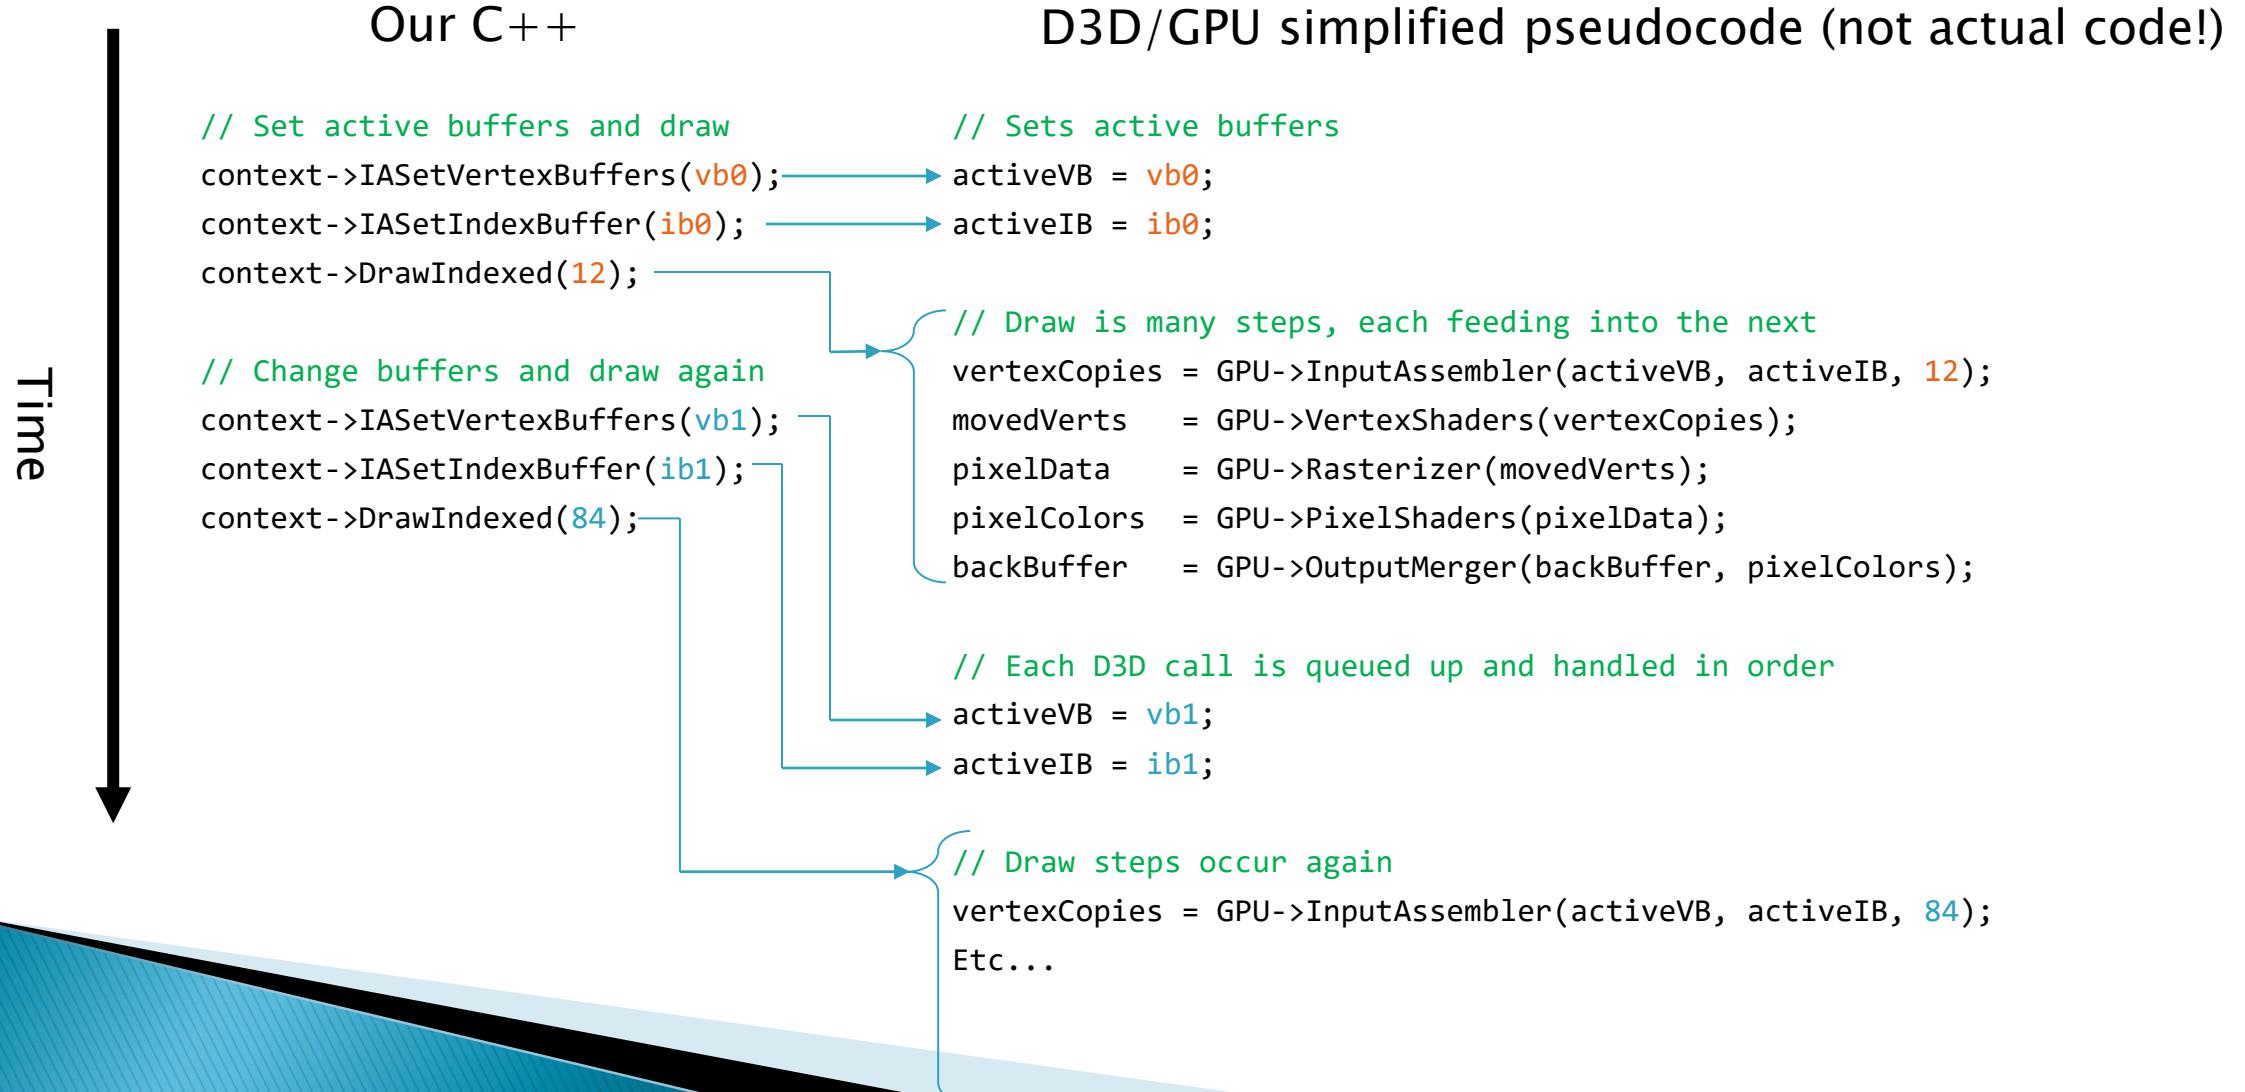
Time



CPU-Side vs. GPU-Side



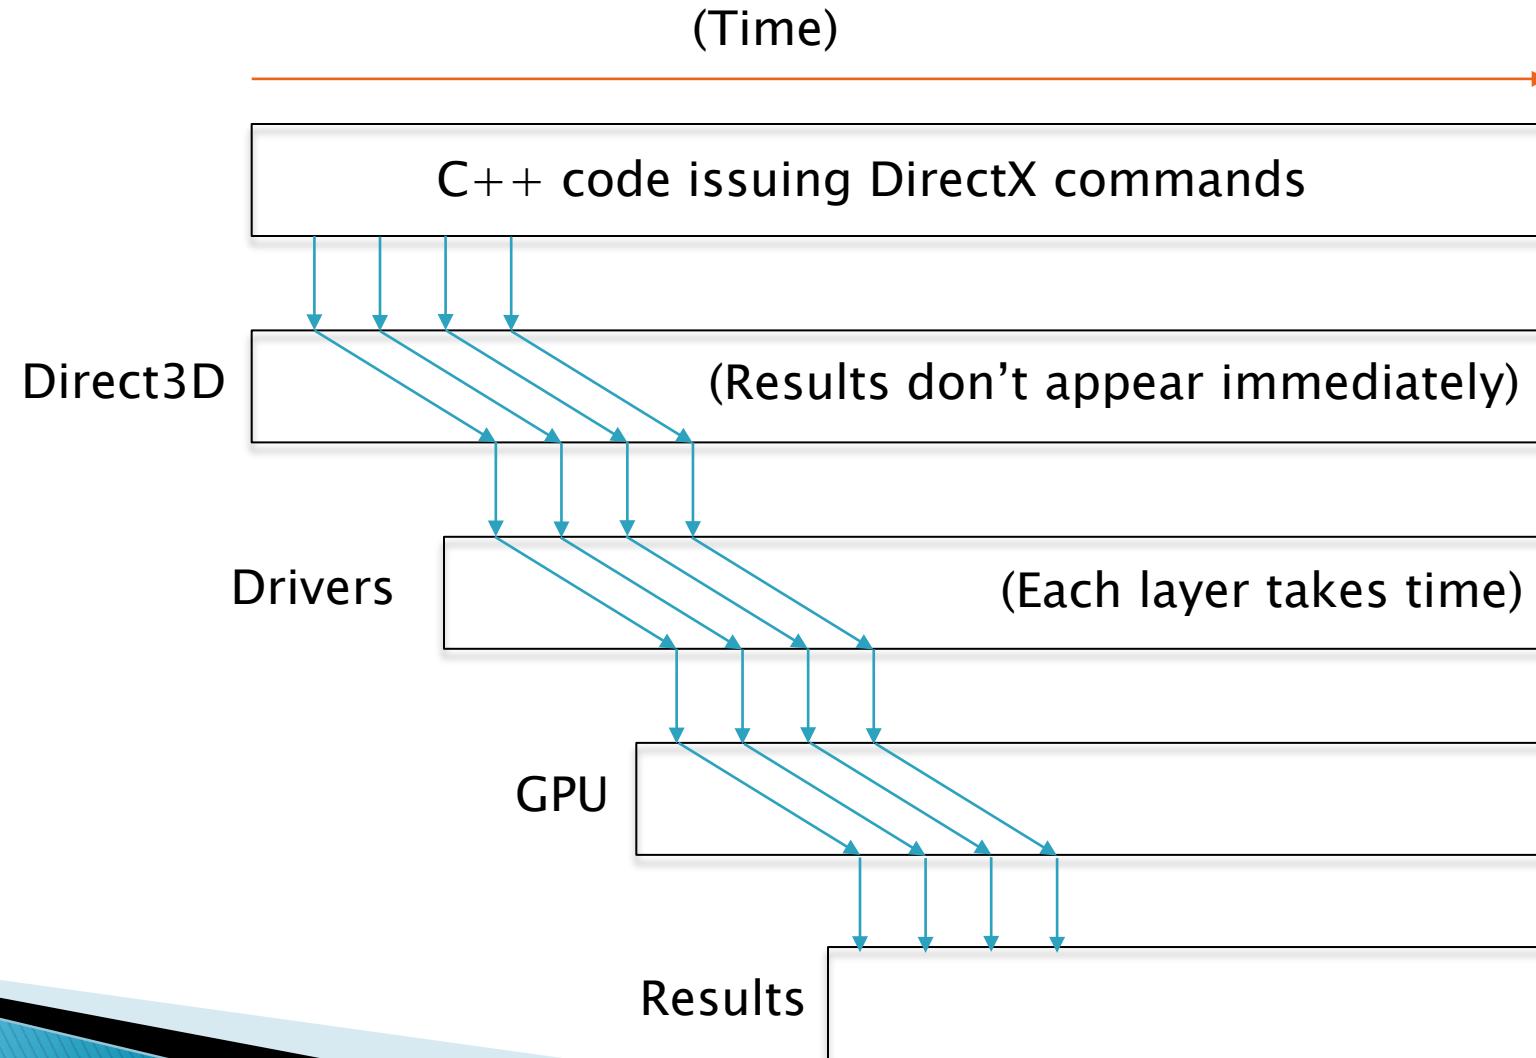
CPU-Side vs. GPU-Side



One-way pipeline

- ▶ Pipeline is optimized in ONE DIRECTION: CPU → GPU
- ▶ After a frame is rendered & displayed, it is **thrown out**
 - Clear the frame
 - Replace with the next one
- ▶ Attempting to get *any* results back causes a **stall**
 - All CPU code must WAIT for GPU to catch up
 - Framerate drops dramatically

One-way pipeline example



CPU readback = Pipeline stall

