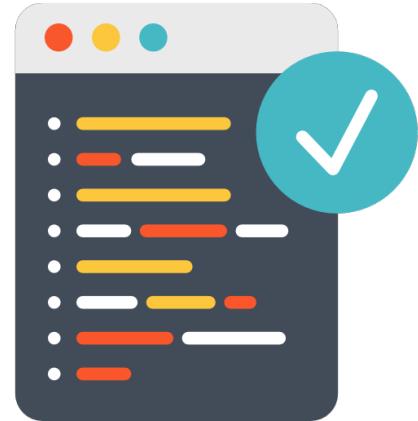


Introduction to Shaders

In Direct3D 11

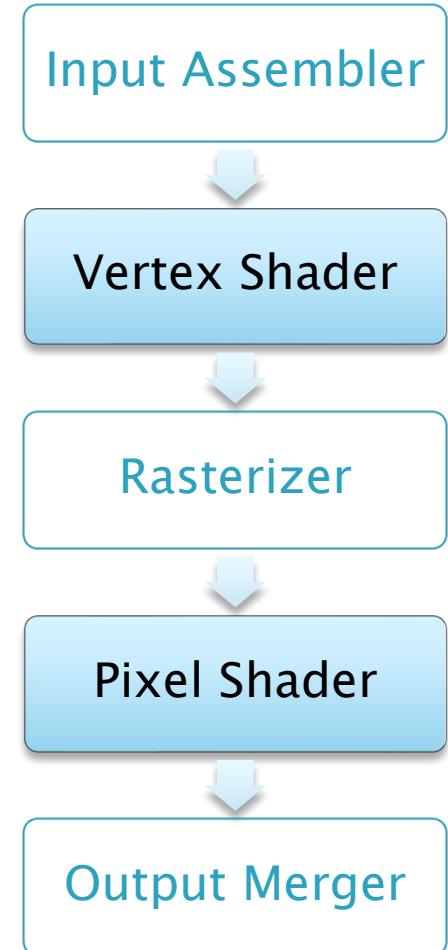
Shaders

- ▶ Small programs executed by the GPU
 - During the rendering pipeline
 - Need to draw something to run a shader!
- ▶ Written in HLSL (for Direct3D)
 - High-Level Shading Language
- ▶ The core of a “programmable pipeline”
 - Required by the pipeline
 - Can’t draw if you don’t have shaders!



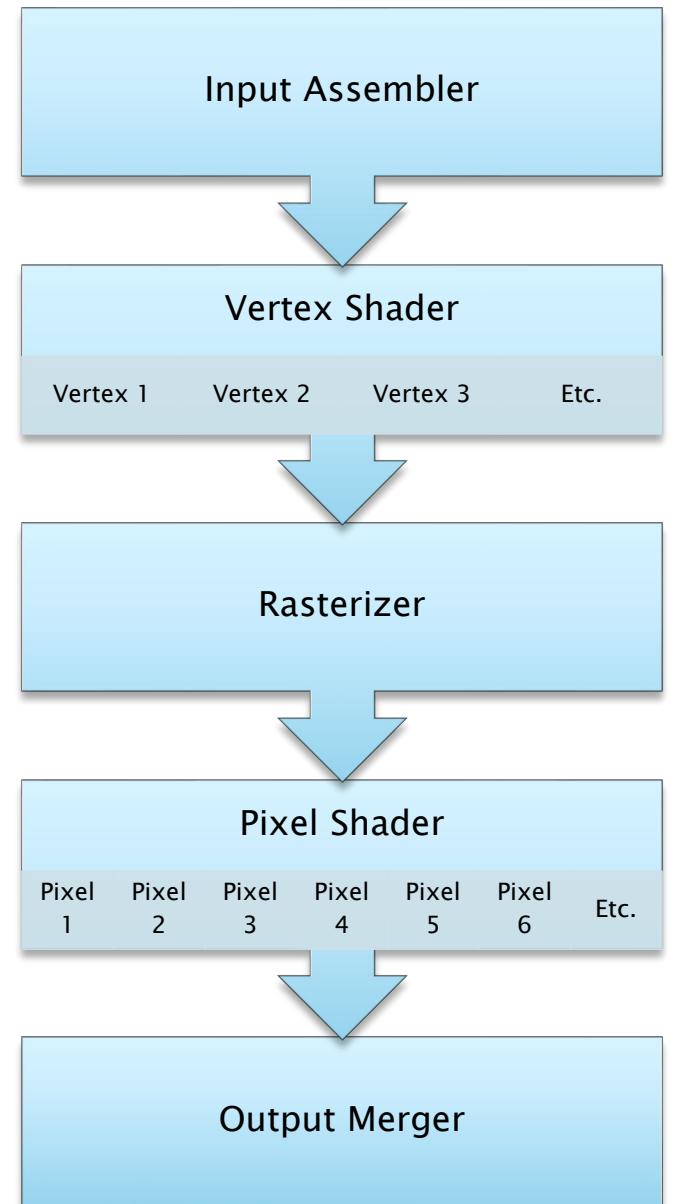
Shaders & the pipeline

- ▶ Two required shader stages during pipeline
 - Vertex Shader
 - Pixel Shader
- ▶ Shaders receive data from previous stages
 - The shader's *input*
- ▶ And must send data to the next stage
 - The shader's *output*



Shaders & parallelism

- ▶ Each shader stage represents multiple *shader instances* running in parallel
 - Copies of the shader program
- ▶ All shader instances run in lock-step
 - All execute the same instruction at once (SIMD!)
 - Each instance is fed different input
 - Allowing instances to produce different output
- ▶ All instances must finish before next stage

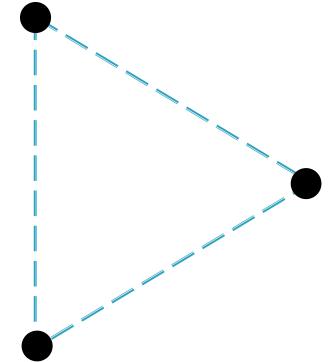


Direct3D shader types

- ▶ **Vertex and pixel shaders required for basic rendering**
- ▶ **Several other varieties**
 - Hull
 - Domain
 - Geometry
 - Compute (not part of the pipeline)
- ▶ **We won't cover these much this semester**
 - Might have a neat demo of compute shaders later on

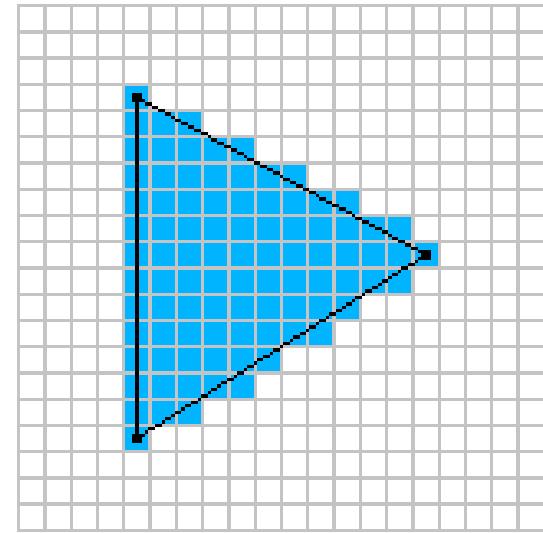
Vertex shader

- ▶ **Main goal:** Transform local vertex position to screen-space position
- ▶ **Inherent Input:** All data related to a single vertex
 - Each instance gets a different vertex from bound vertex buffer
 - From input assembler stage
- ▶ **Minimum output:** 4-component screen-space vector
 - Must be tagged as “SV_POSITION”
 - Can also output any other per-vertex data



Pixel shader

- ▶ **Main goal:** Determine the color of this pixel
- ▶ **Inherent Input:** Interpolated vertex output data
 - Each instance gets slightly different data
 - From rasterizer stage
- ▶ **Minimum output:** 4-component color vector (RGBA)
 - Must be tagged as “SV_TARGET”



Shaders - Details & Syntax



HLSL: High-Level Shading Language

- ▶ HLSL is the shading language for Direct3D
- ▶ C-like language
 - Originally, shaders were written in pure assembly
 - These days: compiled from HLSL to assembly
- ▶ Need to be on the GPU before execution
- ▶ Generally, individual shaders written for specific effects
 - Rather than an *all-in-one mega shader*

“Shader model” = Version

- ▶ The “version” of HLSL you’re using
 - Direct3D 11’s shader model is 5.0
- ▶ Higher shader models generally have:
 - More features
 - Higher instruction limits
- ▶ Can compile shaders to older models
 - For backwards compatibility
 - Purposefully limits language features

Common programming constructs

- ▶ HLSL has:
 - Local variables (defined in a function)
 - Global variables (from constant buffers)
 - Functions
 - Structs
 - Loops
 - If statements
 - Switch statements
- ▶ Note: Branching *may* negatively impact performance!
 - Worst case when different instances take different branches

Common data types, operators & functions

▶ Common data types

- Scalar: int, bool, float
- Vector: float2, float3, float4
- Matrix: float3x3, float4x4 (matrix)

▶ Operators

- Standard mathematical, relational, boolean & bitwise operators

▶ Intrinsic functions

- Basic math and trig: abs(), pow(), min(), max(), sin(), cos(), etc.
- Vector math: reflect(), refract(), dot(), cross()

Vector initialization examples

```
// Basic initialization
float4 color = float4(1.0f, 0.5f, 0.0f, 1.0f);
float2 two    = float2(5.5f, 6.2f);

// Initialize a larger vector with a smaller one
float4 large = float4(two, 1.0f, 1.0f);
float4 again = float4(1.0f, two, 1.0f);
float4 twice = float4(two, two);
```

Vector Components

- ▶ A vector's components may be accessed with:
 - r, g, b, a
 - x, y, z, w
- ▶ Accessing data in colors (or ANY vector):

```
float4 color      = float4(1,0,0,1);
float redAmount   = color.r; // or color.x
float greenAmount = color.y; // or color.g
```

Swizzling

- ▶ Syntactical shortcut for rearranging vector components

```
float4 color      = float4(1.0f, 0.5f, 0.1f, 0.0f);
float4 foSwizzle = color.xxyx; // 1, 1, 0.5, 1
float4 worksFine = color.agbb; // 0, 0.5, 0.1, 0.1
float4 error      = color.xbya; // Can't mix notations
```

Swizzling - Automatic repetition

- ▶ If fewer components than necessary specified, last is repeated
- ▶ When a float4 is expected:
 - .xy really means .xyyy
 - .z really means .zzzz
 - .xxx really means .xxxx
- ▶ Code examples:
 - `float4 color = float4(1, 0.5f, 0.25f, 0);`
 - `float4 allZ = color.z; // 0.25, 0.25, 0.25, 0.25`
 - `float4 someY = color.xy; // 1, 0.5, 0.5, 0.5`

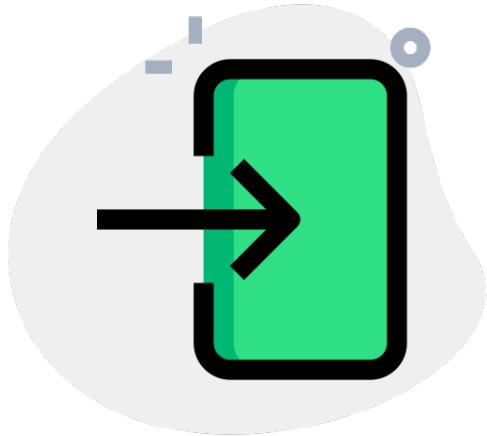
Functions

- ▶ C-like functions
 - Require return types (or void)
 - Can have parameters
 - Support *out* parameters
- ▶ Can write any number of functions
- ▶ Always inline'd by compiler
 - No function call overhead!
 - And no recursion

Shader entry point

- ▶ **main()** function
 - Where a shader's execution begins
- ▶ Must have a return type
 - Represents data to send to next stage of pipeline
 - Each type of shader has unique requirements for its return value(s)
- ▶ Can (and will) have parameters
 - Represents data from previous pipeline stage

Parameters to main()



- ▶ Represent input from previous pipeline stages
 - All shaders will have one or more params
- ▶ Option 1: Comma-separated variables

```
// Pixel shader with 2 parameters
float4 main(float2 uv : TEXCOORD, float4 c : COLOR) : SV_TARGET {...}
```

- ▶ Option 2: Custom struct with multiple members

```
// Pixel shader with 1 struct parameter
float4 main(VS_OUTPUT psInput) : SV_TARGET {...}
```

Return value from main()

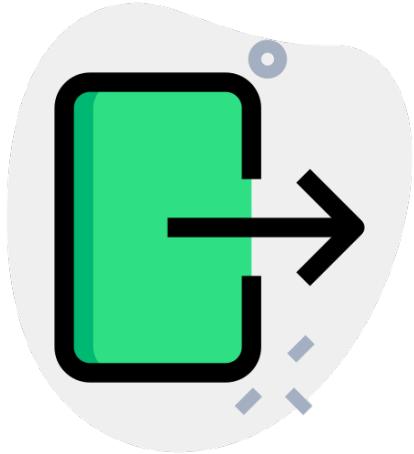
- ▶ Shaders require a return type for output
 - Sends data to next pipeline stage

- ▶ Option 1: A single value

```
// Pixel shader returning one variable worth of data
float4 main(VS_OUTPUT psInput) : SV_TARGET {...}
```

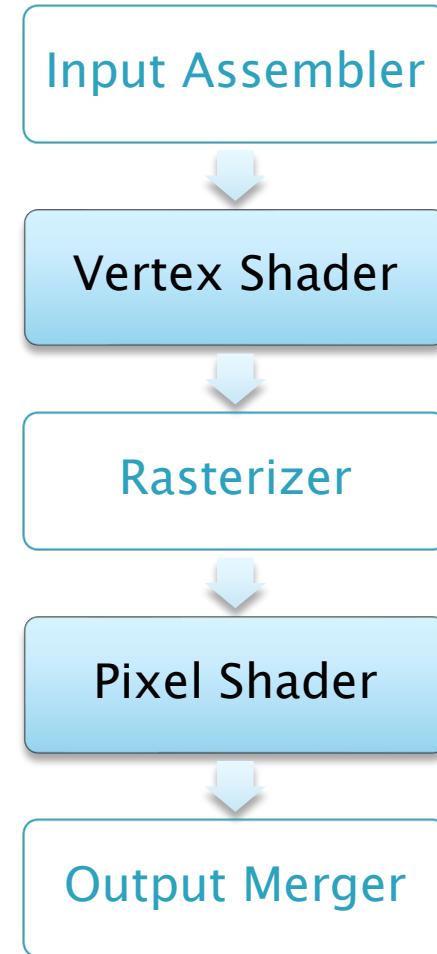
- ▶ Option 2: Custom struct with multiple members

```
// Pixel shader returning a whole struct
PS_OUTPUT main(VS_OUTPUT psInput) {...}
```



Basic shader input & output

- ▶ Vertex shader
 - **Parameters:** Individual per-vertex data
 - **Must return:** Transformed per-vertex data
- ▶ Pixel shader
 - **Parameters:** Per-pixel interpolated VS output
 - **Must return:** Color (a single float4)
- ▶ VS output must match PS input
 - Can define common structs in external files
 - And use #include as necessary



Semantics

- ▶ Tags for shader input & output
 - Additional meta-data for the pipeline
 - Informs Direct3D how to process this data
 - Required on parameters & return types

```
struct VertexShaderInput
{
    float3 localPosition : POSITION;
};
```

Semantic



Required semantics

- ▶ Some semantics are **expected and required**
 - Begin with SV_ (“System Value”)
- ▶ Vertex shader
 - Output *must* include a float4 tagged as SV_POSITION
 - Rasterizer expects this data
- ▶ Pixel shader
 - Output *must* include a float4 tagged as SV_TARGET
 - Output Merger expects this data



Semantic names are flexible



- ▶ `float3 color : HORSE;`
 - Technically works
 - But not great – lacks actual meaning/intent
- ▶ Common conventions
 - POSITION, NORMAL, TANGENT, TEXCOORD, COLOR, etc.
 - These had meaning in Direct3D 9.0
- ▶ When in doubt: use a string that shows intent
 - CAMERAPOSITION is better than LOLWTFPLEASECOMPILE

Shader include files

- ▶ Can create files to be #include'd into shaders for organization
- ▶ Common extension is .hlsl
 - “HLSL Include”
 - But any extension works – it's all just text
- ▶ Must use C-style include guards (no “#pragma once” here)

```
#ifndef UNIQUE_NAME_FOR_THIS_FILE
#define UNIQUE_NAME_FOR_THIS_FILE
<all code goes here>
#endif
```

Global shader data from C++

- ▶ Data that is shared (constant) among all shader instances
- ▶ Constant buffers!
 - Global shader variables
 - Good old cbuffer definitions
- ▶ No semantics on cbuffer variables
 - This is data straight from GPU memory
 - Not “flowing through” pipeline stages

Resources & registers

- ▶ **Registers:** Predefined “slots” for resources
- ▶ Common resource types:
 - Samplers (s registers)
 - Textures (t registers)
 - Constant buffers (b registers)

```
cbuffer externalData : register( b0 )  
{  
    matrix worldMatrix;  
};
```

Available shader data

- ▶ Data for a shader comes from one of three places
 - 1. Input from previous pipeline stage
 - 2. Constant buffers
 - 3. Texture resources
- ▶ In all cases, the data must be on the GPU before calling Draw()
- ▶ It's our job to put it there

Shader Possibilities



What *can* we do with shaders?

- ▶ Technically...anything we want!
 - It's all numbers
 - It's all math
- ▶ Perform lighting equations to mimic photo-realism
- ▶ Procedurally generate abstract art
- ▶ Create music visualizations
- ▶ <https://www.shadertoy.com/> – WebGL shader playground!

What *will* we do with shaders?

- ▶ A game engine usually isn't generating abstract art in shaders
- ▶ Our goal:
 - Mimic surface materials
 - Under various lighting conditions
 - To create photo-realistic 3D worlds
- ▶ Plus a few other interesting effects

How many shaders do we need?

- ▶ Different effects generally have unique shaders
 - Static meshes
 - Animated meshes
 - Photo-realistic materials
 - Stylized materials (e.g. toon shading)
 - The skybox
 - Particle systems
 - Post processing passes (glow, depth of field, etc.)
 - User interface
- ▶ Each of these would use different shaders

Managing Shaders

Managing shaders

- ▶ Each shader gets its own .hlsl file
 - Visual Studio will compile your .hlsl files
 - And do syntax highlighting
 - And give you errors & line numbers if necessary
- ▶ While your application is running, you should:
 - Load pre-compiled shaders at start-up
 - Set (“activate”) individual shaders when necessary
 - Combine them in whatever ways make sense

Combining Shaders

- ▶ Shaders must be compatible to work together
- ▶ You can't just throw any two (or more) shaders together
- ▶ At a minimum, the following must match:
 - Output of your vertex shader
 - Input of your pixel shader

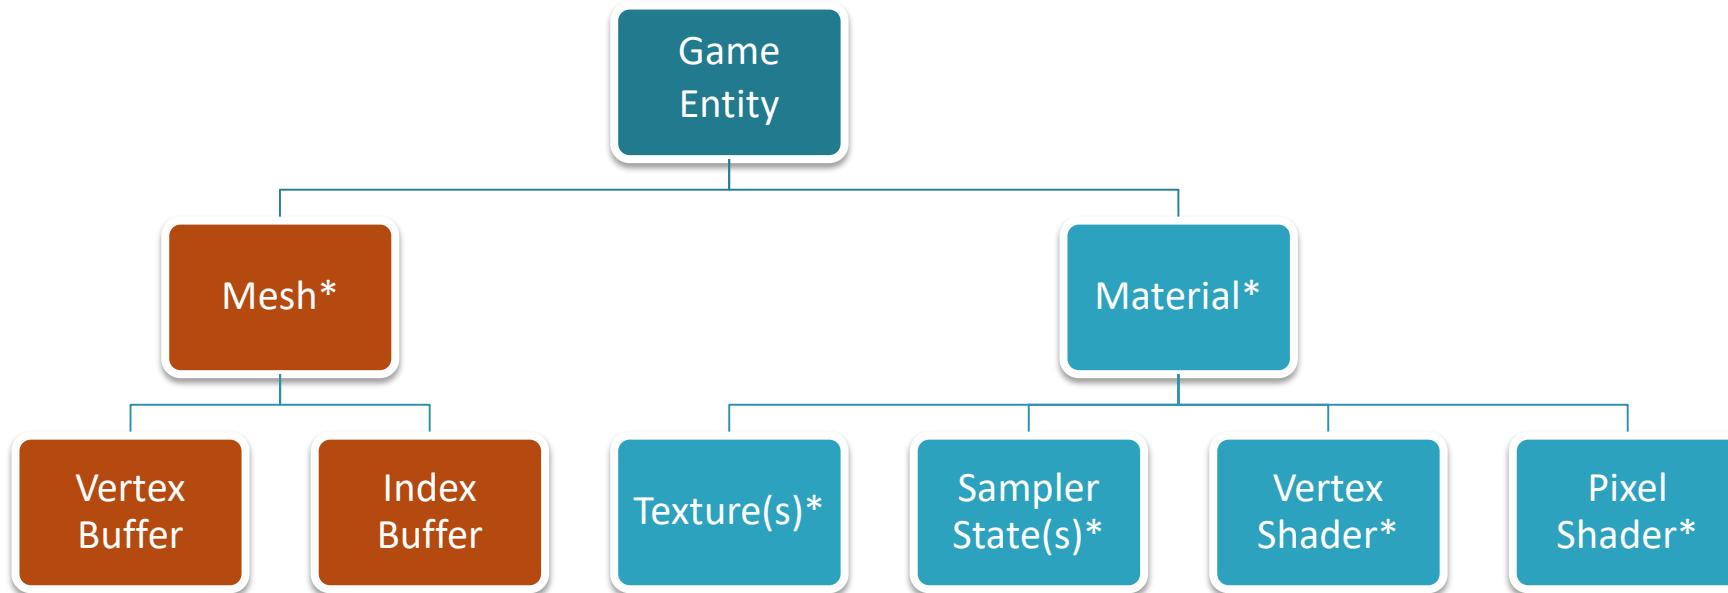
Drawing with Shaders

- ▶ When drawing, you need:
 - A vertex shader
 - A pixel shader
- ▶ You might also want:
 - One or more *textures*
 - One or more *samplers* (options for using textures)
 - External data (numbers) – colors, lights, time, etc.

Material System

- ▶ What's a material?
 - Combination of shaders, textures and other data
 - Describes the look of a mesh
 - Think of materials in Unity as an example
- ▶ Material definition should include shaders
- ▶ When “using” a material
 - Activate material’s shaders
 - Draw one or more meshes

Basic game entity hierarchy



Example Shaders

From the Starter Code

Example shaders

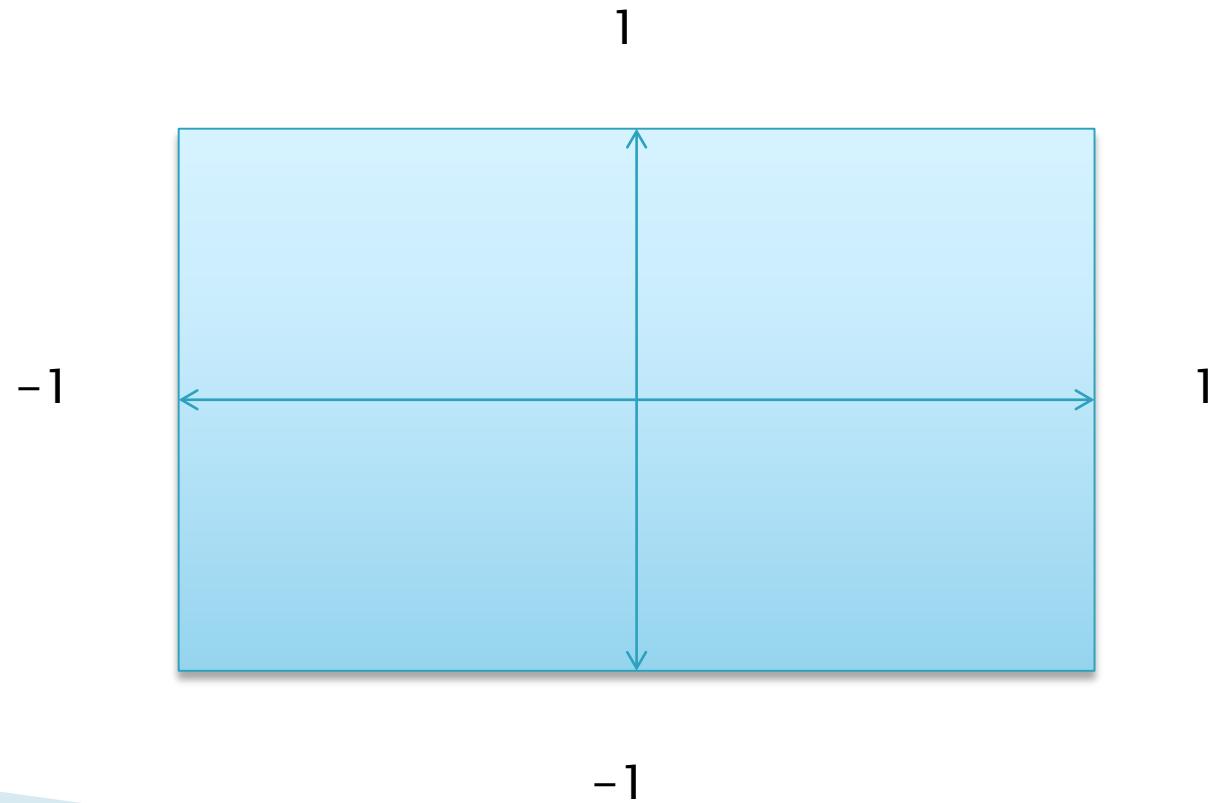
- ▶ We'll look at the starter project shader code
 - VertexShader.hlsl – Vertex shader code file
 - PixelShader.hlsl – Pixel shader code file
- ▶ Each is compiled by Visual Studio
 - At build time
 - Outputs .cso files

Vertex shader

- ▶ Performs per-vertex processing
- ▶ At a minimum:
 - Output vertex position in homogenous clip space
- ▶ Homogenous clip space
 - Normalized device coordinates
 - Resolution-agnostic coordinate system
 - –1 to 1 on each axis
- ▶ Can also output other data

Normalized device coordinates

- Vertex shader output screen coordinates:



VS: Constant buffer

- Constant buffer (cbuffer) holds external data
 - “Per model” data can be set from CPU-side

```
// Constant Buffer
cbuffer externalData : register( b0 )
{
    matrix world;
    matrix view;
    matrix projection;
};
```

Binds this variable to a specific constant buffer index

- **b** means Constant Buffer
- **0** is the index

- Built-in data type
- Same as float4x4

VS: Input struct

```
// Represents a single vertex worth of data
// - Must match vertex layout in your buffers
struct VertexShaderInput
{
    float3 localPosition : POSITION;
    float4 color       : COLOR;
};
```

- Built-in data types
- Multi-component vectors

- These are called “semantics”
 - Define the usage of the variables
 - Some names have meaning
 - Others can be custom
- Can also have an index – Examples:
 - COLOR0, TEXCOORD1, etc.
- HLSL handles some types differently

VS: Output struct

- ▶ “VertexToPixel” – my own naming convention
 - Used as the VS output & PS input struct

```
// Represents data we send down the pipeline
// - Should match pixel shader's input
struct VertexToPixel
{
    float4 screenPosition    : SV_POSITION; // SV = System Value (required)
    float4 color              : COLOR;
};
```

VS: Entry point

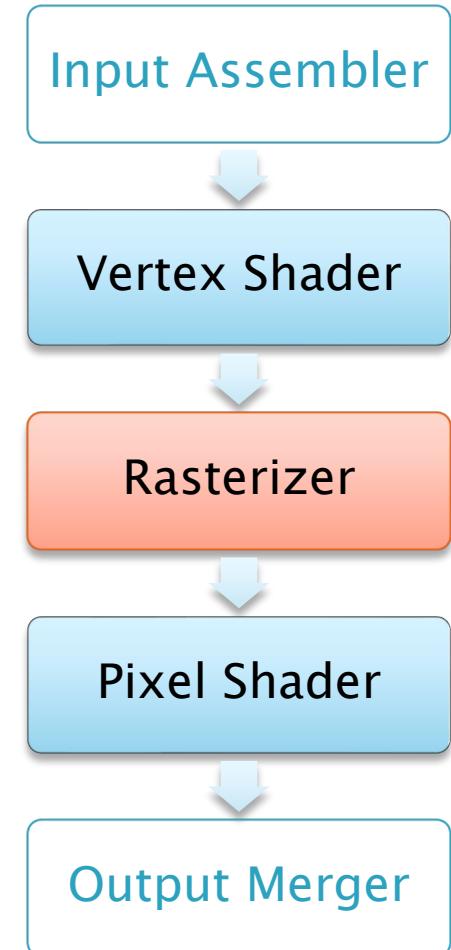
```
VertexToPixel main( VertexShaderInput input )
{
    // Create output variable
    VertexToPixel output;

    // Pass data through (in correct format)
    output.screenPosition = float4(input.localPosition, 1.0f);
    output.color = input.color;

    // Return entire struct (pass to rasterizer)
    return output;
}
```

Between VS and PS

- ▶ Rasterization occurs
- ▶ Rasterizer turns triangles into pixels
 - After the vertex shader is run on all vertices
- ▶ Rasterizer interpolates all VS output
 - Each pixel gets an altered version of VS output
 - Sometimes called “per-pixel” data



Pixel Shader

- ▶ Performs per-pixel processing
- ▶ Input struct must match VS output struct
- ▶ At a minimum, must output a color

Pixel Shader - Input Struct

- ▶ Could #include external file defining this struct

```
// Defines the input to this pixel shader
// - Represents data we expect to receive
//   from earlier pipeline stages
struct VertexToPixel
{
    float4 screenPosition      : SV_POSITION;
    float4 color              : COLOR;
};
```

PS – Actual Shader “Program”

This pixel shader returns
a single color as a
4-component vector

This pixel shader's output data has
a semantic: SV_TARGET

Results go to current render target

```
// Entry point for this pixel shader
float4 main(VertexToPixel input) : SV_TARGET
{
    return input.color;
}
```