

Lighting & Shading



What do we mean by
“lighting” and “shading”?



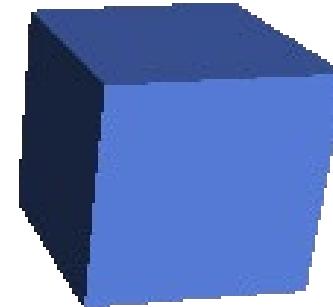
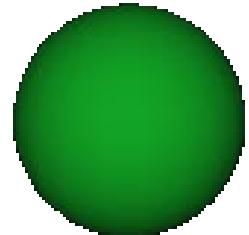
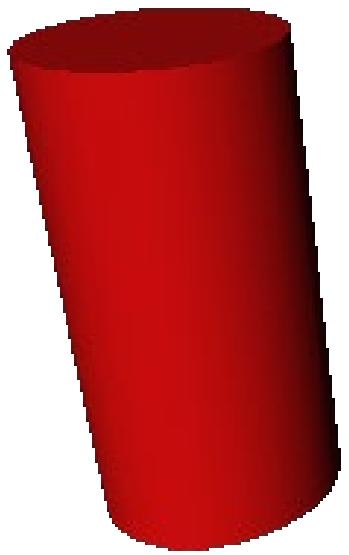
Lighting vs. Shading

- ▶ **Lighting:** Approximating how light interacts with a surface
 - Mathematical equations
 - Based on physics
- ▶ **Shading:** Applying lighting results to the pixels of a mesh
 - Running the equations
 - Combining with surface materials
 - Getting the results on the screen
- ▶ Terms sometimes used interchangeably

What we have so far

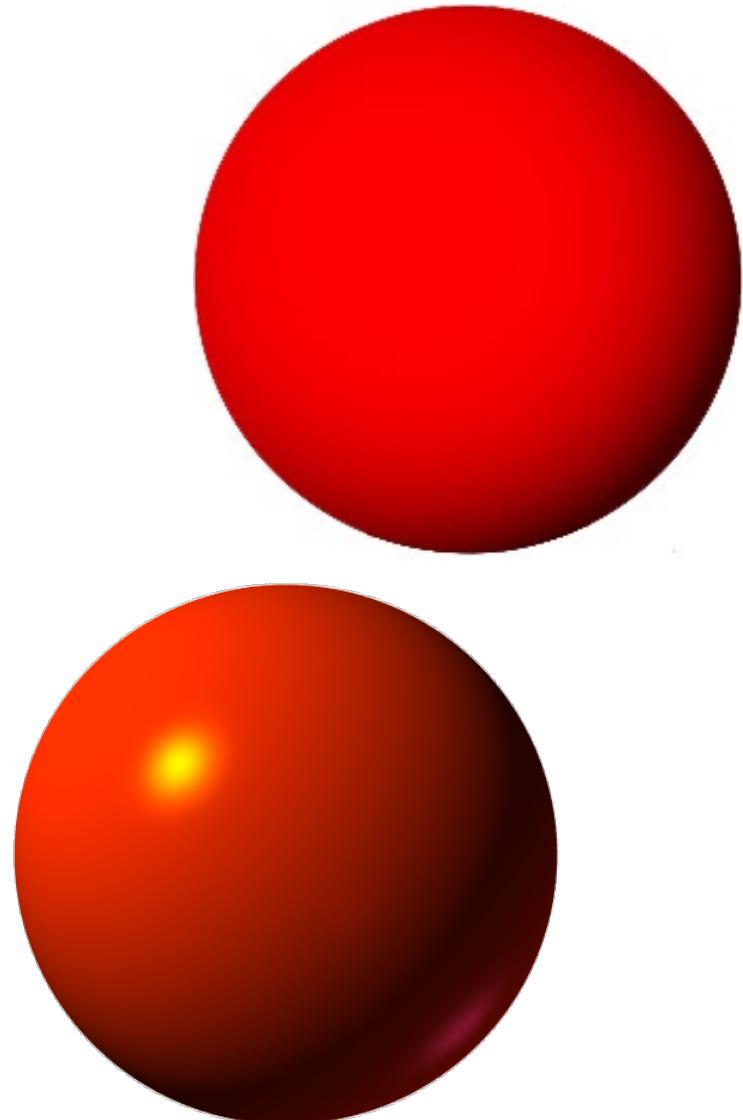


What we want (to start...)



Lighting

- ▶ We aren't *simulating* actual light
- ▶ We're *approximating* it
- ▶ Looking for plausible results
- ▶ Which are quick to calculate
- ▶ Need “simple” math functions to describe light/surface interactions



What is really happening with light?

- ▶ Actually lighting is very complex
- ▶ Incredible amount bouncing around
- ▶ Reflections everywhere
- ▶ Too much to calculate in real time!
- ▶ We focus on major components



Colors in a computer

- ▶ Colors stored as three numeric components (channels)
 - Red, Green & Blue
 - Each is 0–255 (or 0.0f – 1.0f in a shader)
 - ~16.7 million combinations

Red (1.0, 0.0, 0.0)

Blue (0.0, 0.0, 1.0)

Purple (1.0, 0.0, 1.0)

White (1.0, 1.0, 1.0)

Black (0.0, 0.0, 0.0)

Grey (0.5, 0.5, 0.5)

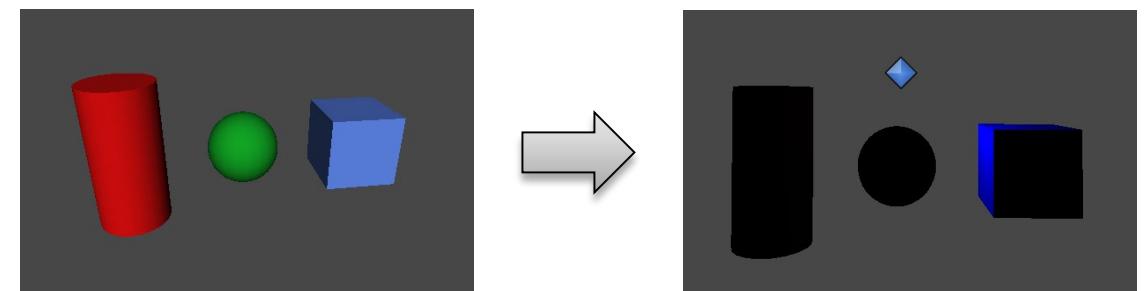
Combining colors

▶ Adding colors

- Often used when applying multiple lights
- More light = brighter!
- Perceptual color may change (blue + red = purple)

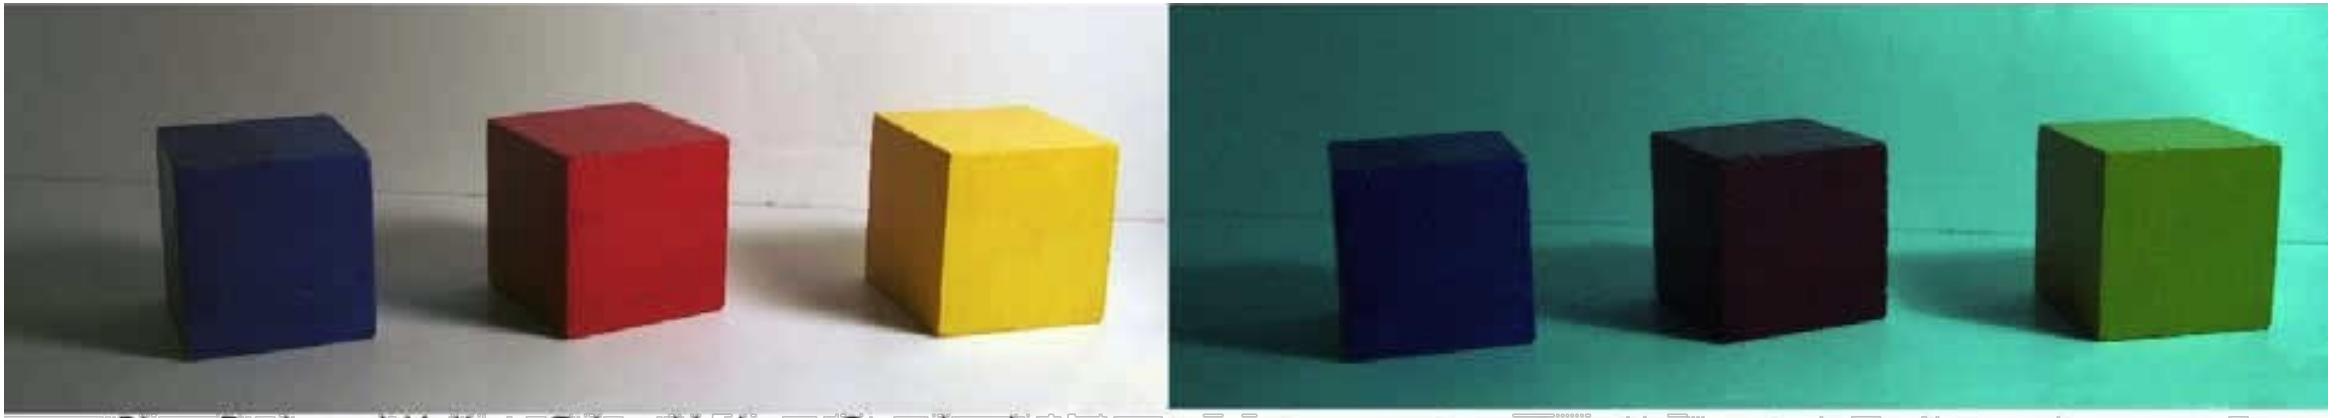
▶ Multiplying colors

- Often used when tinting colors
- Light striking a surface is a tint
- Blue * red = black



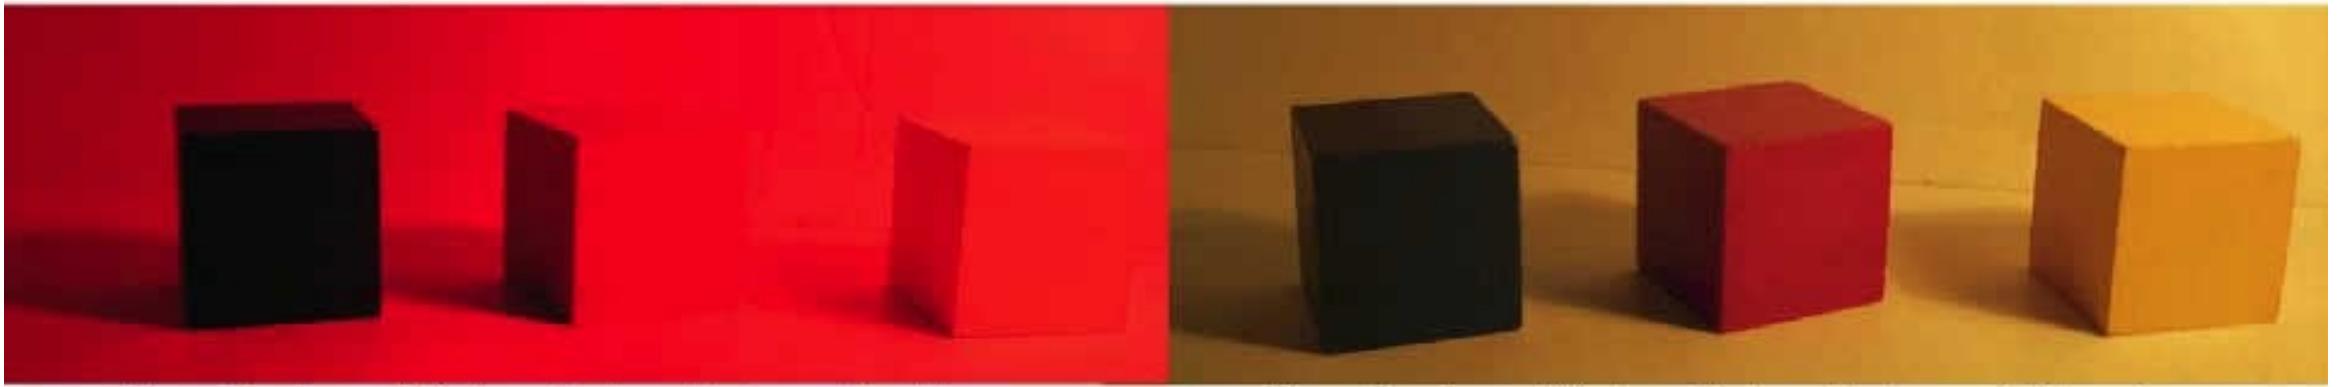
▶ Color * Scalar → Brighten or darken

Real example of lights & surfaces



Blue, Red, and Yellow Cubes Under a Standard Desk Lamp Light Source

Blue, Red, and Yellow Cubes Under a Blue Lamp



Blue, Red, and Yellow Cubes Under a Red Lamp

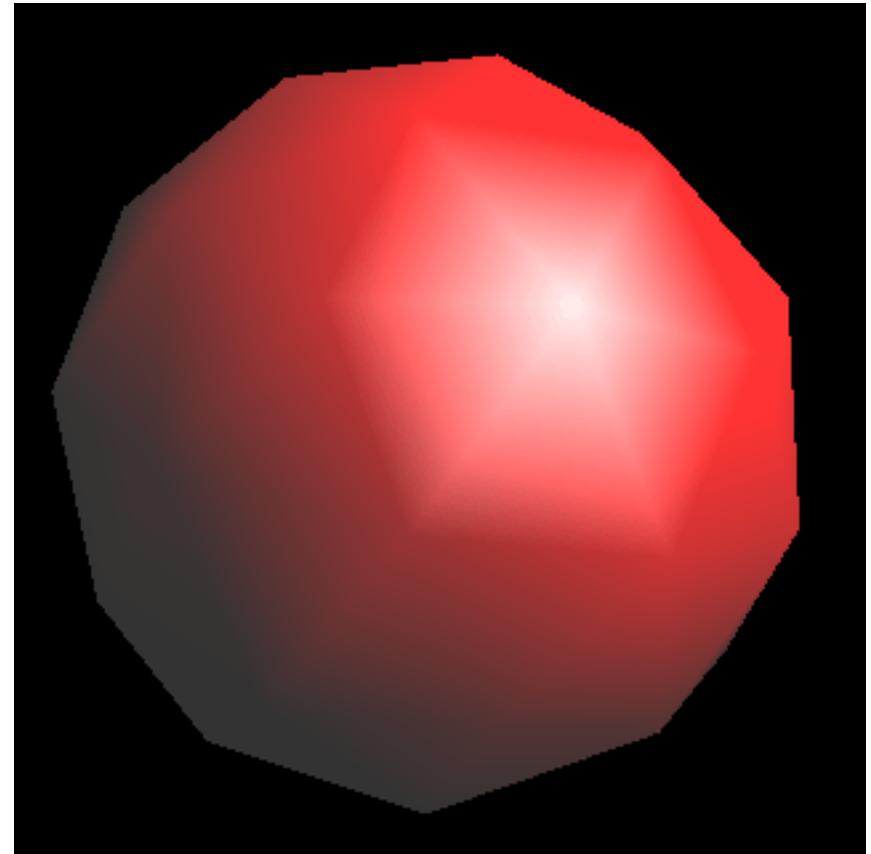
Blue, Red, and Yellow Cubes Under a Yellow Lamp

Quick aside – Where do we put lighting code?

- ▶ It happens in a shader – but which one?
- ▶ Could technically do it in either
 - Lighting in vertex shader = “per-vertex lighting”
 - Lighting in pixel shader = “per-pixel lighting”
- ▶ Which is best?

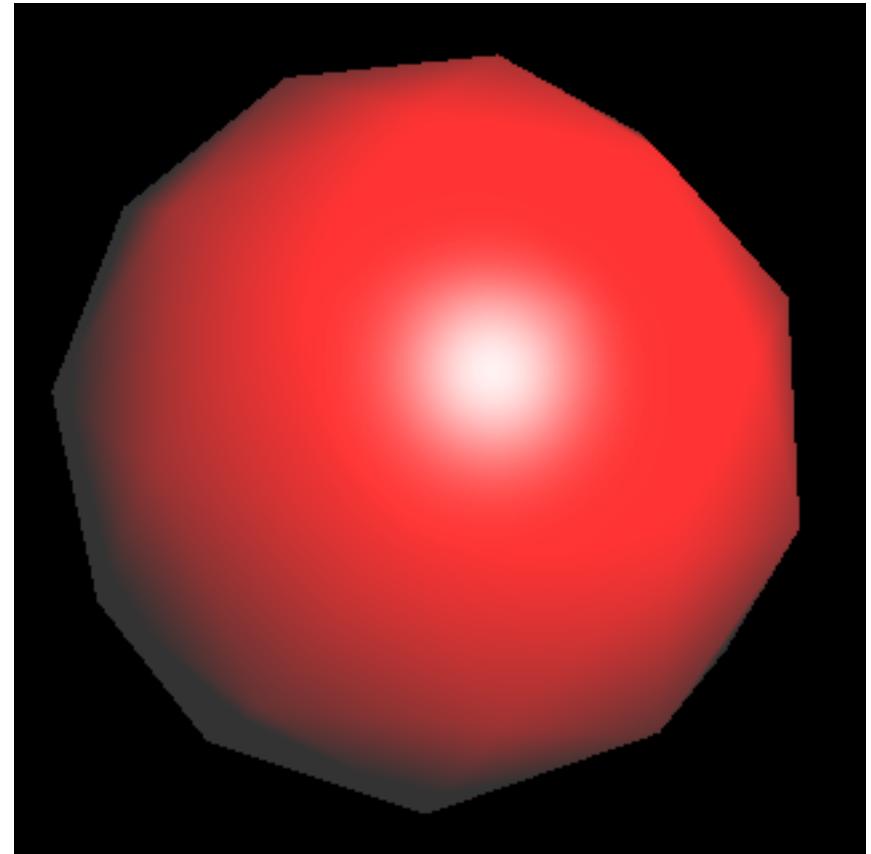
Per-vertex lighting

- ▶ All lighting done in vertex shader
 - Results interpolated by rasterizer
 - Known as Gouraud shading
- ▶ Results in very obvious geometry
- ▶ Faster, but *ugly*
 - Fewer vertices than pixels
 - Fewer overall calculations



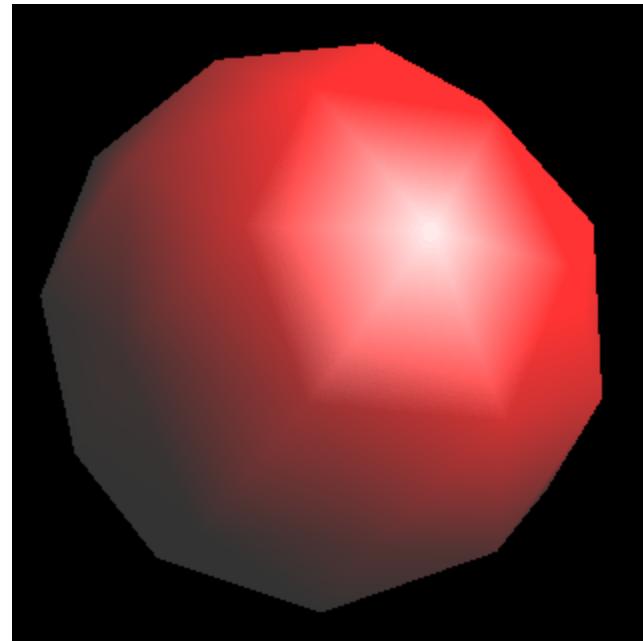
Per-pixel lighting

- ▶ Lighting done at the *pixel level*
 - Pass normal to pixel shader
 - Rasterizer interpolates normal
- ▶ Much more accurate!
 - Changes non-linearly across each triangle
 - Small differences noticeable
- ▶ The only way to fly!

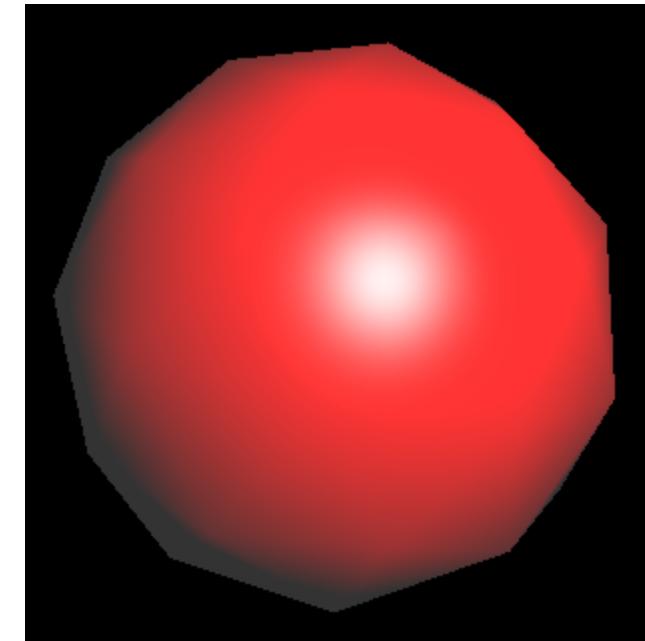


Comparison

Per-Vertex Lighting



Per-Pixel Lighting

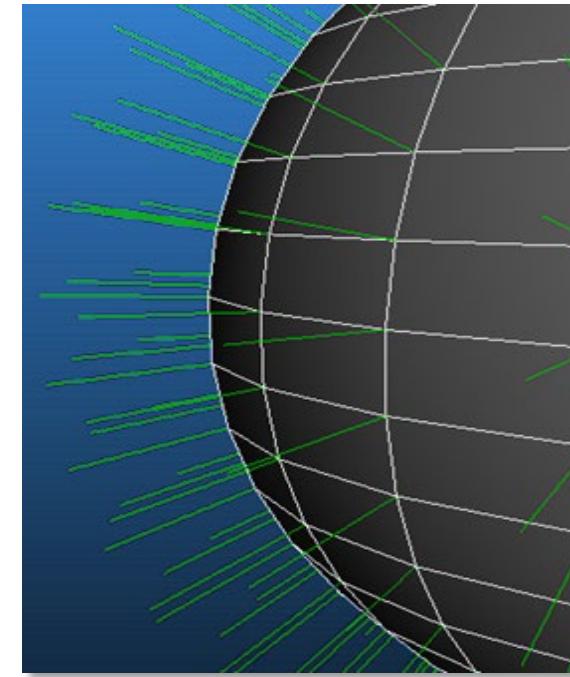


Shading Requirements



Requirement 1: *Surface information*

- ▶ **Normals**
 - Direction surface faces
 - Generally “outward”
 - Defined at each vertex
- ▶ **Shininess**
 - Material property
 - Shiny $\leftarrow \rightarrow$ Matte
- ▶ **Color**
 - Solid RGB color or texture
 - Tinted by light

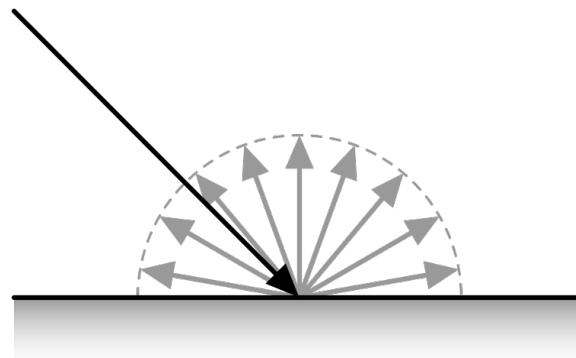


Requirement 2: *Light source*

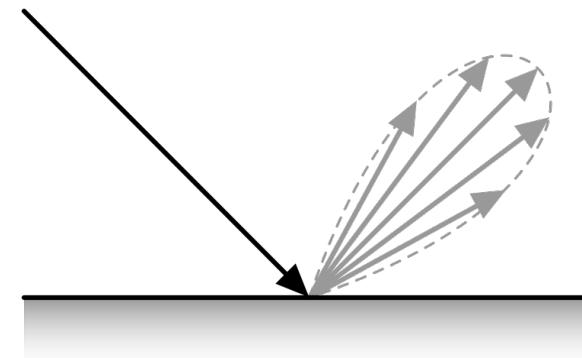
- ▶ **Position** of the light source
 - If it has one
- ▶ **Direction** light is traveling
 - From light source to surface
- ▶ **Color** – single RGB color
- ▶ **Intensity** – a simple multiplier

Requirement 3: *Lighting equations*

- ▶ Bidirectional Reflectance Distribution Functions (BRDFs)
- ▶ Functions that describe how light is reflected



Diffusion
(Lambert)



Glossy Reflection
(Specular / Phong)

1: Surface Information

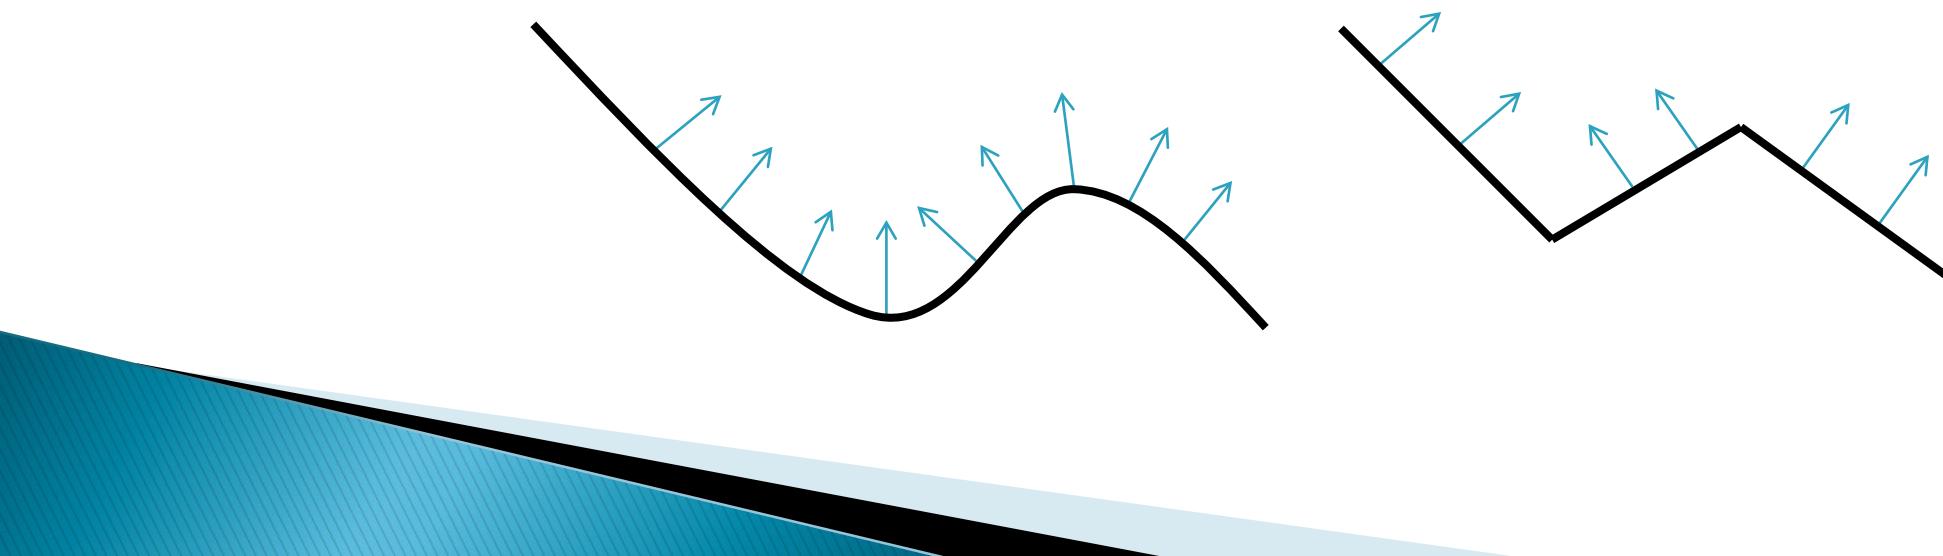


Surface information

- ▶ Already have a “surface” – a mesh
- ▶ Need the following:
 - Normals – Defined in the 3D model file
 - Color – Material property (pick a color or use a texture)
 - Shininess – Material property (pick an amount)
- ▶ Let’s take a closer look at normals

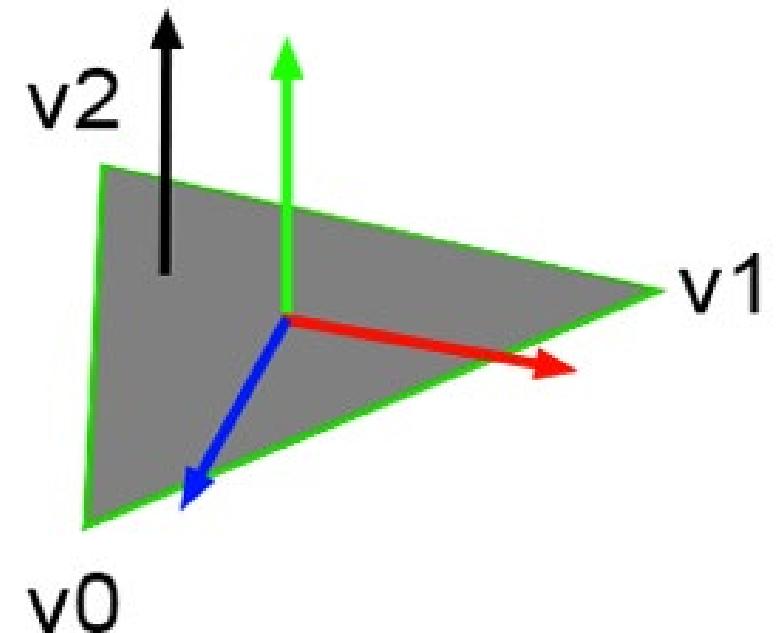
Normals

- ▶ A vector that is perpendicular to the tangent plane to a surface at a particular point
- ▶ Which way the surface “faces” at that point



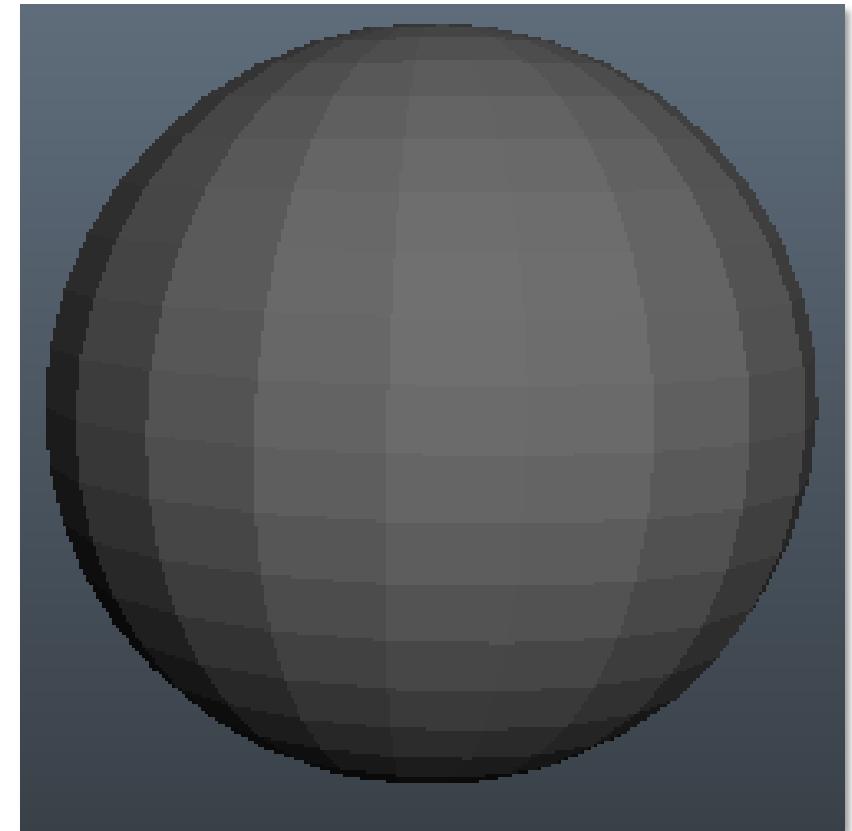
Calculating normal of a triangle

- ▶ Cross product
 - Cross two edge of the triangle
 - $(v_0, v_1) \times (v_0, v_2)$
- ▶ This gives a uniform normal for the entire face
- ▶ Is that what we want?

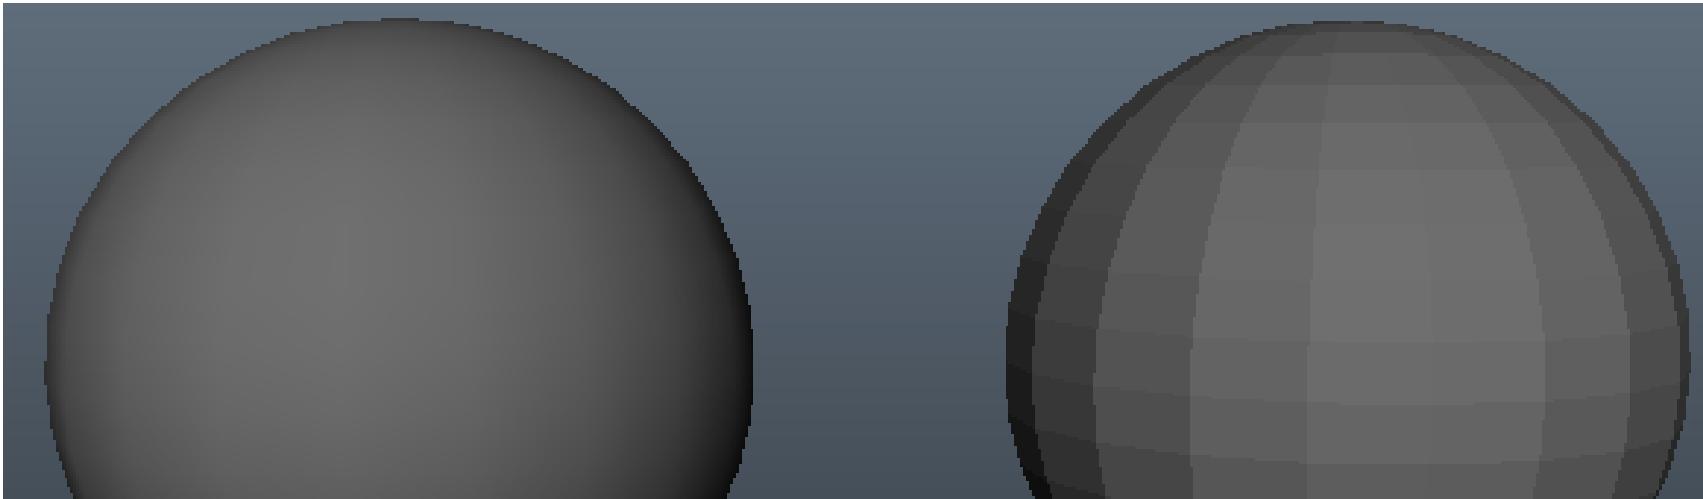


Uniform face normals

- ▶ When each pixel of a triangle has the same normal
- ▶ “Hard shading”
- ▶ Results in flat-looking triangles
- ▶ And hard edges between neighboring triangles



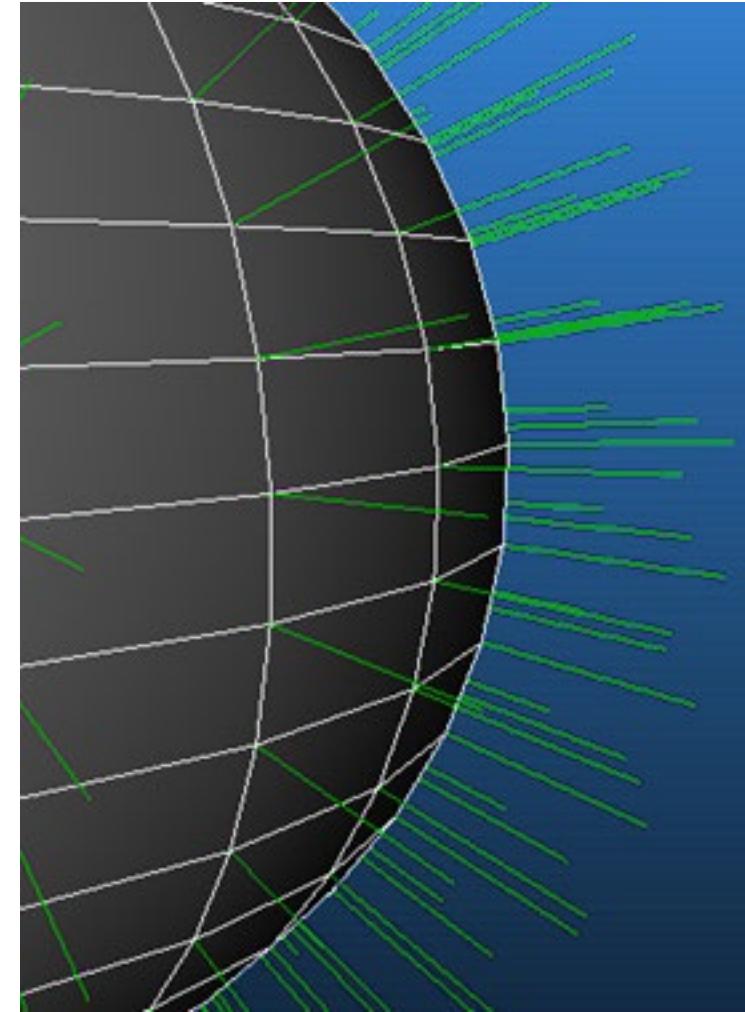
Smooth vs. Hard Shading



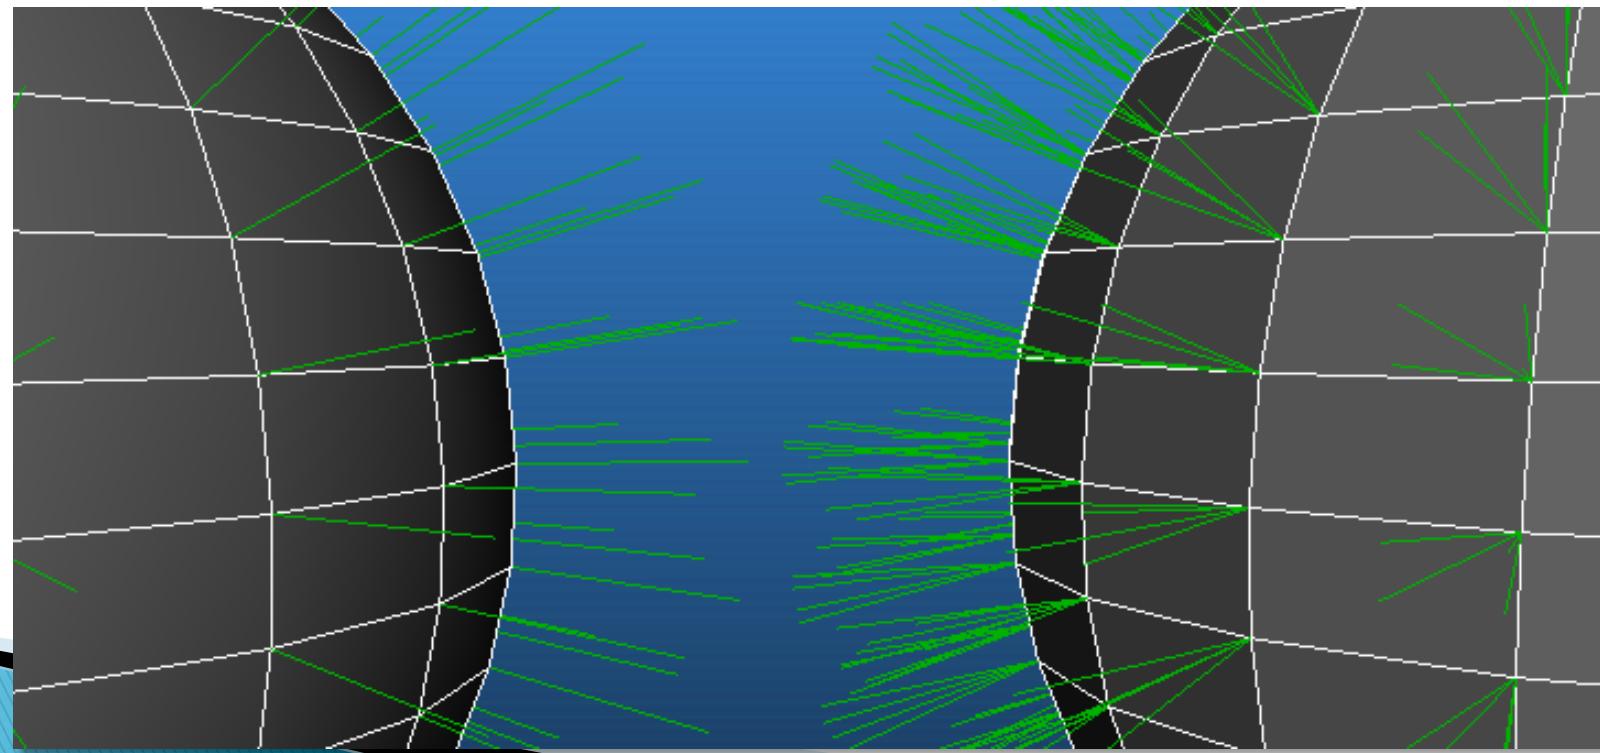
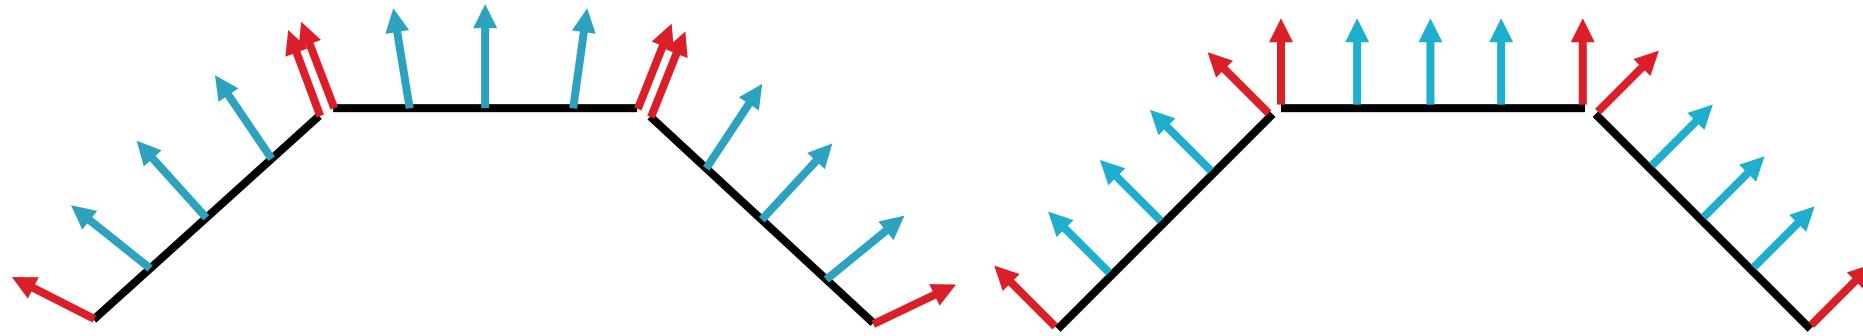
- ▶ Normal is *different* at each pixel of a face
- ▶ Each vertex of a triangle has a *different* normal
- ▶ Normal is the *same* at each pixel of a face
- ▶ Each vertex of a triangle has the *same* normal

Normals

- ▶ Stored at each vertex
- ▶ Define which way the surface “faces”
- ▶ Normal of each *pixel* is interpolated from normals at *vertices*



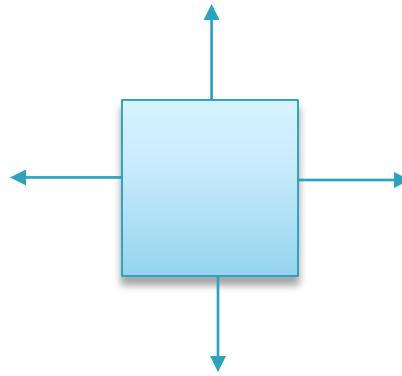
Smooth vs. Hard Shading



Normals in shaders

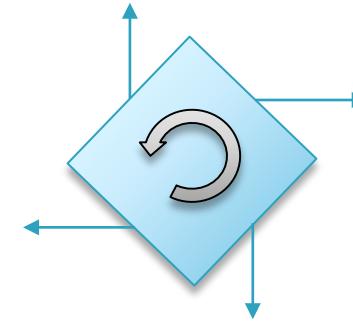
- ▶ VS must send normals to pixel shader
 - Update VertexToPixel
 - `float3 normal : NORMAL;`
- ▶ Just pass normals through? Fine if there's *no rotation*
- ▶ If an object rotates, its normals must be rotated, too!

Rotations & normals

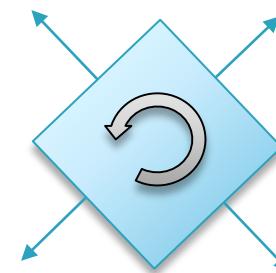


Object with normals
(No rotation)

Only rotating vertices
Normals are incorrect



Rotating both
Normals are correct



Normals in world space

- ▶ Multiply the input normal by world matrix
 - Works fine if there's rotation
 - Works fine if there's *uniform* scaling
- ▶ Problematic if there's translation!
- ▶ Shouldn't translate normal!

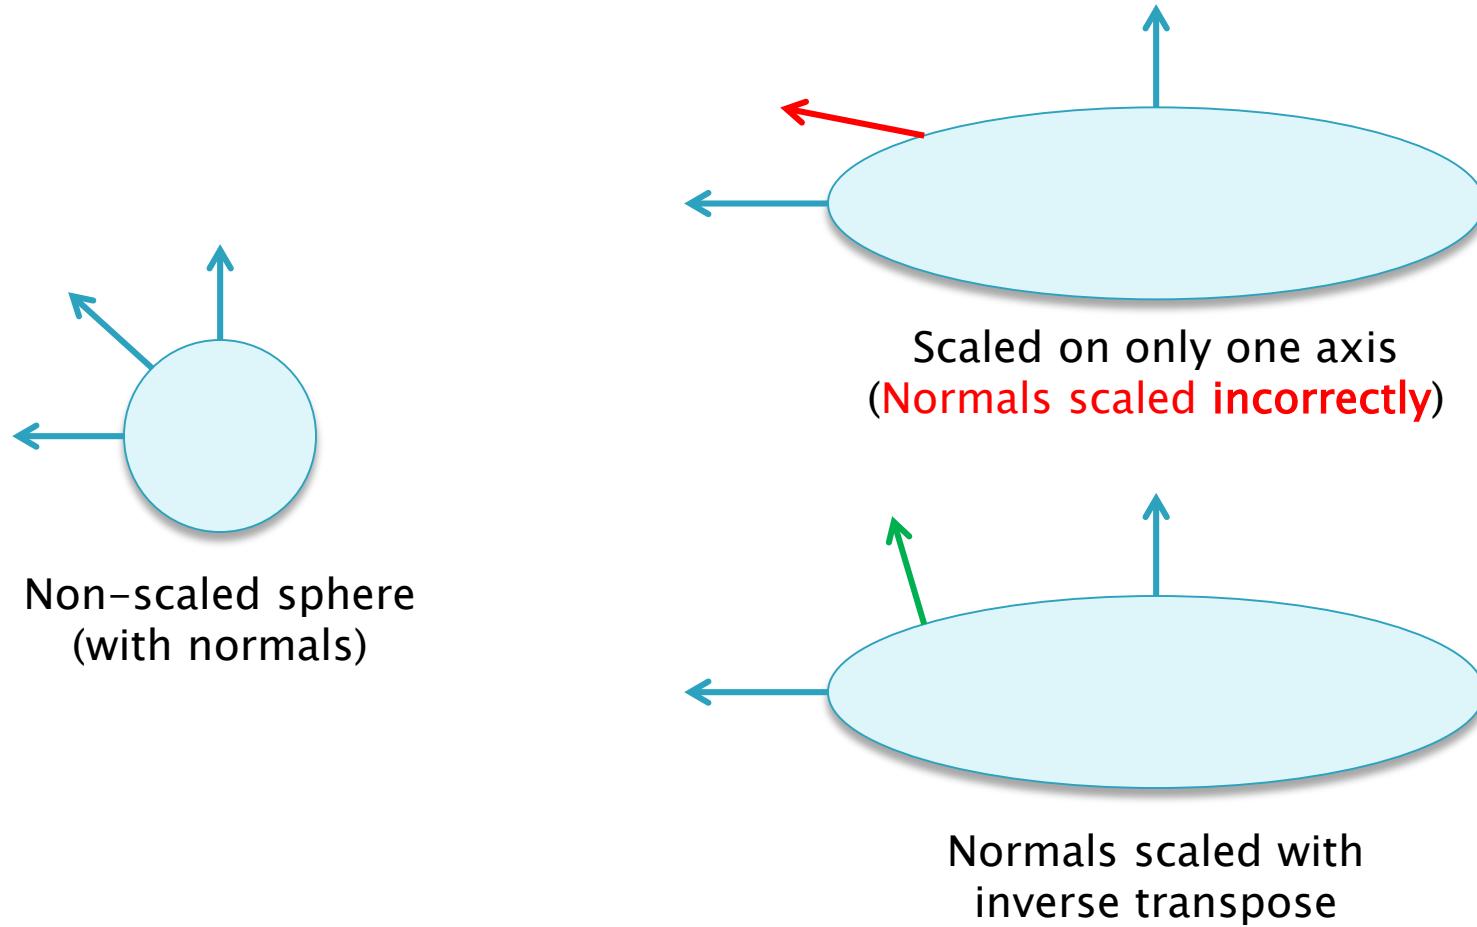
Ignoring translation

- ▶ We just want scaling and rotation
- ▶ Cast the world matrix to a float3x3
 - Keeps top-left 3x3 portion
 - That's the part with scaling and rotation
- ▶ Normal is 3-component anyway
 - So this makes the math easy

Normals & non-uniform scale

- ▶ Problem when multiplying:
 - Normal
 - World matrix with a *non-uniform scale*
- ▶ Non-Uniform Scale
 - Scale with a different value on at least one axis
- ▶ Incorrect results for world space normals

Scaling Normals



Fixes for non-uniform scale

- ▶ **Solution 1:** Don't use non-uniform scales!
- ▶ **Solution 2:** Use the *inverse transpose* of the world matrix when transforming normals
 - Rotation remains the same
 - Scale becomes inverted
- ▶ Requires sending one more matrix to shader
 - Already part of your transform class!

Shader code for normals

```
// VERTEX SHADER -----
// Cast matrix to a 3x3 - removing translation
// Then normalize result
output.normal = mul(input.normal, (float3x3)worldInvTrans);
output.normal = normalize(output.normal);

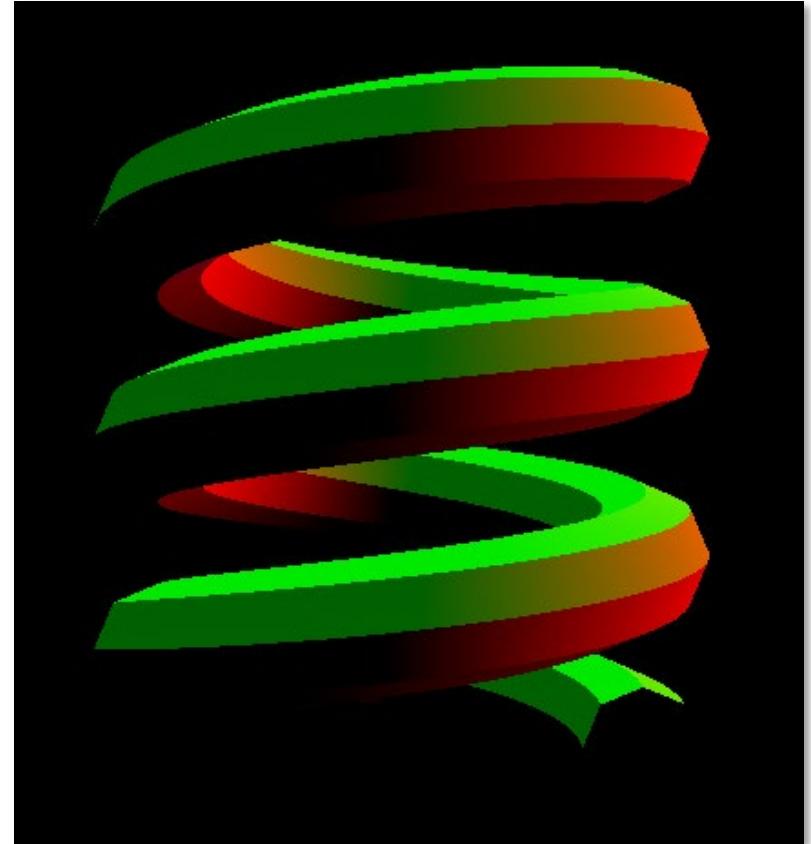
// PIXEL SHADER -----
// Must re-normalize interpolated normals! Interpolation results in
// non-uniform vectors, since each component is interpolated independently.
input.normal = normalize(input.normal);
```

Normals in the Pixel Shader

- ▶ Normal passed in from vertex shader
 - But...how do we know if it worked?
 - How do we check the data?

- ▶ Return the normal as if it were a color!
 - Normal is $\text{float3}(x,y,z)$
 - Return as our (r,g,b)

```
// Temporarily for testing...
return float4(input.normal, 1);
```



2: Light Sources



Light sources

- ▶ What about potential sources of light?
 - How do we simulate a large source of light like the sun?
 - Is everything just a light bulb?
 - What about flash lights?

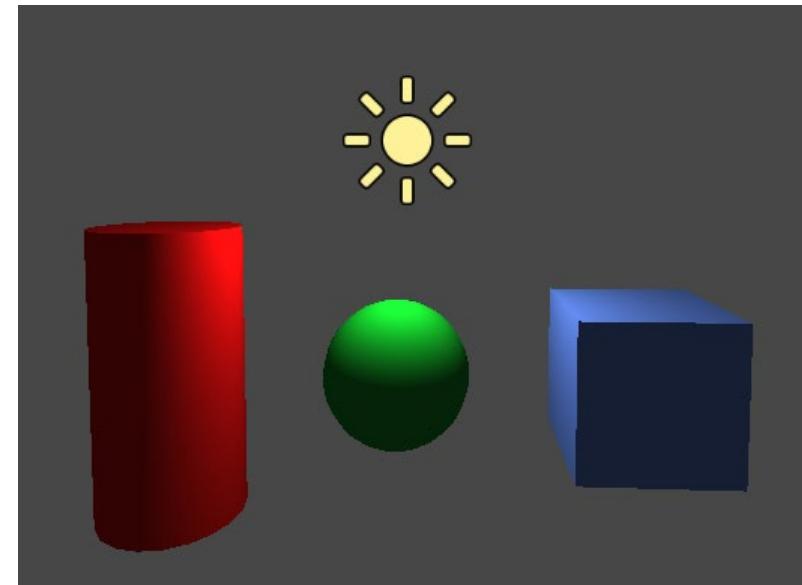
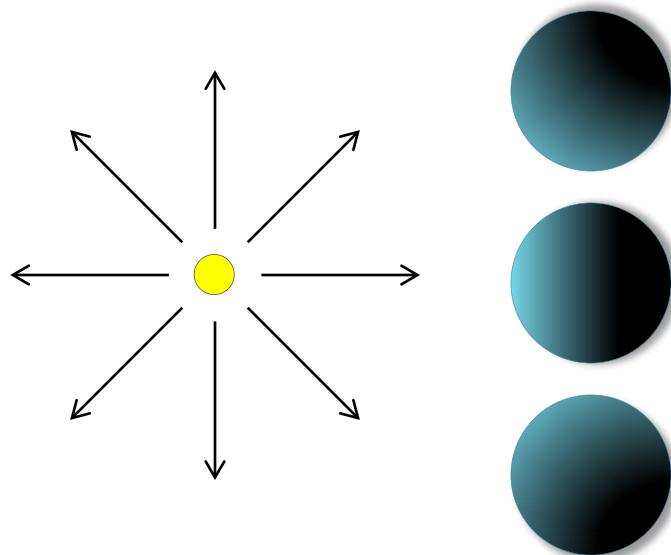
- ▶ We use *punctual* light sources
 - “Pertaining to or of the nature of a point”

Light source types

- ▶ **Directional Light**
 - All rays are parallel
 - Often simulates sunlight
- ▶ **Point Light**
 - Like a light bulb
 - Light “emits” from a point (in all directions)
- ▶ **Spot Light**
 - Like a flash light
 - Light “emits” from a point in a cone shape

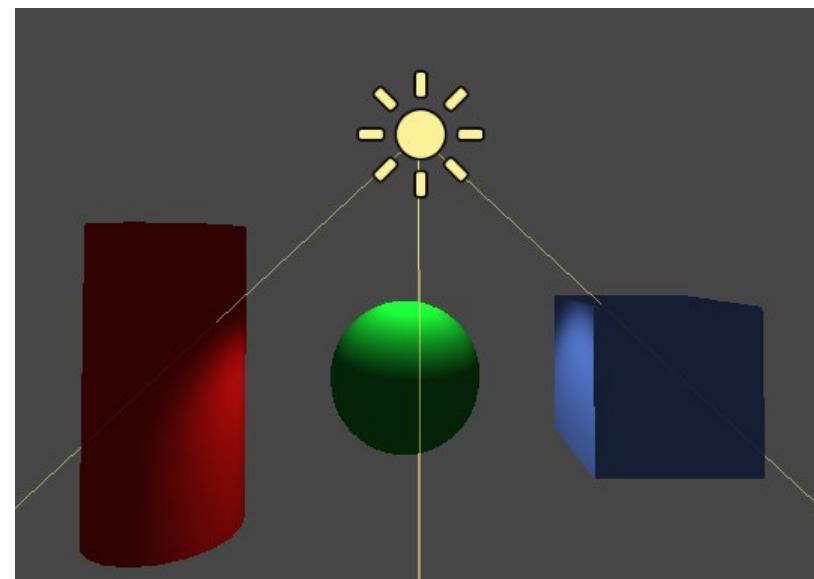
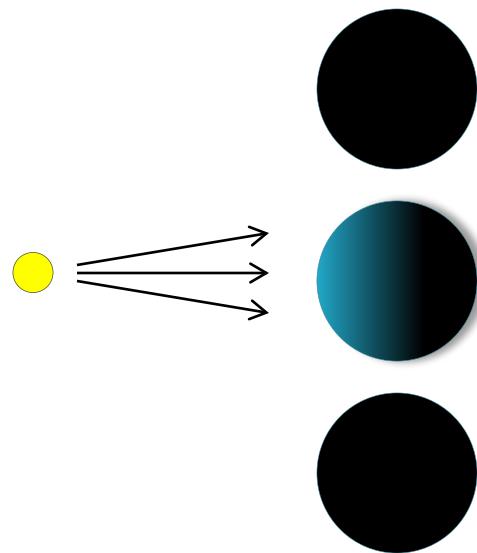
Point light

- ▶ All light rays come from a single point
 - Direction is calculated at each pixel
 - Vector between the light and the surface
 - Must pass *world* (not screen) position into pixel shader



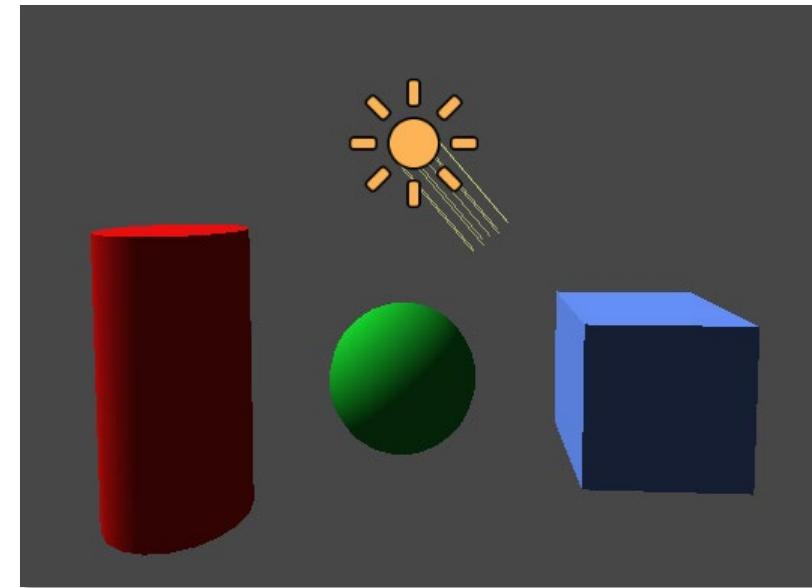
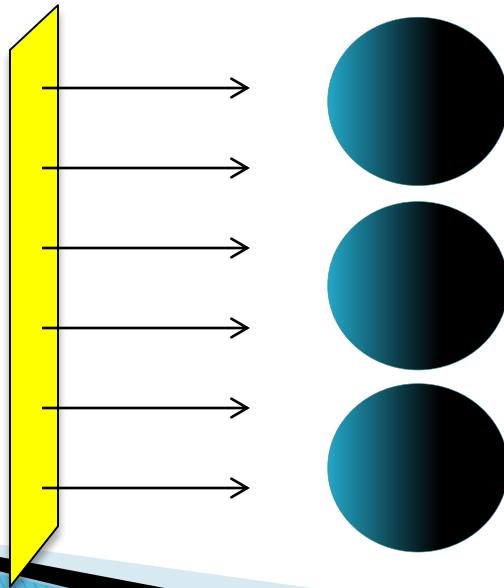
Spot light

- ▶ Like a flash light – Conical shape
- ▶ Similar to a point light, with restrictions on the angle



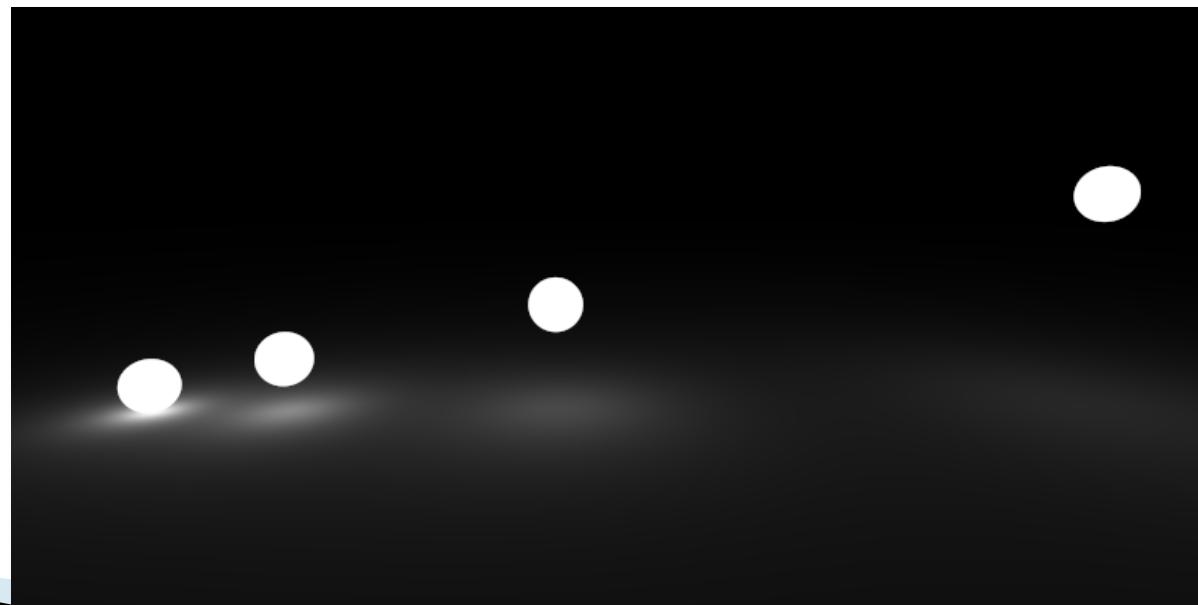
Directional light

- ▶ All light “travels” in the same direction (parallel)
- ▶ Uniform light direction (no calculation necessary)
 - Directional lights have no position!



Attenuation

- ▶ Light intensity should **attenuate** (weaken) with distance
- ▶ Otherwise light is always equally bright regardless of distance
- ▶ More distance = less light



Which lights should attenuate?



▶ Point & Spot lights

- Have a position
- Need a range, too
- Can properly attenuate



▶ Directional lights

- All light in same direction
- No “source position”
- Assumes the light goes on forever
- At uniform intensity
- Never attenuate

3. Lighting Equations

BRDFs – Our lighting equations

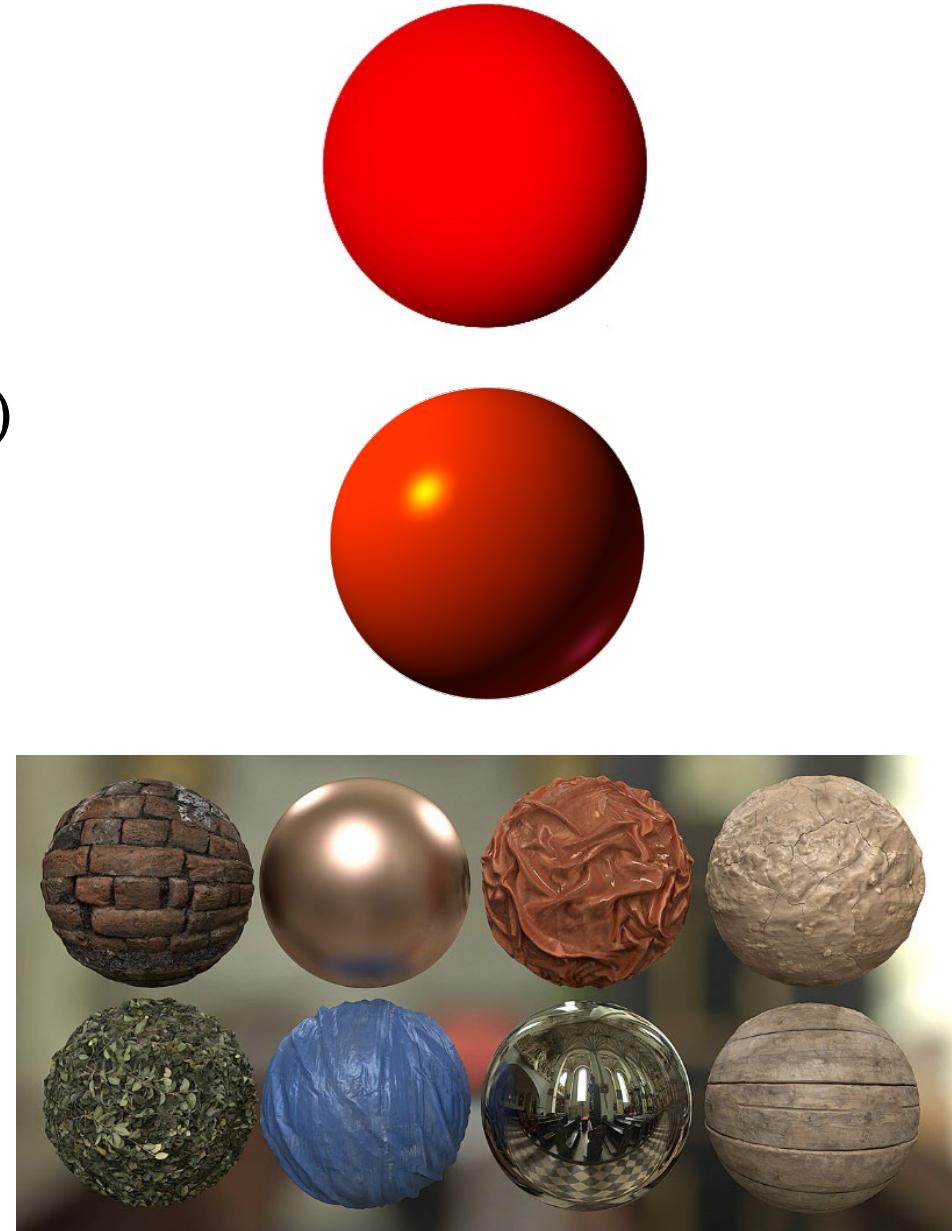
- ▶ Mathematical functions to calculate how much light reflects
- ▶ Bidirectional Reflectance Distribution Functions (BRDFs)
- ▶ Approximates how light reflects, given:
 - The incoming light direction
 - The surface normal
 - Our eye position (sometimes)

Where do BRDF's come from?

- ▶ Modeled after the physical world
- ▶ Each approximates a single property of light
 - Diffusion
 - Specularity
 - Anisotropy
 - Etc.
- ▶ We'll use a combination to make objects look realistically lit

Examples of BRDFs

- ▶ **Lambertian:** Basic, diffuse lighting
- ▶ **Phong:** Specularity (bright reflections)
 - **Blinn–Phong:** Similar to Phong
 - Less accurate but faster to compute
- ▶ **Cook–Torrance (Microfacet)**
 - Physically-based
 - Specular BRDF
 - Handles more surface properties



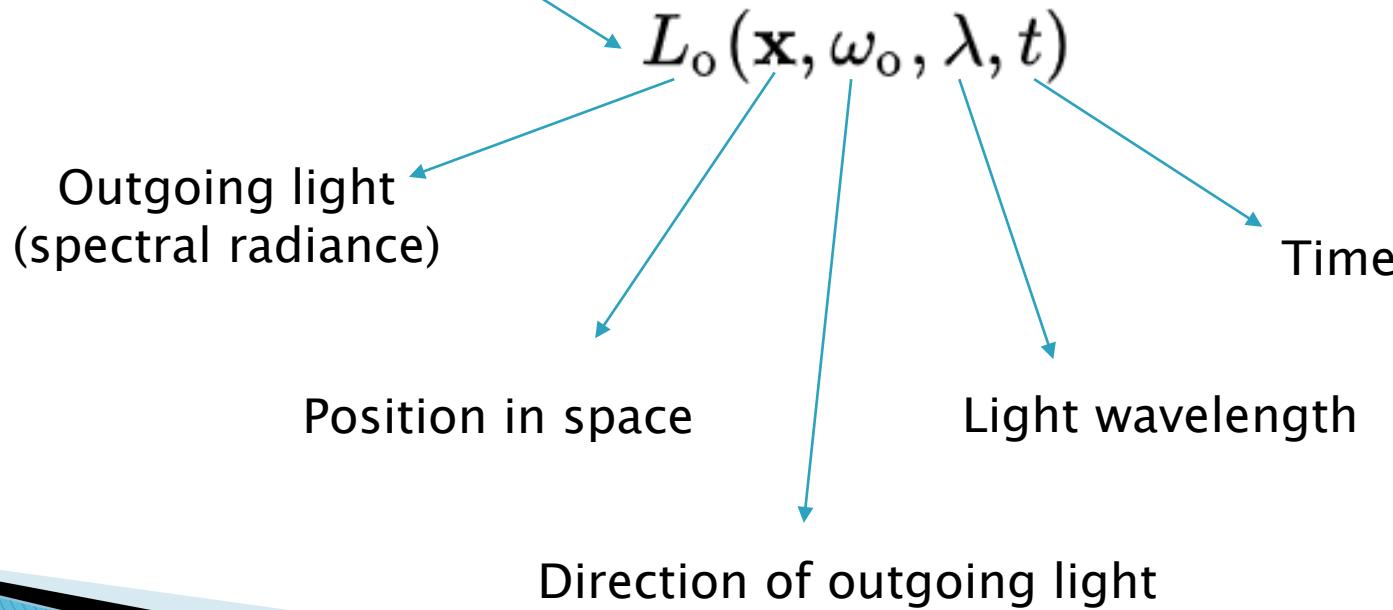
“The Rendering Equation”

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

- ▶ Introduced in 1986
 - Literally called “The Rendering Equation”
 - Original papers [here](#) and [here](#)
- ▶ Calculates the radiance leaving a point as the sum of emitted and reflected radiance
- ▶ Spoiler: We’re not doing all this math!!!

Deciphering the Rendering Equation: Output

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$



Deciphering: Emitted light

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

The diagram illustrates the components of the emitted light spectral radiance, $L_e(\mathbf{x}, \omega_o, \lambda, t)$. A blue bracket groups the first term, $L_e(\mathbf{x}, \omega_o, \lambda, t)$, and the integral term. Five blue arrows point from these components to their respective labels: "Emitted light (spectral radiance)" for the first term, "Position in space" for the spatial coordinate \mathbf{x} , "Light wavelength" for the wavelength λ , "Time" for the time variable t , and "Direction of outgoing light" for the integral over the domain Ω .

Emitted light (spectral radiance)

Position in space

Light wavelength

Time

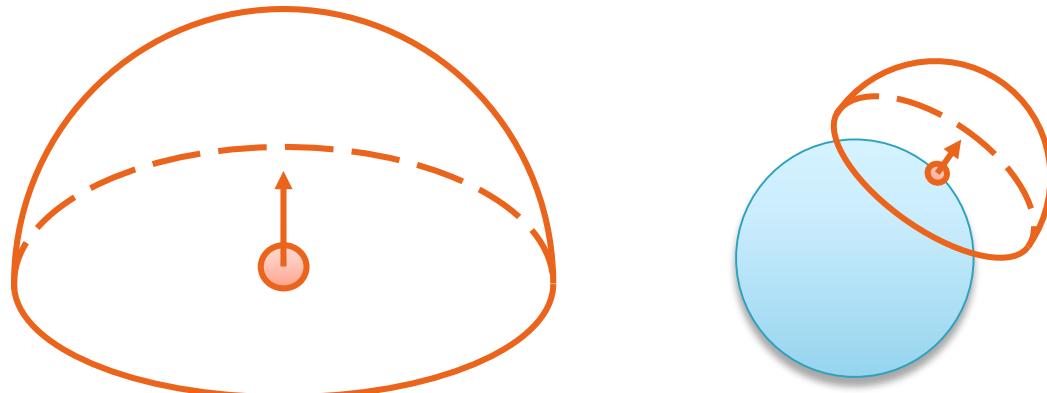
Direction of outgoing light

Deciphering: All incoming light

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

\int_{Ω}
Integral over Ω

Ω is a **hemisphere** encompassing
directions of all incoming light (ω_i)
that can strike this position in space



Deciphering: BRDF

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

$f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t)$

BRDF

Position in space

Negative direction*
of incoming light
(direction TO the light)

* All directions over Ω

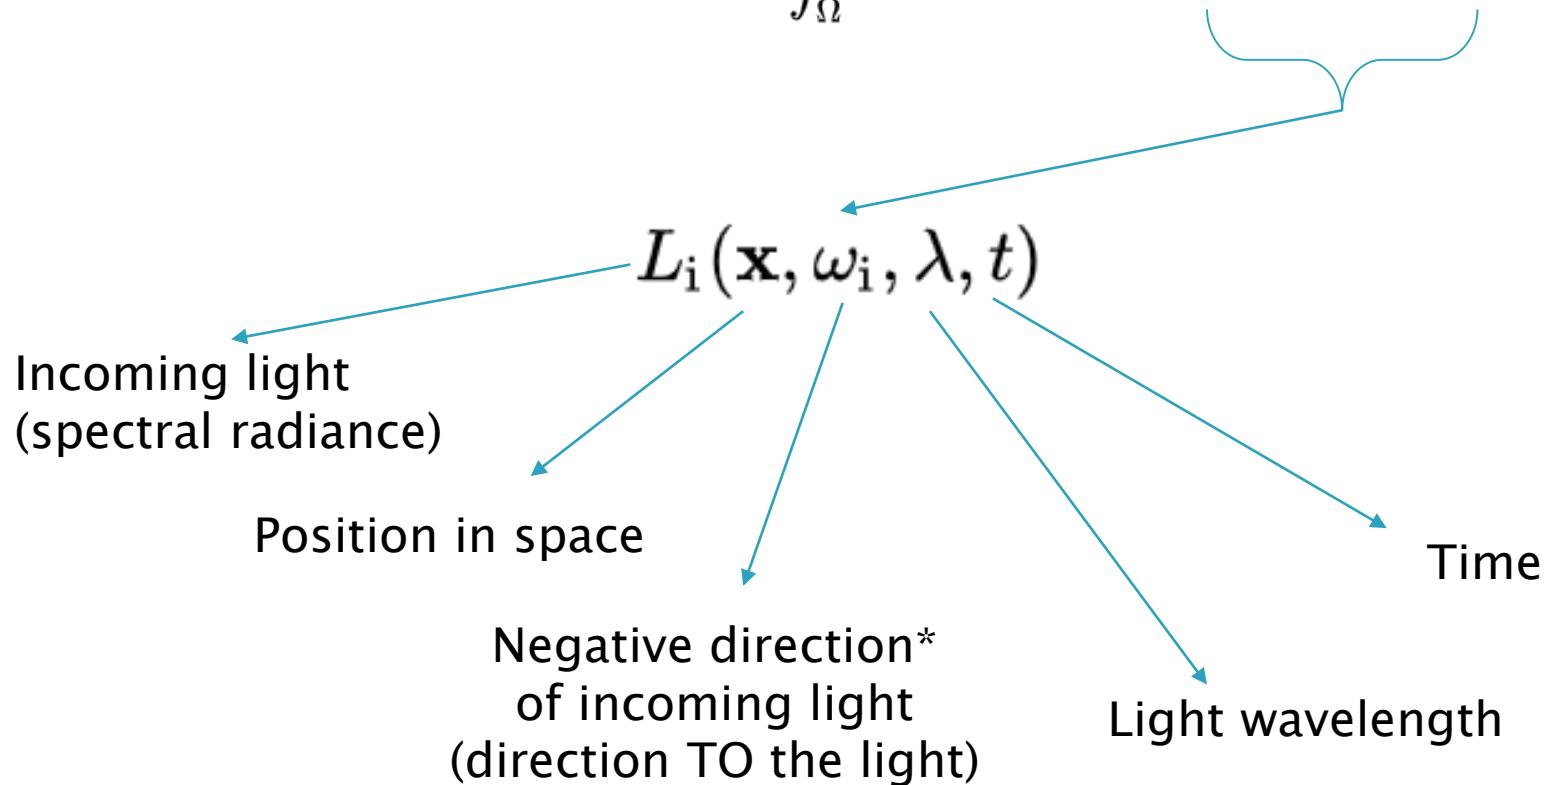
Light wavelength

Time

Direction of outgoing light

Deciphering: Incoming light

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$



* All directions over Ω

Deciphering: The rest

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

Negative direction*
of incoming light
(direction TO the light)

* All directions over Ω

Dot product

$(\omega_i \cdot \mathbf{n}) d\omega_i$

normal

Differential of negative
direction of incoming light
(solid angle)

The Rendering Equation for us

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

- ▶ We're obviously not using the whole thing
- ▶ Where are we simplifying?

The Rendering Equation for us

$$L_o(\mathbf{x}, \omega_o, \cancel{\lambda}, \cancel{t}) = \cancel{L_e(\mathbf{x}, \omega_o, \lambda, t)} + \cancel{\int_{\Omega}} f_r(\mathbf{x}, \omega_i, \omega_o, \cancel{\lambda}, \cancel{t}) L_i(\mathbf{x}, \omega_i, \cancel{\lambda}, \cancel{t}) (\omega_i \cdot \mathbf{n}) d\omega_i$$

- ▶ **Wavelength?** We assume RGB are similar
- ▶ **Time?** Single point in time
- ▶ **Emitted light?** Nope (but could be emissive texture)
- ▶ **Integral over a hemisphere?**
 - WAY too expensive for a shader!
 - We only handle light coming directly from a light source
 - And only a single direction per light
- ▶ Don't need that **solid angle** anymore, either

The Rendering Equation – What's left?

$$L_o(\mathbf{x}, \omega_o) = f_r(\mathbf{x}, \omega_i, \omega_o) L_i(\mathbf{x}, \omega_i)(\omega_i \cdot \mathbf{n})$$

BRDF (f), light source (L) and a dot product

- ▶ However, need to account for *all lights* in a scene
 - Multiple lights linearly increase brightness
 - So this is a simple sum of all results

$$L_o(\mathbf{x}, \omega_o) = \sum_{light=1}^{numLights} f_r(\mathbf{x}, \omega_i, \omega_o) L_i(\mathbf{x}, \omega_i)(\omega_i \cdot \mathbf{n})$$

The Rendering Equation as pseudocode

$$L_o(\mathbf{x}, \omega_o) = \sum_{\text{light} = 1}^{\text{numLights}} f_r(\mathbf{x}, \omega_i, \omega_o) L_i(\mathbf{x}, \omega_i) (\omega_i \cdot \mathbf{n})$$

- ▶ HLSL pseudocode based on the above formula:

```
float3 total = float3(0,0,0);
for (int i = 0; i < numLights; i++)
{
    Light light = lights[i];
    total += BRDF(light, surface) * LightColor(light) * dot(-light.Direction, normal);
}
return total;
```

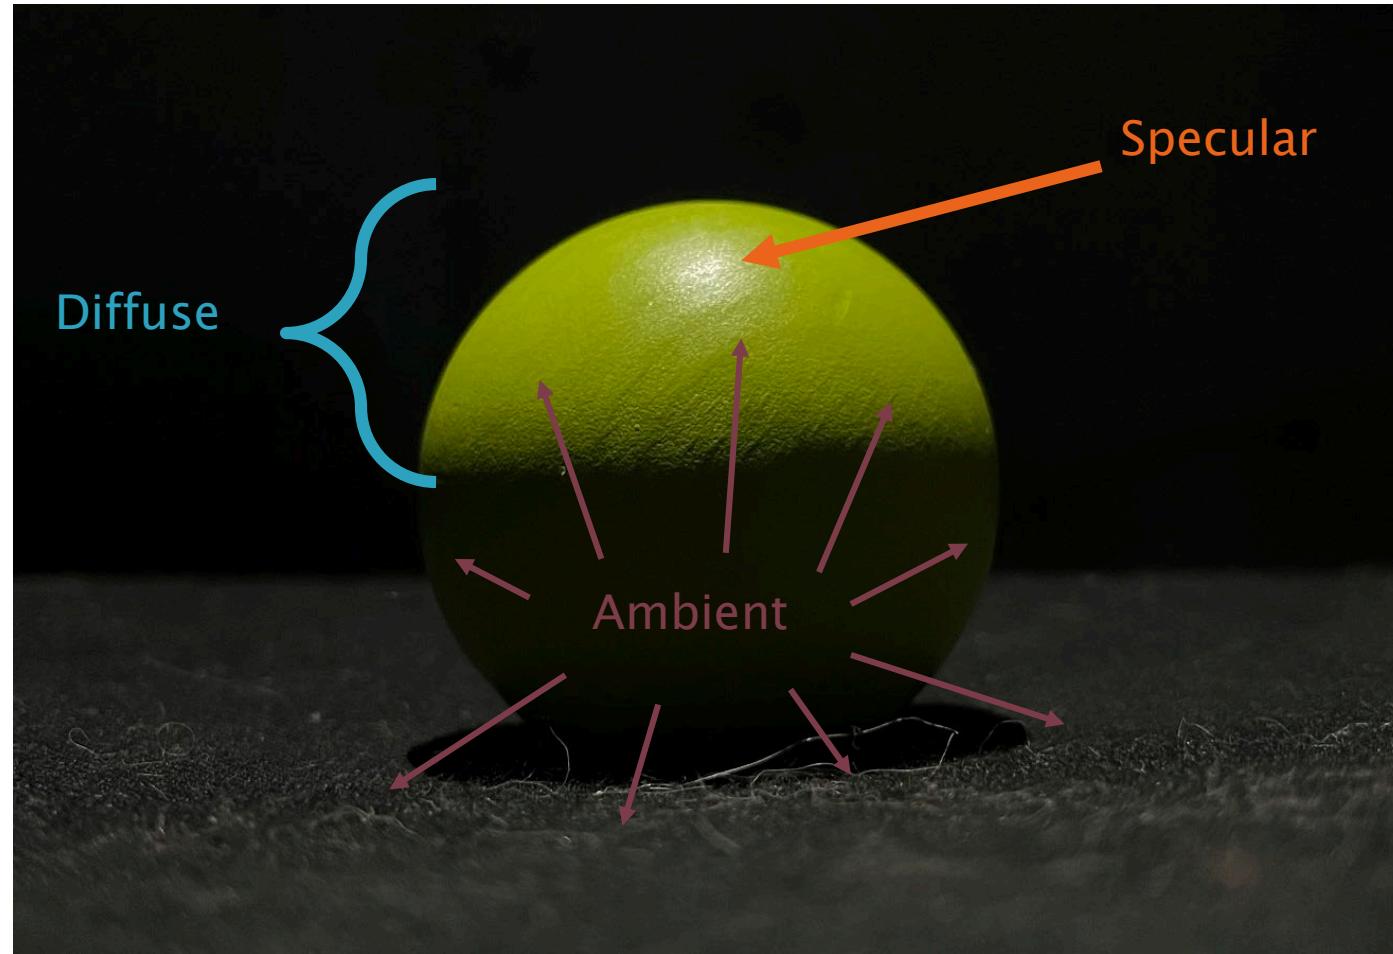
Lighting Approximations

And their BRDFs

Photo - What should we approximate?

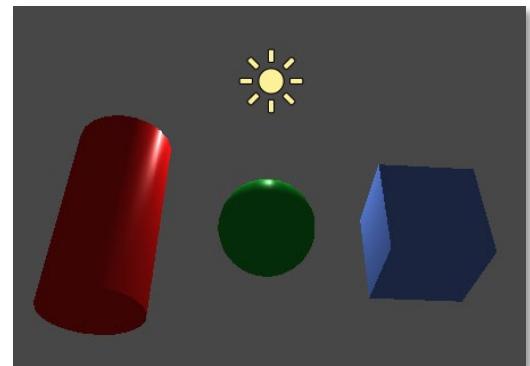
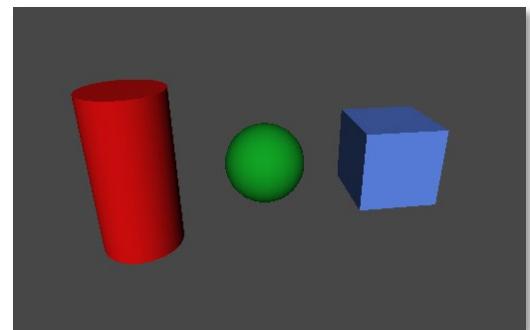
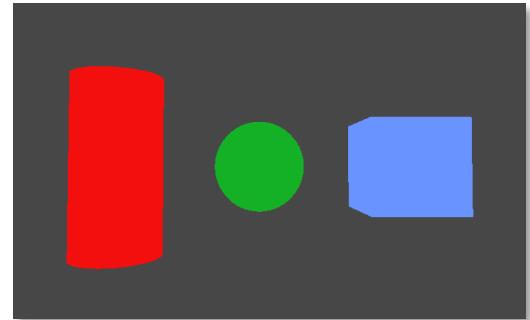


Classic computer graphics light approximations



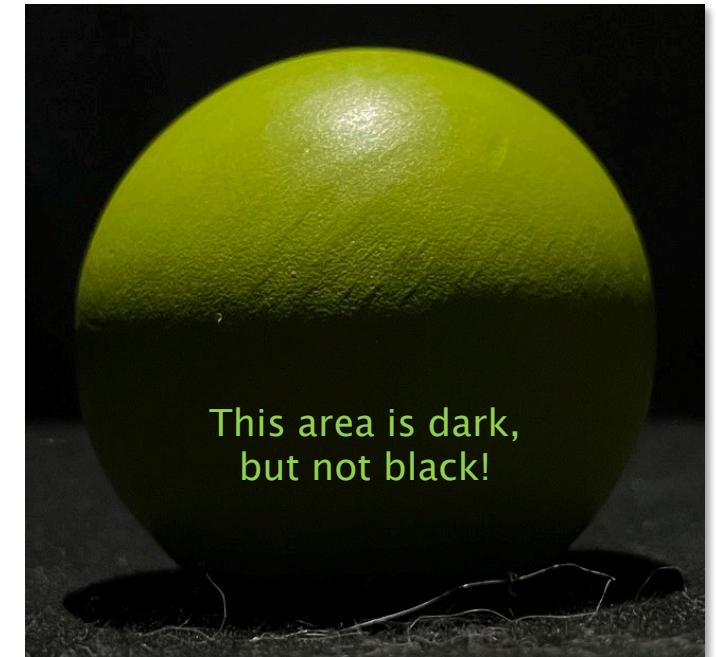
The big three

- ▶ Ambient
 - Approximates light bouncing around
 - *Incredibly* simplistic model of complex system
- ▶ Diffuse
 - Reflectance (brightness) of the surface
- ▶ Specular
 - Reflection of the light source



Ambient

- ▶ Approximation of light bouncing around
- ▶ Minimum light level & color of scene
- ▶ Implementation:
 - Pick a subtle color that matches the scene
 - That's your ambient color
- ▶ No actual “calculation” here
 - Gotta go fast!



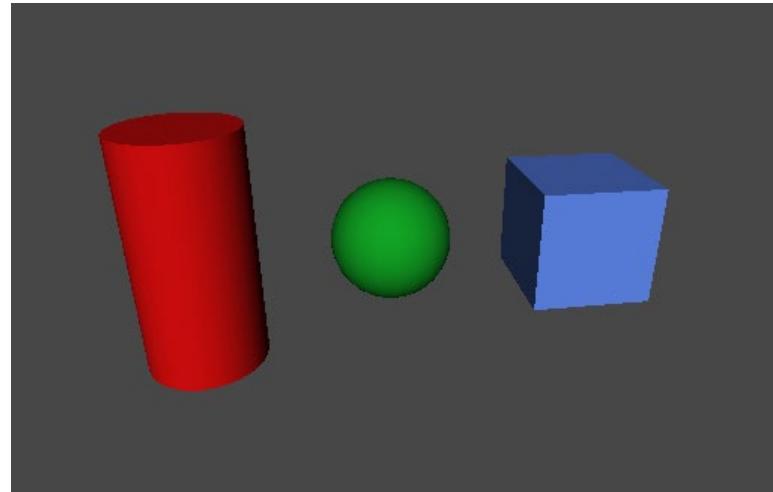
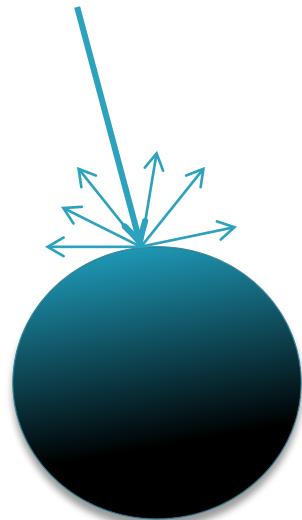
Shader - Ambient implementation

```
// Pretty simple
float3 ambientTerm = ambientColor * surfaceColor;

// ambientColor is coming from a constant buffer
// surfaceColor is currently object's color tint
//   - Will eventually come from a texture
```

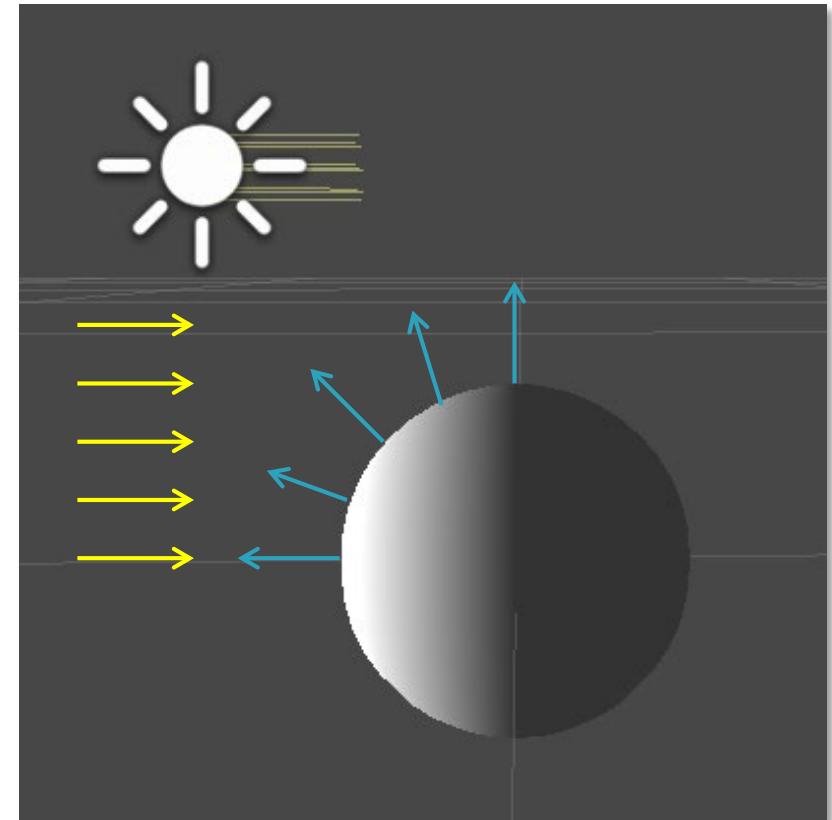
Diffuse light

- ▶ Hits & spreads out in all directions
- ▶ Looks the same from any viewing angle
 - Camera's position doesn't matter
 - Only light direction and surface normal matter



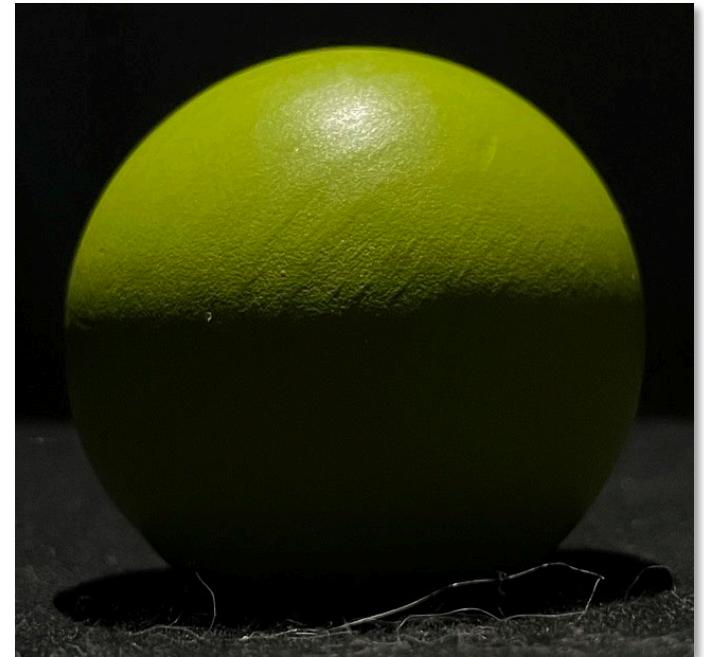
Diffuse lighting

- ▶ Compare surface normal & light direction
- ▶ Brighter where light hits “head on”
- ▶ Less bright where light “glances off”
- ▶ Lambert’s Cosine Law
 - “ $N \cdot L$ ”



Diffuse - Lambertian reflectance

- ▶ Approximates light leaving a perfectly diffuse surface
- ▶ BRDF: $I_D = \mathbf{L} \cdot \mathbf{N} C I_L$.
 - \mathbf{L} is normalized vector from surface to light
 - \mathbf{N} is surface normal
 - C is surface color
 - I_L is light color & intensity
- ▶ Implementation:
 - Dot product between two unit vectors
 - Multiply by light (color & intensity) and surface color!



Shader - Diffuse implementation

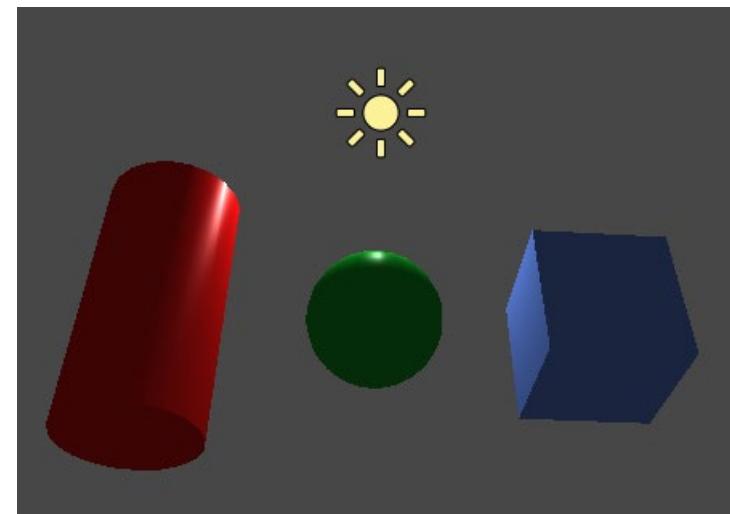
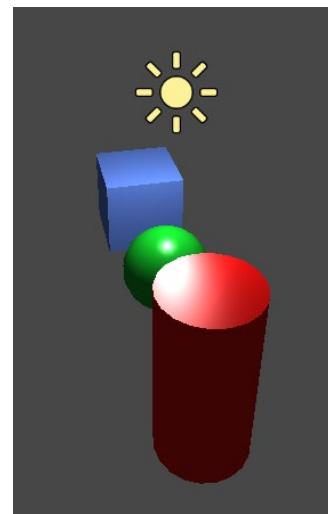
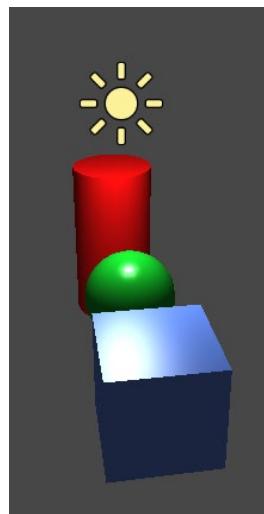
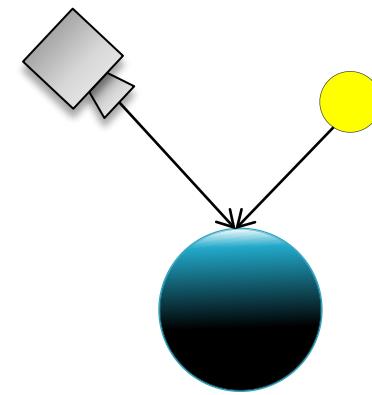
```
// Get direction TO the light (opposite of light's direction)
float3 dirToLight = -lightDirection;

// Lambertian reflectance is really just a dot product
float3 diffuseTerm =
    saturate(dot(input.normal, dirToLight)) * // Diffuse intensity, clamped 0-1
    lightColor * lightIntensity * // Light's overall
    surfaceColor; // Light tints the surface

// normal is pixel shader input (from rasterizer)
// light's direction, color & intensity from cbuffer
// surfaceColor is the material's color tint (will eventually come from a texture)
```

Specular light

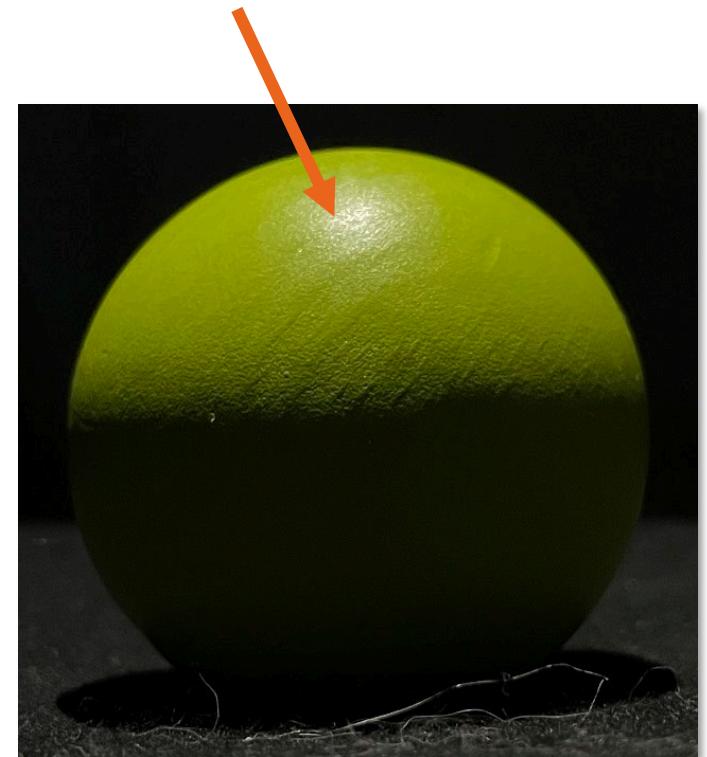
- ▶ Simulates “shine” on a surface
 - Basic reflection of the light source
- ▶ Looks different based on your viewing angle



Specular – Phong

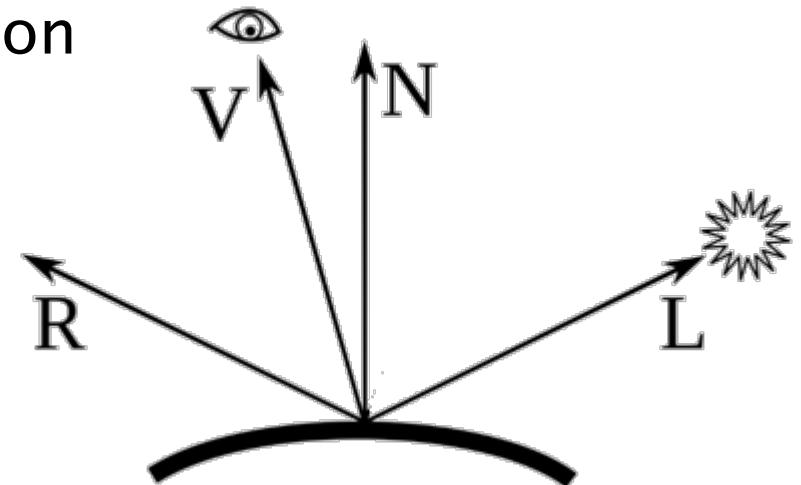
- ▶ Approximates shiny reflection
- ▶ BRDF: $k_s(\hat{R}_m \cdot \hat{V})^\alpha i_{m,s}$
 - k_s is specular constant (1.0 usually)
 - R_m is normalized light reflection vector
 - V is normalized vector from surface to viewer
 - α is material's shininess constant
 - i_m is light intensity
 - i_s is specular intensity
- ▶ Implementation
 - Dot product of vectors, raised to a power
 - Multiply by intensities

Just the shiny bits



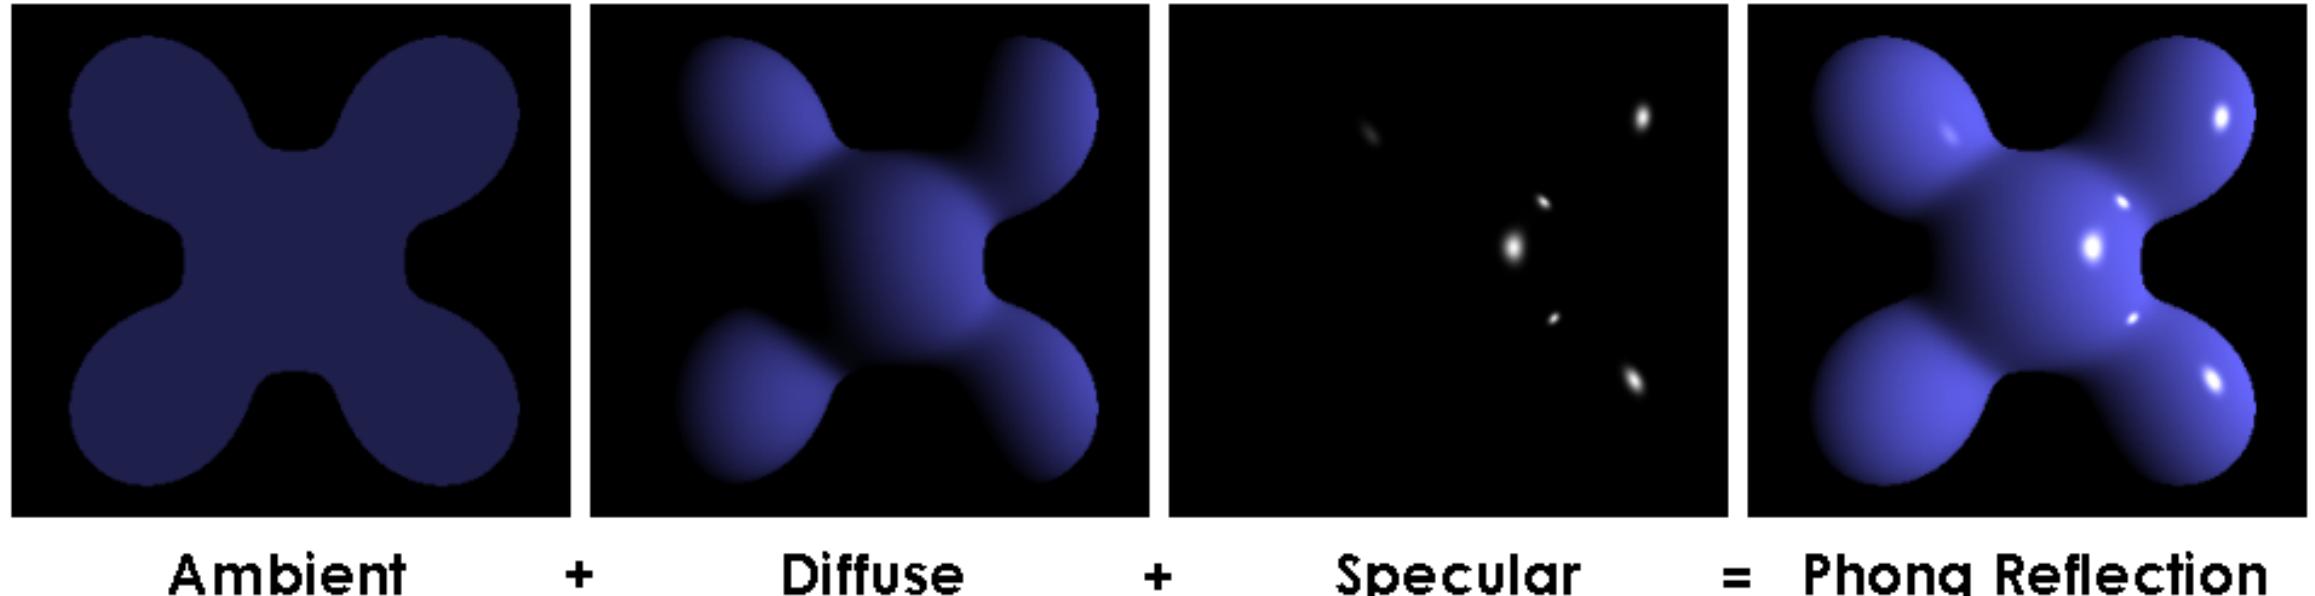
Phong reflection model

- ▶ Calculate the reflection vector (R)
 - Based on the Normal (N) and Light (L) vectors
- ▶ Compare reflection vector to view vector
 - Use dot product to get cosine of angle
 - Within certain angle, we should “see” reflection
- ▶ Raise result to a power for a smooth falloff



Phong reflectance results visualized

- Three separate approximations, added together



Shader – Phong implementation

```
// Calculate the reflection of the INCOMING light dir  
// using the surface normal (ensure both are normalized)  
float3 refl = reflect( lightDir, input.normal );  
  
// Calculate cosine of the reflection and camera vector  
// - saturate() ensures it's between 0 and 1  
float RdotV = saturate( dot( refl, dirToCamera ) );  
  
// Raise the that to a power equal to some “shininess”  
// factor - try large numbers here (64, 200, etc.)  
float3 specularTerm = pow(RdotV, shininess) *  
    lightColor * lightIntensity * surfaceColor;
```

Total light for a single pixel

- Overall calculation looks like this:

$$I_p = i_a + \sum_{m \in \text{lights}} ((\hat{L}_m \cdot \hat{N}) i_{m,d} + (\hat{R}_m \cdot \hat{V})^\alpha i_{m,s})$$

Final light intensity for a pixel Ambient term Add up the following for each light

Diffuse (Lambert) Specular (Phong)

Shader – Final shading value

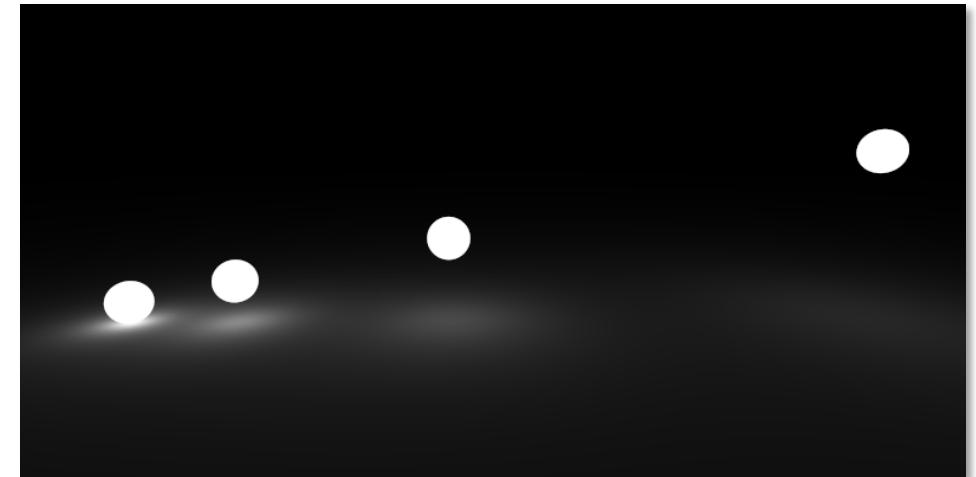
```
float3 totalShading = ambientTerm;  
  
// Do this once PER LIGHT  
// Each light will have its own diffuse and specular result  
totalShading += diffuseTerm + specularTerm;  
  
// Return this from our pixel shader  
return float4(totalShading, 1);
```

Attenuation



Attenuation

- ▶ Light intensity should **attenuate** (weaken) further from source
 - Otherwise light is always equally intense
 - Greater distance = less light
- ▶ Attenuate point & spot lights
- ▶ Can't attenuate directional
 - Has no position
 - Can't measure distance
- ▶ But how much?



Attenuation functions

- ▶ Many exist, all have the same effect:
 - Result decreases as distance increases
- ▶ $1.0 / (1.0 + a * \text{dist} + b * \text{dist} * \text{dist})$
 - Classic computer graphics attenuation function
 - “a” is linear falloff value
 - “b” is quadratic falloff value
 - “dist” is distance between light and pixel
 - This one *isn't great* for our purposes!

Common attenuation example

- ▶ Tweakable and looks fine, but...
 - Parameters don't represent exact range
 - Doesn't hit zero at a finite distance



Better attenuation function?

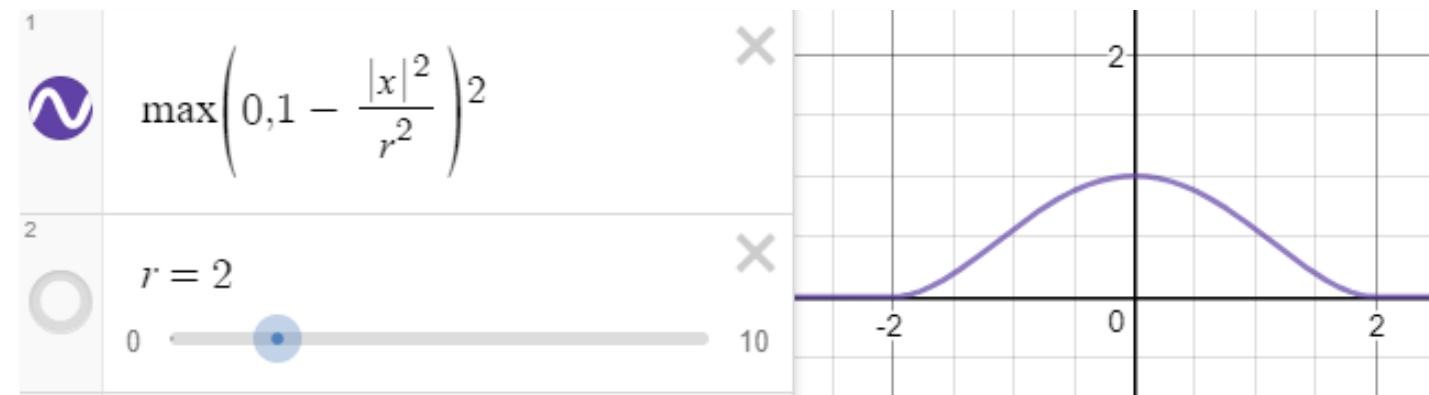
- ▶ We need a *range-based* function
- ▶ An attenuation function that...
 - Goes to zero at a *specific distance* (range)
 - Still has a nice smooth falloff
 - Ideally isn't linear

$$\max(0, 1.0 - \text{dist}^2 / \text{range}^2)^2$$

Range-based attenuation

- ▶ Example in HLSL
 - dist is actual distance between pixel and light
 - radius is light's "range"

```
att = saturate(1.0-dist*dist/(radius*radius));  
att *= att;
```



Applying attenuation

```
float3 totalShading = ambientTerm;  
  
// Do this once PER LIGHT  
// Each light will have its own diffuse, specular and atten.  
totalShading += (diffuseTerm + specularTerm) * attenuation;  
  
// Return this from our pixel shader  
return float4(totalShading, 1);
```

Representing Lights in Code



Sending light data to the shader

- ▶ A single light is made up of multiple pieces of data

Directional Light:

- Color
- Intensity
- Direction

Point Light:

- Color
- Intensity
- Position
- Range

Spot Light:

- Color
- Intensity
- Position
- Direction
- Range
- Spot Falloff

- ▶ How do we represent and send all this data to the shader?

Sending light data...poorly

```
cbuffer ExternalData : register(b0)
{
    float3 dirLight1Color;
    float  dirLight1Intensity;
    float3 dirLight1Direction;
    float3 dirLight2Color;
    float  dirLight2Intensity;
    float3 dirLight2Direction;
    float3 pointLight1Color;
    float  pointLight1Intensity;
    float3 pointLight1Position;
    float  pointLight1Range;
    // And on and on and on...
}
```

Please, I beg you, no...

Use a common struct for all lights!

```
#define LIGHT_TYPE_DIR      0  
#define LIGHT_TYPE_POINT    1  
#define LIGHT_TYPE_SPOT     2
```

Can use these in your code
when checking a light's Type

```
// Make a matching struct and #defines on CPU (C++) side  
struct LightStruct  
{  
    int      Type;  
    float3   Direction;      // 16 bytes  
    float    Range;  
    float3   Position;       // 32 bytes  
    float    Intensity;  
    float3   Color;          // 48 bytes  
    float    SpotFalloff;  
    float3   Padding;         // 64 bytes  
};
```

Has all properties to represent
any light we want, and a Type
so we can dynamically handle
each one.

Padding at the end ensures it
doesn't cross the 16-byte
boundary.

Sending light data with structs

```
cbuffer ExternalData : register(b0)
```

```
{
```

```
    Light dirLight1;
```

```
    Light dirLight2;
```

```
    Light dirLight3;
```

```
    Light pointLight1;
```

```
    Light pointLight2;
```

```
}
```

You know, an array would make this even easier!

More on that in the future

```
// Meanwhile, in C++...
```

```
ps->SetData("dirLight1", &myLight, sizeof(Light));
```

Spot Light Example



Spot Light Calculation

```
// Start with the POINT LIGHT calculation

// Calculate the falloff from the center of the
// spot light by comparing the spot light's
// true direction (center of the cone) to the
// direction from the light to this pixel.
// Remember: LightDir is direction to THIS PIXEL.
float angleFromCenter = max(dot(-lightDir, light.direction), 0.0f);

// Raise to a power for a nice “falloff”.
// Multiply the diffuse and specular results by this:
float spotAmount = pow(angleFromCenter, light.spotPower);
```