

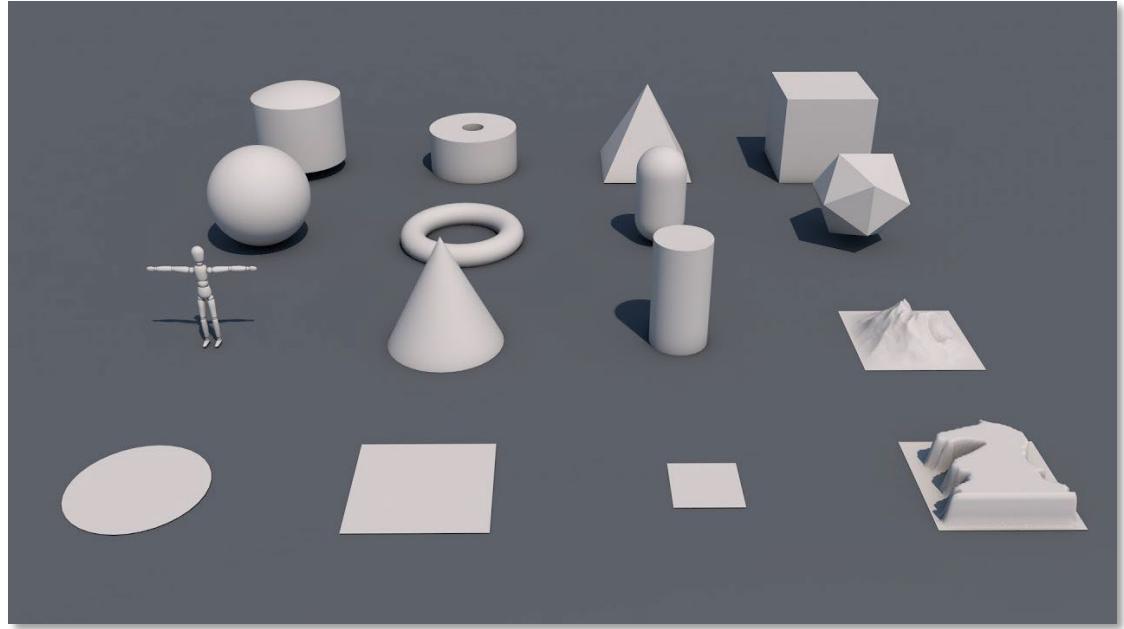
Real-Time Graphics

An overview

Computer graphics

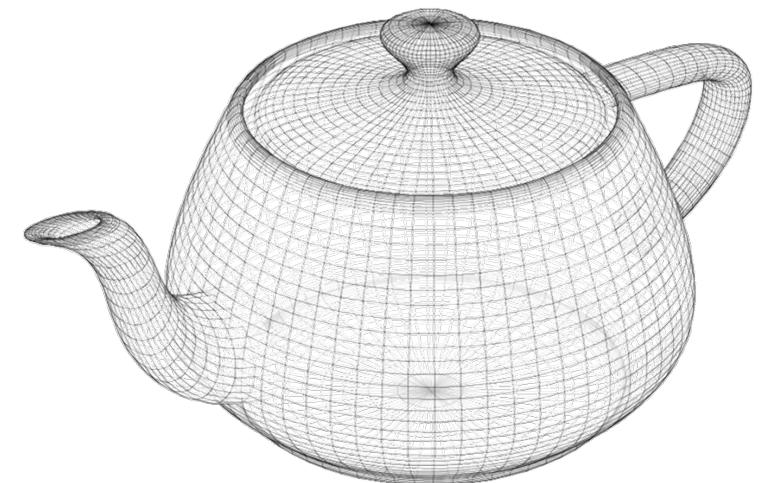
- ▶ Encompasses many topics
 - Games
 - Film
 - Photography
 - Image processing
 - Computer vision
 - Etc.

- ▶ Our focus: real-time graphics for games
 - Rendering a 3D world at 60+ frames per second
 - Using modern GPU hardware



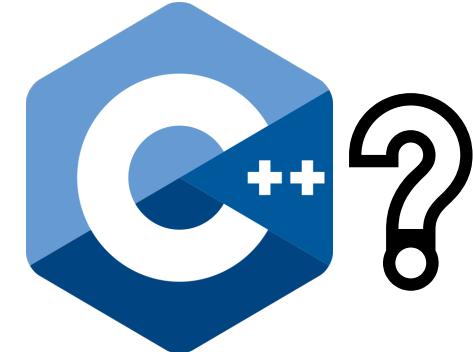
3D rendering

- ▶ Generating an image from one or more 3D models
- ▶ 3D models comprised of triangles
 - Each triangle has 3 vertices
 - Each vertex contains surface information
- ▶ Special hardware not *strictly* necessary
 - Can be done completely in software
 - C++, C#, Javascript, any language



3D rendering for games

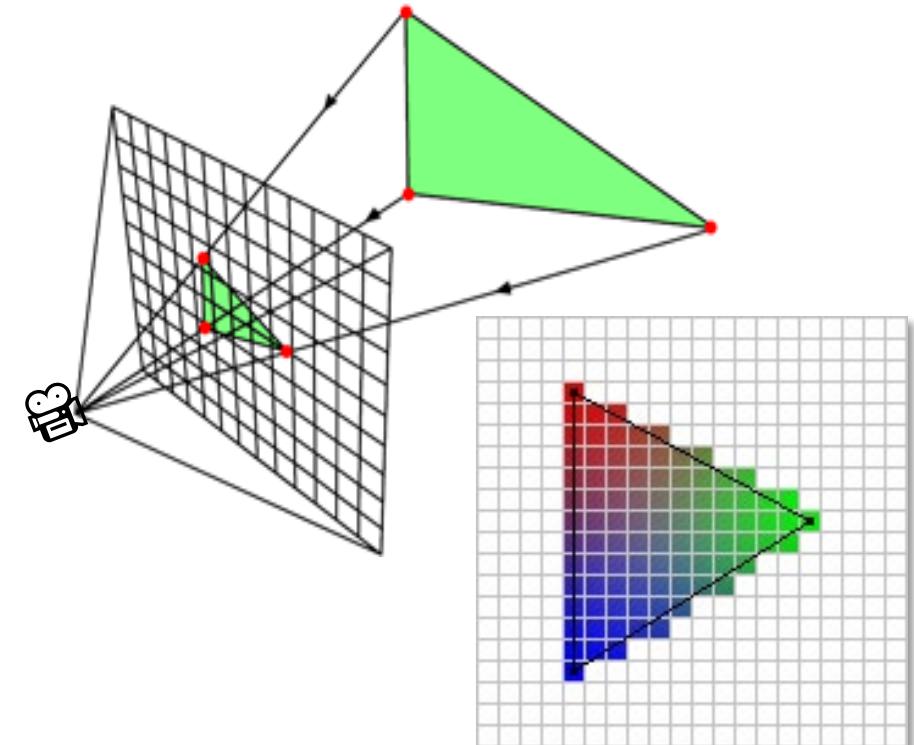
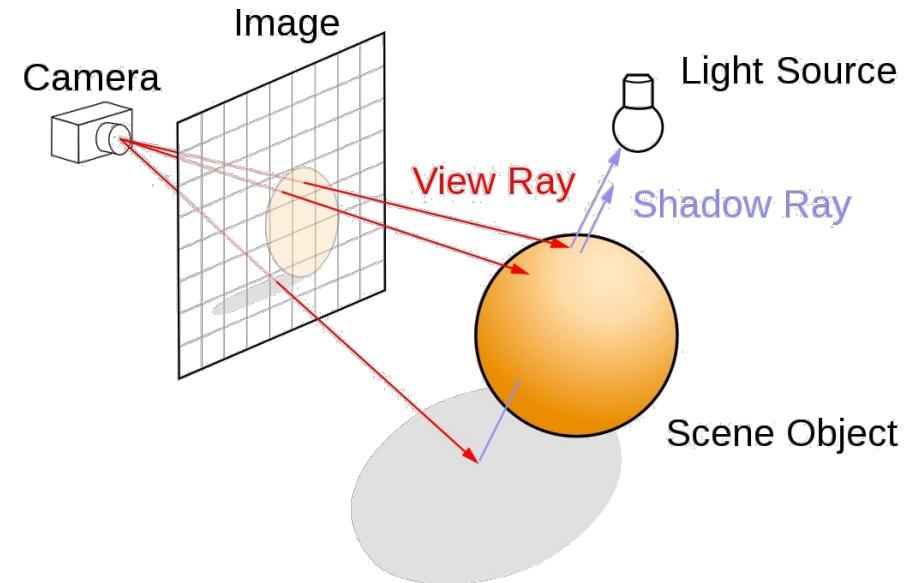
- ▶ We *can* perform rendering in software
 - Often called “offline rendering”
 - But games generally *don’t*
 - Too slow!
- ▶ Games utilize features of modern GPUs
 - Speeds up the rendering process
 - Allows for real-time rendering
 - “Interactive computer graphics”



Rendering techniques

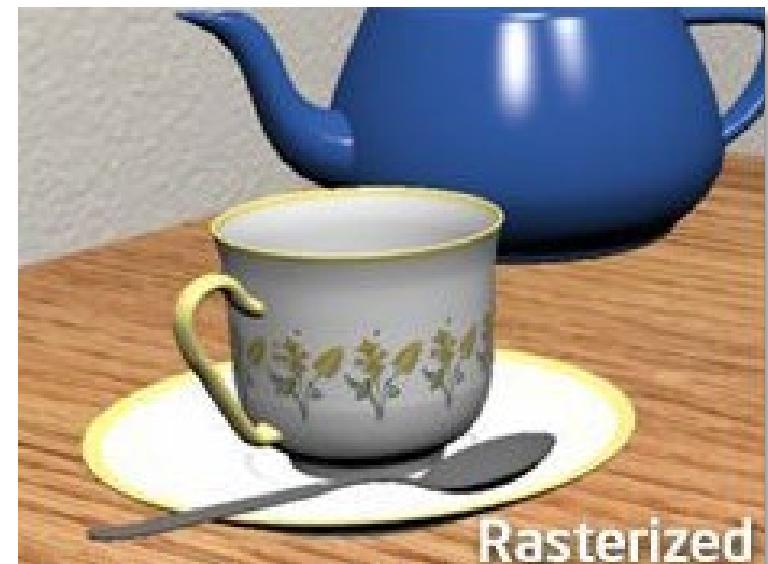
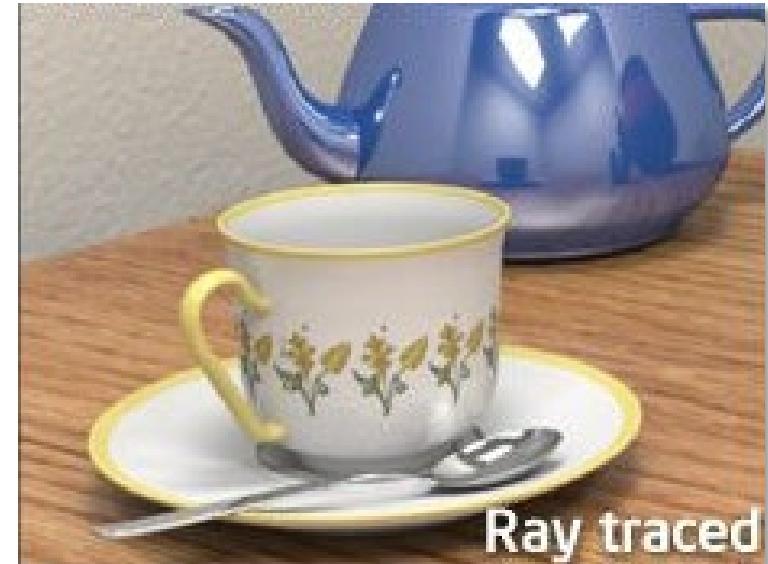
- ▶ Raytracing
 - Calculate rays through screen pixels
 - Perform ray/shape intersections
 - Use surface info from closest hit
 - Can recursively trace new rays

- ▶ Rasterization
 - Project vertices onto screen pixels
 - Determine pixels within triangle's bounds
 - Interpolate vertex data for each pixel



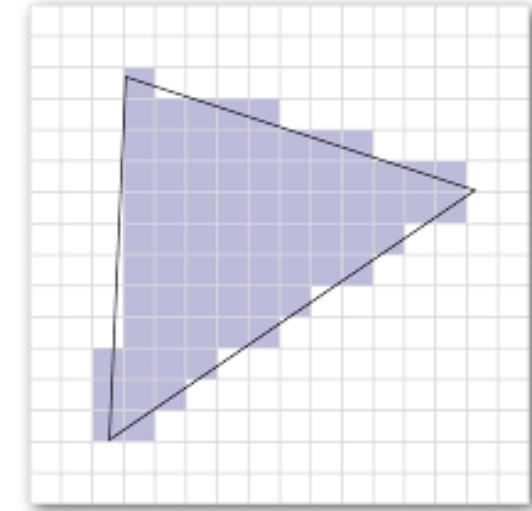
Raytracing vs. rasterization

- ▶ Both can produce high quality images
- ▶ Raytracing (slower)
 - Has access to entire scene
 - Inherently handles reflections
- ▶ Rasterization (faster)
 - One object at a time
 - Only has surface information
 - (Additional techniques can improve results)



For speed, rasterization wins

- ▶ Generally, rasterization is faster than ray tracing
- ▶ GPUs have long been built for rasterization
 - Contain specialized raster hardware
 - Can rasterize many, many triangles quickly
- ▶ All modern graphics APIs handle rasterization
- ▶ Until recently, every game out there was rasterized

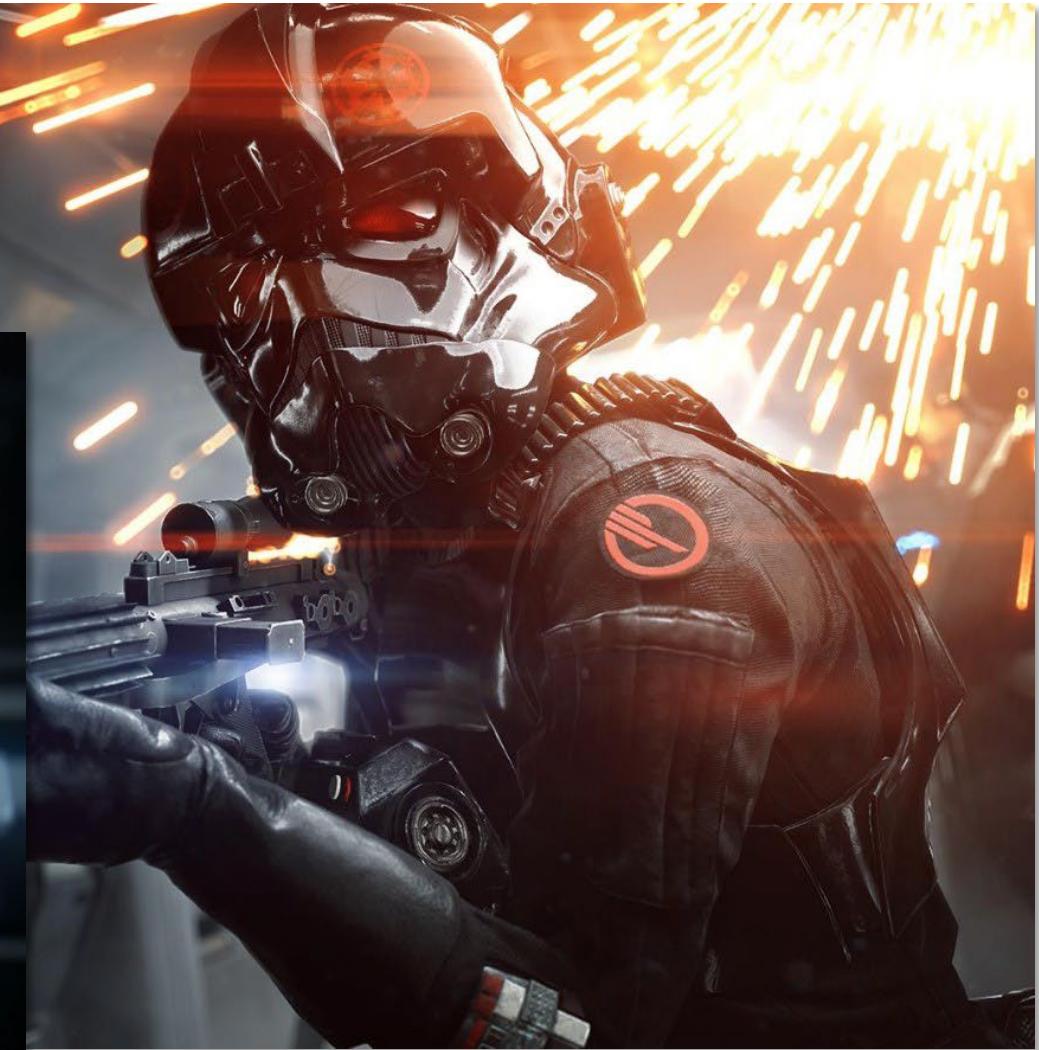


Sidenote: Real-time raytracing?



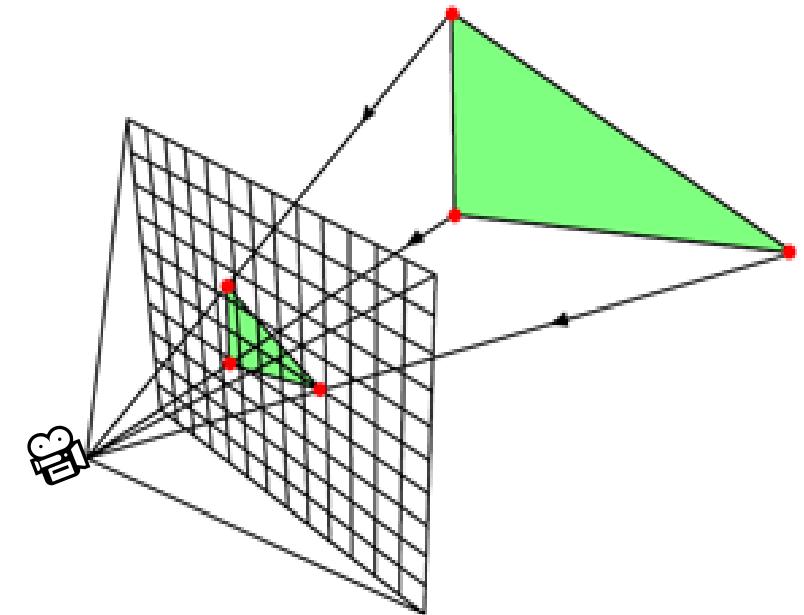
- ▶ Hardware accelerated raytracing
- ▶ Really impressive
 - But not a silver bullet
 - Still expensive in terms of performance
 - Often used in conjunction with rasterization
- ▶ Requires a pretty deep understanding of a low-level API
 - Such as DirectX 12 or Vulkan
 - And beefy hardware

So is rasterization bad? Not at all



Okay, rasterization it is. How is it done?

- ▶ Take one or more triangles
 - Vertices contain surface information
 - Color, surface normal, texture coords, etc.
- ▶ Project vertices onto a plane
 - The image we're creating
 - Also the near plane of our 3D camera
 - Need some matrix math here
- ▶ Determine which pixels are covered
- ▶ Interpolate triangle data and use it to color the pixels



General pseudocode of rasterization

- ▶ Foreach triangle
 - Project 3 vertices onto 2D plane
- ▶ Foreach projected triangle
 - Rasterize into pixels
- ▶ Foreach rasterized pixel
 - Determine color
 - Textures, lighting equations, etc.
- ▶ Lots of O(n) loops here 😬

- What if triangles share vertices?
 - Maybe we can denote unique vertices?
-
- What if two triangles overlap?
 - Like the front and back of a sphere
-
- Should this color replace other colors we've already rendered?
 - What about transparency?
 - What if we've already rendered something closer than this object?

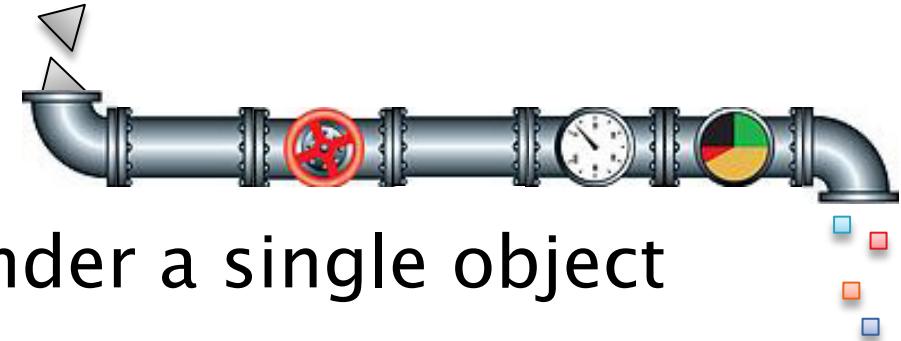
Let's label these steps

- ▶ 1. Input organization → handle duplicate vertices, etc.
- ▶ 2. Vertex transformation → project vertices onto plane
- ▶ 3. Rasterization → determine which pixels are covered
- ▶ 4. Pixel coloring → what color is the surface at this pixel?
- ▶ 5. Merge → Combine colors based on occlusion, blending, etc.

- ▶ This is known as a “rendering pipeline”
 - Input is triangles
 - Output is pixel colors



Rendering pipeline



- ▶ The set of steps to go through to render a single object
- ▶ What is “a single object”?
 - One distinct mesh of triangles
 - With a single surface material
- ▶ What if I want to render multiple objects?
 - Do it again! (mesh → rasterize → pixels)
 - And again (mesh → rasterize → pixels)
 - And again (mesh → rasterize → pixels)

- Sidenote: Does your 3D model use two separate textures?
 - Those are *two different materials*
 - Must split the model and render each independently

Speeding up a rendering pipeline

- ▶ Must go through all 5 steps *in order*
 - How can we speed them up?
 - Either do less work or get the work done faster
- ▶ It's a bunch of loops, right?
 - Foreach triangle...
 - Foreach pixel...
 - Etc.
- ▶ Can we do the work of each loop in parallel?
 - Process ALL triangles at once
 - Then process ALL pixels at once

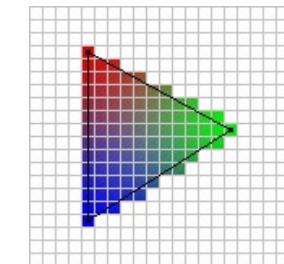
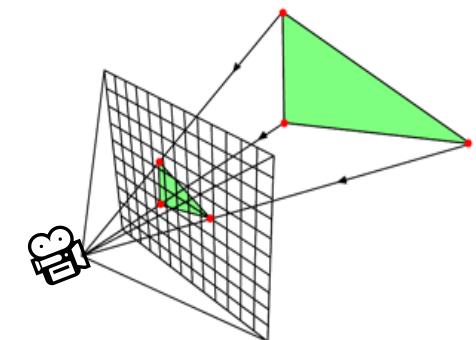


GPUs to the rescue

- ▶ Each pipeline step can be parallelized on the GPU
 - Built for highly parallel number crunching
 - And primitive rasterization (points, lines & triangles)

- ▶ For example:
 - Each vertex can be transformed simultaneously
 - Vertex A's transformation does not affect Vertex B's

 - Each pixel can be colored simultaneously
 - Color of Pixel A does not affect color of Pixel B



Sidenote: Doing less work?

- ▶ Using a GPU is obviously a big win
- ▶ But doing less work overall is also a great optimization
 - **Frustum culling** → don't draw what the camera can't see
 - **Batching** → draw multiple of the same thing at once
 - **Material IDs** → pre-process meshes and encode material IDs at each vertex, then dynamically index textures at runtime
- ▶ Advanced topics – for now, we'll just draw everything

How do we “use a GPU”?

- ▶ Directly accessing the hardware isn’t really feasible
- ▶ Using the drivers directly isn’t, either
 - Each driver is different
 - And has very specific requirements
- ▶ Need an Application Programming Interface (API) for graphics
 - A standardized way of using a GPU
 - That takes care of talking to the underlying hardware
 - And makes our lives (a little) easier

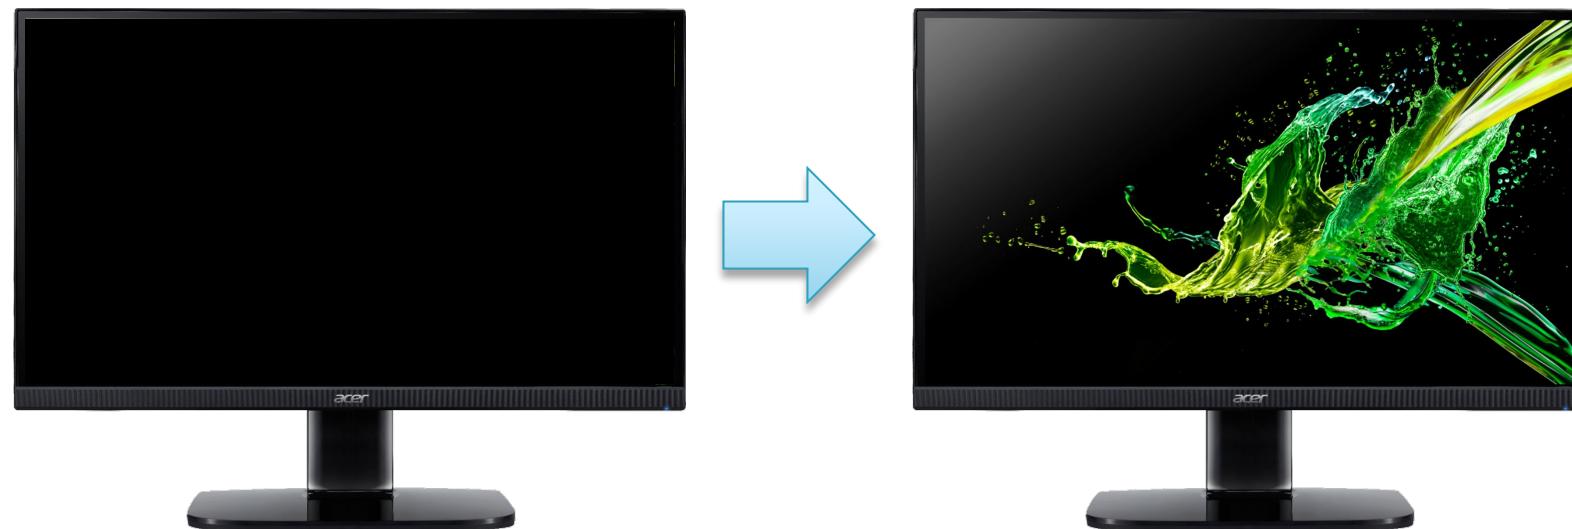
Graphics API

- ▶ System for utilizing graphics hardware for real-time rendering
- ▶ Several exist
 - DX11, DX12, OpenGL, Vulkan, Metal
 - Each is syntactically different
- ▶ Each has the same goal → Put something on the screen!
 - And mostly have the same features
 - Since they target the same hardware



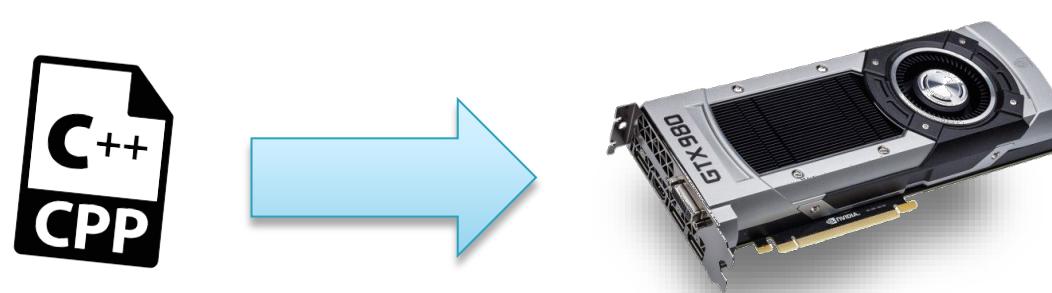
What's a Graphics API do?

- ▶ They all have the same goal: Rendering
 - Putting something on the screen
 - Also called “drawing”



What's a Graphics API really?

- ▶ C++ code library
- ▶ Allows an application to utilize built-in features of our graphics hardware (GPUs) to greatly speed up rendering



- ▶ Technically...
 - Our **C++ Code** → **Graphics API** → **Graphics Drivers** → **GPU Hardware**

Graphics API organization



- ▶ All APIs have a similar high-level organization
- ▶ GPU Resources → Data in GPU memory used while rendering
- ▶ Graphics State → Options for rendering & active resources
- ▶ Issuing Commands → Changing state, start rendering, etc.
- ▶ Rendering Pipeline → Steps to actually render

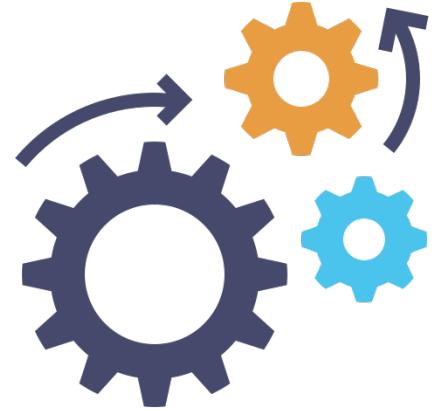
Graphics APIs & GPUs

- ▶ Graphics APIs help us...
 - **Manage** the current rendering state (options)
 - **Store** data properly on the GPU
 - **Invoke** specific graphics functionality
 - **Synchronize** GPU with our application
 - And much more

- ▶ All of this goes into properly rendering
 - And rendering quickly
 - To maintain interactive framerates



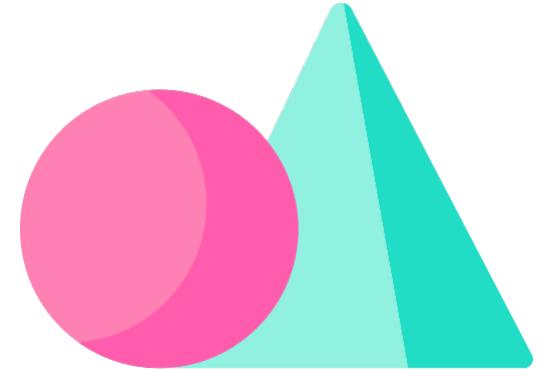
Graphics API usage



- ▶ All APIs have similar overarching steps for usage
 - VASTLY differ in the details & syntax
- ▶ **Initialize API** → Turn it on, bind to a window
- ▶ **Create Resources** → Load models & textures, set up state objects
 - Done during start-up as these are slow!
 - Will have many resources in GPU memory at once
 - But only use a few at a time
- ▶ **Render one frame** → Part of game loop
 - Repeat!

Rendering a frame

- ▶ **Clear output** → reset to a solid color
- ▶ Foreach entity we want to render...
 - Set graphics state for this object
 - Change rendering options
 - Set active triangles
 - Set active textures
 - Set active shaders
 - Issue **draw()** command → run rendering pipeline w/ active gfx state
- ▶ **Present** → show results of the frame to the user



Easy, right?

- ▶ Using a graphics API is easier than the alternative
 - Dealing with the drivers ourselves
- ▶ But they can be complex
- ▶ We're using DirectX 11 in this class
 - It's powerful
 - It's mature (many, many examples exist)
 - It's easier than lower-level APIs like DX12/Vulkan
 - We can learn the basics and do cool stuff with it quickly

