

# Skyboxes & Cube Maps



# The Sky



# What's a skybox?

- ▶ The infinitely-far-away backdrop to a 3D scene
- ▶ Always “behind” the game world



# Mixing with geometry

- ▶ Skybox is all the way in the back
- ▶ Not...
  - The dragon
  - Northern Lights



# We never get closer to the sky

- ▶ Otherwise it would look like we're approaching a wall



# Moving up doesn't make clouds larger



# Sky appears to “move” with you



# Sky can be animated and/or layered



# Should lights affect the sky?

- ▶ Nope!
- ▶ Just use raw texture sample
- ▶ Sky should appear to *provide* light to a scene



# What about clouds?

- ▶ Some modern games do render volumetric clouds
  - In front of the skybox



(Outside the scope of this lecture)

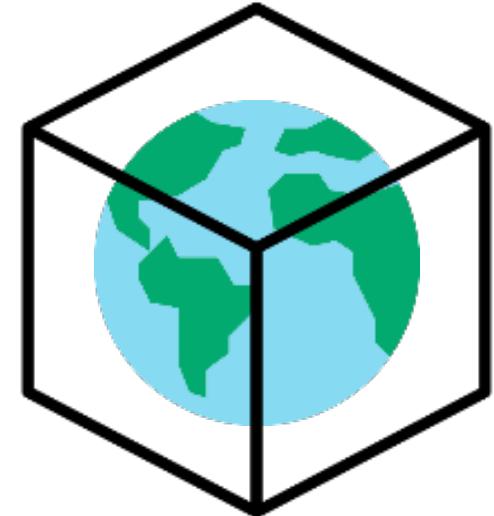
# Skybox Requirements

# Sky rendering overview

- ▶ To render the sky, we'll need:
  - A mesh
  - A texture
  - Specialized “sky” shaders
- ▶ The sky isn't a regular “game entity”
  - Lights don't affect it
  - Doesn't need a transform
  - There is an optimal time to render the sky

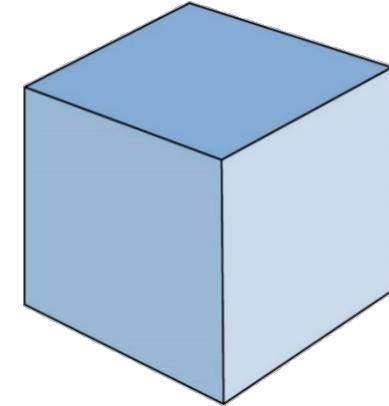
# Sky mesh: An actual box!

- ▶ The sky should fill the screen
  - Need to rasterize geometry “all around us”
  - Need a shape that fully surrounds us
  
- ▶ A literal box!
  - Cubes are nice and simple (12 triangles)
  - Not actually using its UVs
  - Just need a direction to each pixel in 3D space



# Sky texture?

- ▶ Need to texture our sky mesh
  - Seamlessly
  - Should look like a continuous image
- ▶ How do we do this with a single 2D texture?
- ▶ Spoilers
  - We generally *don't*
  - We need another type of GPU texture resource: *a cube map*



# Cube Maps

Six textures treated like they're a cube

# Cube maps

- ▶ A special type of GPU texture resource
- ▶ Internally, contains six Texture2Ds
  - Each of which represents one “face” of a cube
- ▶ Sometimes referred to as “Skyboxes”
  - This is not entirely accurate
  - Though it is a common colloquialism
  - Cube maps can be used for more than just the sky

# Cube map example

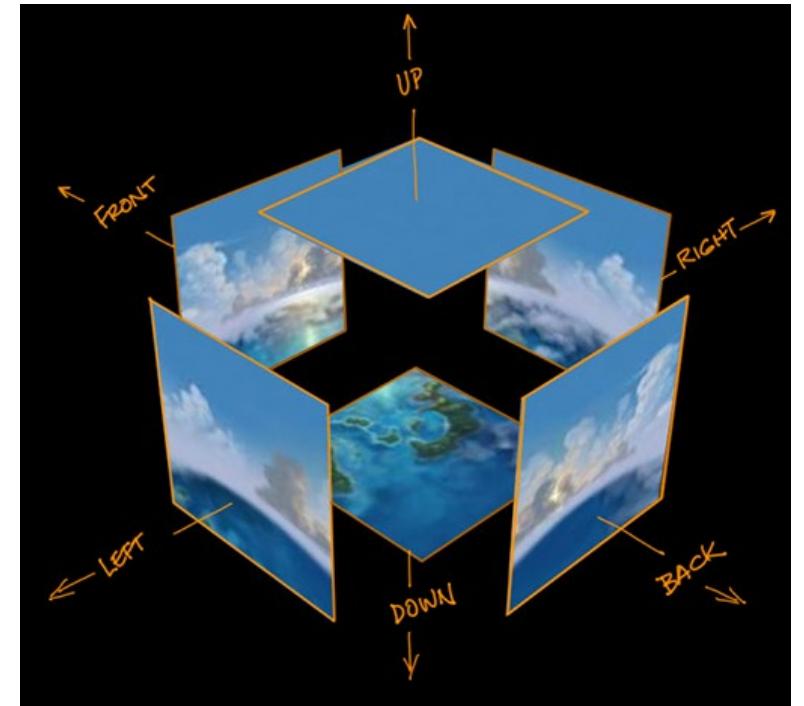
- ▶ 6 textures



- ▶ Cube Map containing those textures

- ▶ Specific face requirements

- Each texture edge must match neighbors
- Each texture must represent a 90° fov
- Otherwise it won't look seamless



# Cube map storage (file) options

- ▶ Many ways to organize these textures

- 6 separate image files
- 1 file with a horizontal layout
- 1 file with a “cross” layout

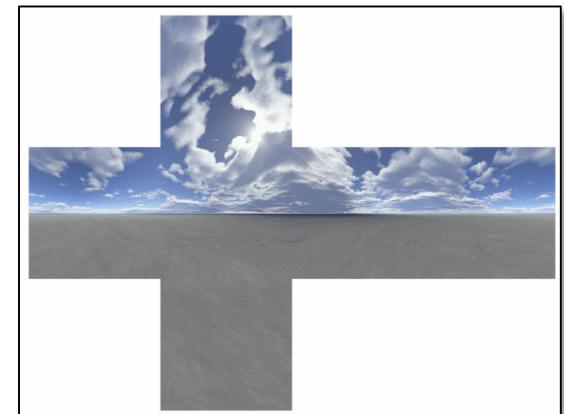


- ▶ All of these are feasible

- With the right code
- To “chop” them up



- ▶ We'll use the “6 image” approach



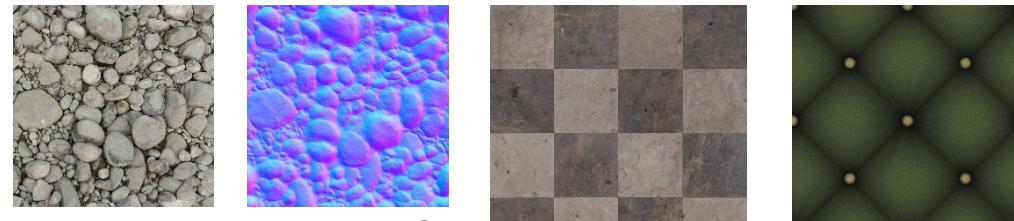
# Cube maps on the GPU

- ▶ Reminder: Textures can have 1 or more sub-resources
  - Can treat sub-resources as “array elements”
  - And flag the resource as a cube map instead of an array
- ▶ Note: This is an explicit GPU resource type!
  - Texture2DArray
  - NOT a C++ array of Texture2Ds
- ▶ We'll look at how to create this through code soon

# GPU texture example

Several individual 2D textures →

GPU Memory



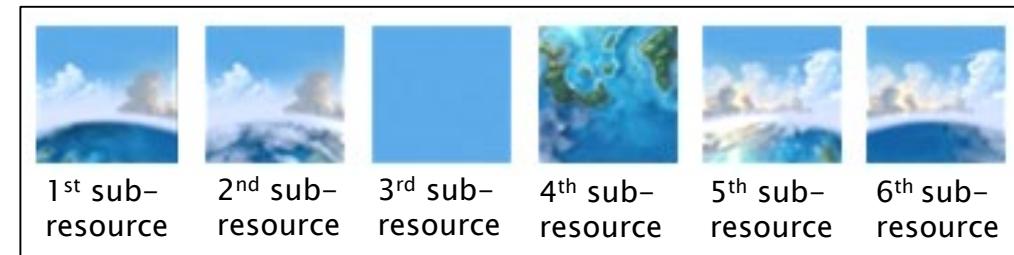
Texture 1

Texture 2

Texture 3

Texture 4

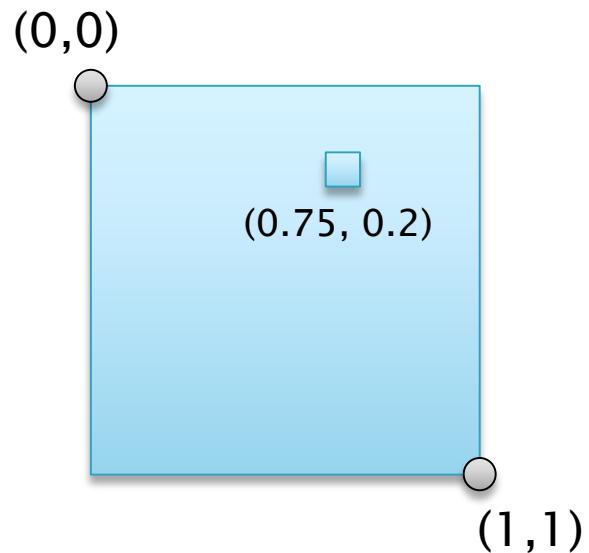
Single cube map texture  
with 6 sub-resources →



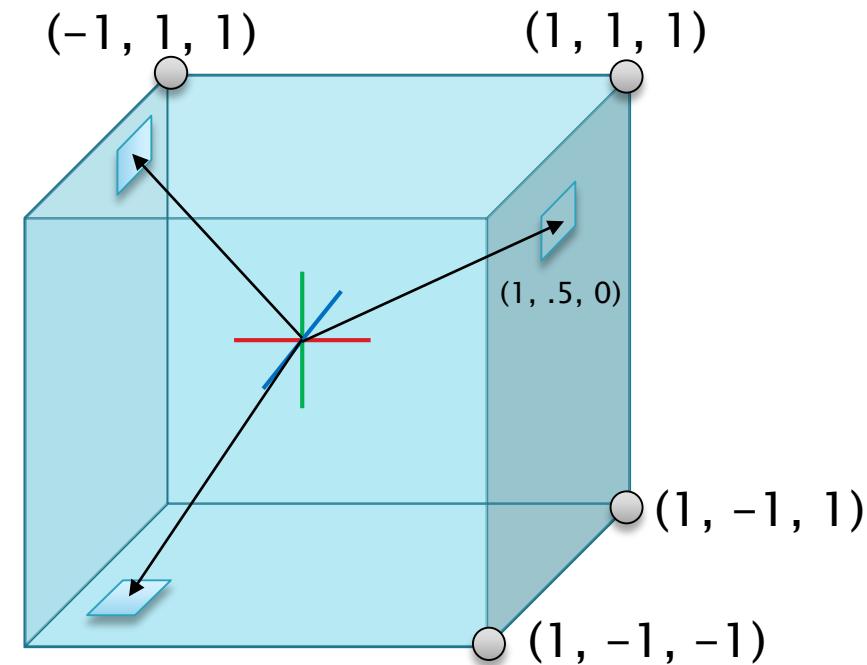
Texture 5

# Sampling: Texture2Ds vs. Cube Maps

- CubeMaps are sampled using a *direction* in space



Texture 2D  
(Uses 2D coordinates)



Texture Cube  
(Uses a direction)

# Cube maps in HLSL

- ▶ Specify `TextureCube` in shader

```
Texture2D  SurfaceTexture : register(t0);
TextureCube SkyboxTexture : register(t1);
```

- ▶ Sampling a `TextureCube` requires a 3D *direction*

```
float3 direction = // Some useful direction for this pixel
float4 skyColor = SkyboxTexture.Sample(BasicSampler, direction);
```

- ▶ Imagine being in the center of the cube
  - Given a direction in 3D space...
  - Sample the cube map's color in that direction

# Uses for cube maps

- ▶ Skies
- ▶ Environment mapping
  - “Pre-baked” reflections
- ▶ Omnidirectional shadows
  - More on this when we get to shadow mapping
- ▶ And more!

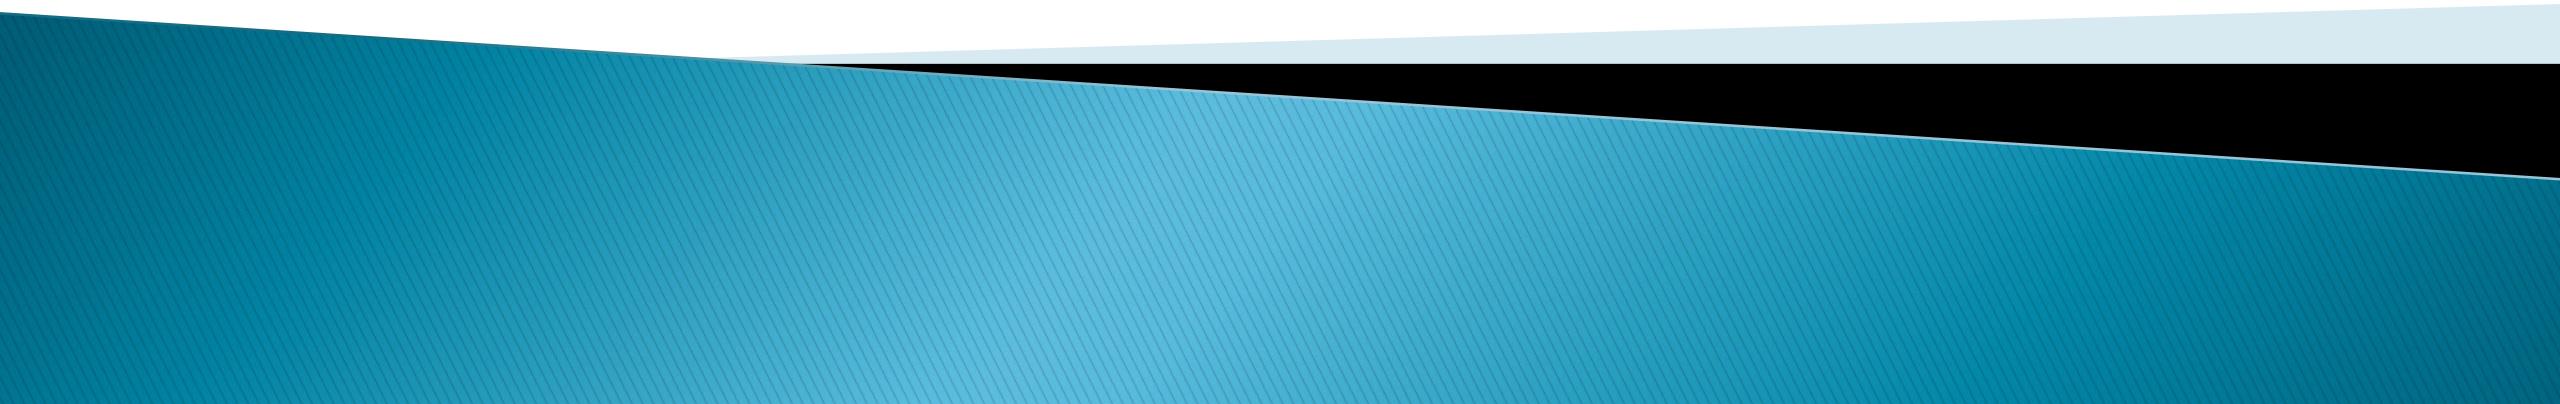


# Quick note: Why not just UV map a cube?

- ▶ Texture filtering may result in obvious lines at geometry borders!
- ▶ Cube maps don't have this issue
- ▶ GPUs properly filter cube map "neighbors"



# Actually Rendering the Sky



# A few issues to tackle

- ▶ Need to render the inside of the cube
- ▶ Where should the cube be in our 3D scene?
- ▶ How do we make the cube appear “behind” everything?
- ▶ When should we render the sky?
- ▶ How do we get the direction to each pixel?

# Rendering the inside

- ▶ By default, the pipeline only rasterizes “front faces”
  - We see the outside of objects
  - But not the insides
  - Useful optimization
- ▶ We can change this behavior
- ▶ Need a D3D API object that contains rasterizer options

# ID3D11Rasterizer State

- ▶ Define the object

```
Microsoft::WRL::ComPtr<ID3D11RasterizerState> skyRasterState;
```

- ▶ Create during initialization

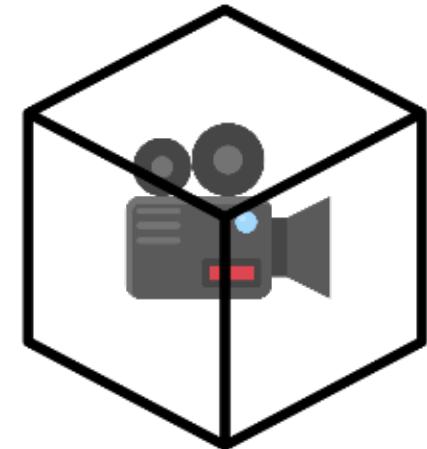
```
D3D11_RASTERIZER_DESC rastDesc = {};
rastDesc.CullMode = D3D11_CULL_FRONT; // Draw the inside instead of the outside!
rastDesc.FillMode = D3D11_FILL_SOLID;
rastDesc.DepthClipEnable = true;
device->CreateRasterizerState(&rastDesc, skyRasterState.GetAddressOf());
```

- ▶ Set/unset during Draw()

```
context->RSSetState(skyRasterState.Get());
// Draw the skybox here!
context->RSSetState(0); // Null (or 0) restores the default state
```

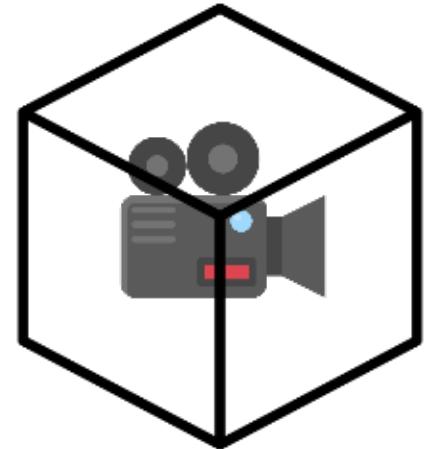
# Where is the cube?

- ▶ Skybox should be centered on camera
  - Always the same distance
  - We never get closer to the skybox
  
- ▶ How do we ensure two objects are at the same position?
  - Move *both*?
  
- ▶ Even better: move *neither*!
  - Keep them both at the origin
  - Results will be the same



# Centering the cube

- ▶ View matrix reminder
  - Makes everything *relative to the camera*
  - Essentially a “reverse transformation”
- ▶ If we move an object to the camera, view matrix *undoes that*
- ▶ If you have a box on your head...
  - It looks the same no matter where you are
  - Same idea here



# Vert shader: Center skybox on camera

```
// Use this when rendering skybox geometry
// - Removes translation portion of camera's view matrix
// - And does NOT use a world matrix for skybox geometry
// - End result: Sky always appears centered on camera

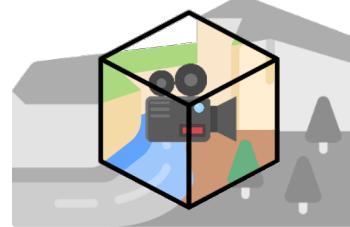
// Create a version of the view matrix without translation
matrix viewNoTranslation = view;
viewNoTranslation._14 = 0;
viewNoTranslation._24 = 0;
viewNoTranslation._34 = 0;

// Calculate output position w/ just view & projection
matrix viewProj = mul(projection, viewNoTranslation);
output.position = mul(viewProj, float4(input.position, 1.0f));
```

# Ensuring the cube is behind everything

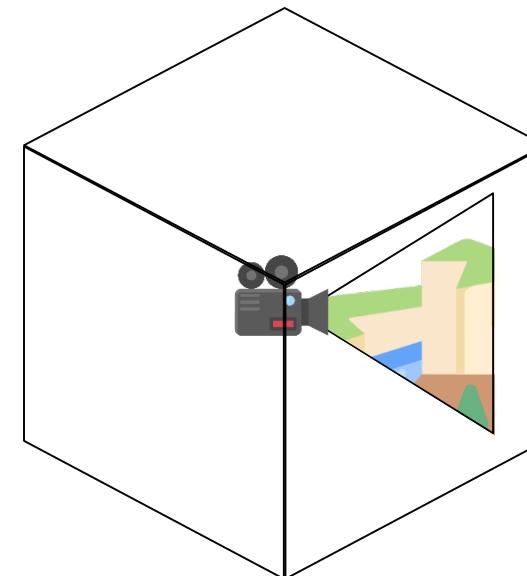
- ▶ If the box is too small

- It blocks the world
  - Like a wall around us



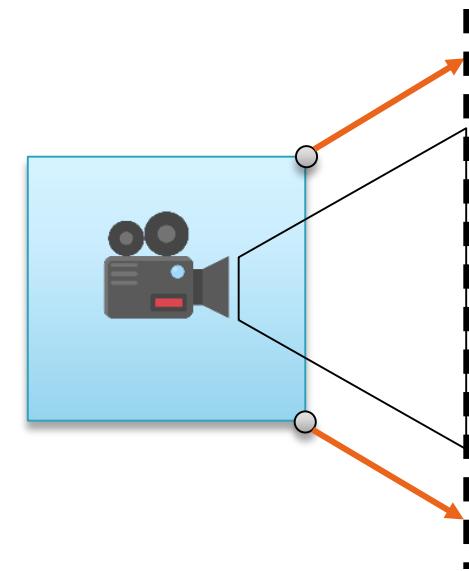
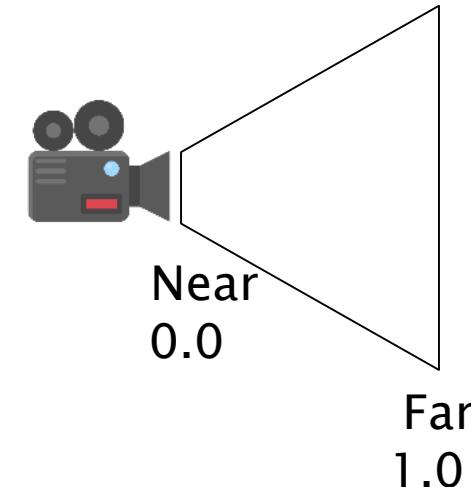
- ▶ If the box is too big

- It's beyond our camera's far plane
  - Can't actually see it



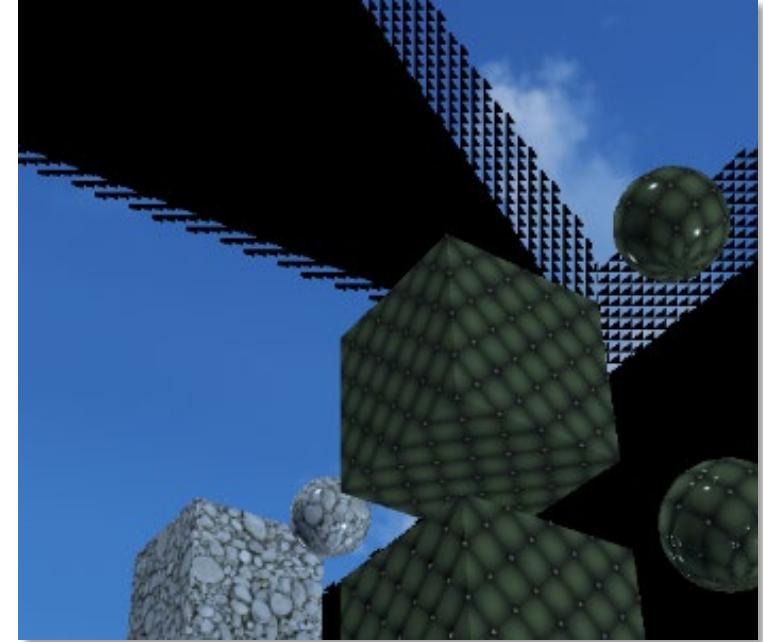
# Reminder: Depth & clip planes

- ▶ Post-projection Z value is *depth*
  - Distance between near and far clip planes
  - Depth values between 0 and 1 are “visible”
- ▶ Largest visible depth value? 1.0
- ▶ How do we ensure that?
  - We overwrite depth in the vertex shader
  - Essentially “pushing” the vertices onto far plane
  - Set  $Z = W$ , so  $Z/W$  division yields a depth of 1.0



# A new depth issue

- ▶ We might run into the following issue
  - Parts of the sky are missing!
  - And they flicker wildly
- ▶ What's going on?
- ▶ When we clear the depth buffer, we set all values to 1.0
- ▶ When we draw the skybox, we set depths to 1.0
- ▶ By default, pipeline *ignores* pixels w/ depth == depth buffer
  - We still see a few due to floating point rounding errors



# Change depth state to accept more pixels

- ▶ Define the object

```
Microsoft::WRL::ComPtr<ID3D11DepthStencilState> skyDepthState;
```

- ▶ Create during initialization

```
// Depth state so that we accept pixels with a depth <= 1
D3D11_DEPTH_STENCIL_DESC depthDesc = {};
depthDesc.DepthEnable      = true;
depthDesc.DepthFunc        = D3D11_COMPARISON_LESS_EQUAL;
depthDesc.DepthWriteMask   = D3D11_DEPTH_WRITE_MASK_ALL;
device->CreateDepthStencilState(&depthDesc, skyDepthState.GetAddressOf());
```

- ▶ Set/unset during Draw()

```
context->OMSetDepthStencilState(skyDepthState.Get(), 0);
// Draw skybox here! (And also set/unset rasterizer state)
context->OMSetDepthStencilState(0, 0);
```

# When do we render the sky?

- ▶ Before other objects?
  - Pro: Easy to implement & probably fast enough
  - Con: Could result in lots of *overdraw*
- ▶ Even better: After all other (opaque) objects!
  - Rasterizer will skip pixels based on depth buffer
  - All scene objects will have depth  $< 1$
  - Only shade pixels we can see!
- ▶ Note: Transparent objects need to come after skybox

# Direction to each pixel for cubemap sampling

- ▶ Sky mesh's verts ideally centered around origin
  - They're basically directions from the origin
  - Vertex position == sampling direction
- ▶ Pass that vector from vertex shader to pixel shader
  - `vsOutput.sampleDirection = input.localPosition;`
- ▶ Pixel shader is extremely simple
  - Sample cube map in given direction
  - Return that color

# **Environment Mapping**

Using a cube map for reflections



# Environment mapping

- ▶ If we have a cube map that represents “the world around us”
- ▶ We can use it on reflective surfaces



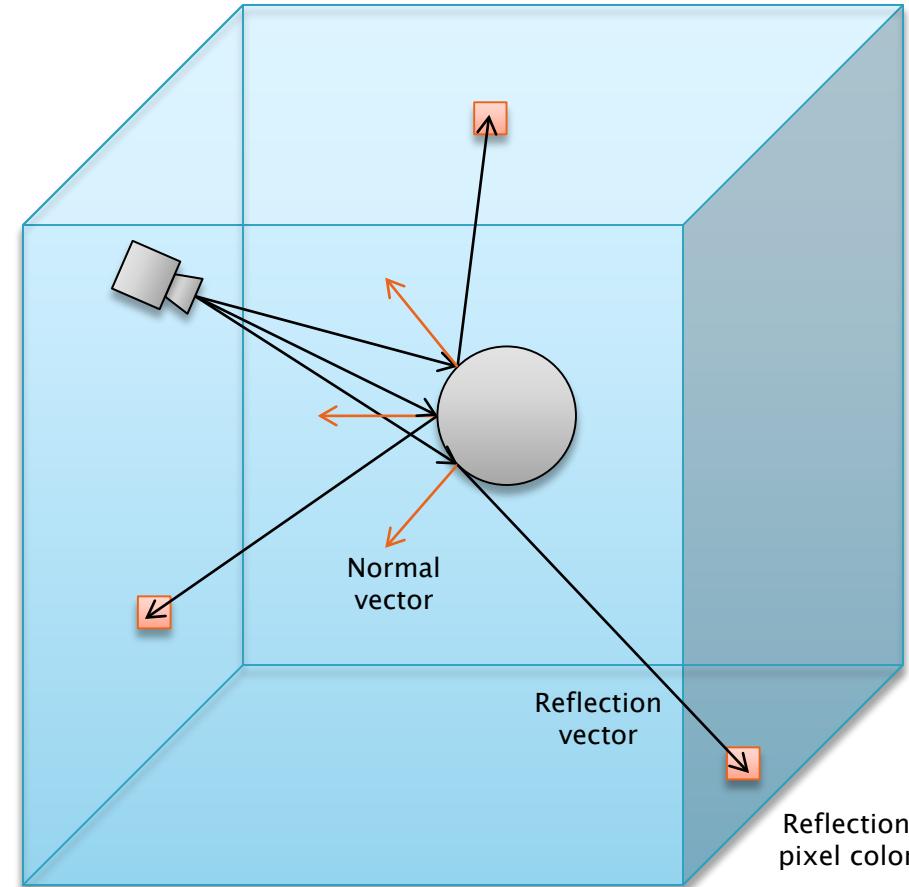
# Reflections

- ▶ Use the same map for the sky and reflections
- ▶ Looks like objects reflect the environment!



# Calculating a reflection

- ▶ Figure out the direction from the camera to the current pixel
- ▶ Calculate the reflection vector
- ▶ Sample cube map in that direction



# Reflection sampling

- ▶ Remember: Cube Map sampling uses a direction
- ▶ Calculate direction from camera to vertex
  - Do this all in world-space!
  - Camera position in world space
  - Vert position in world space: `vertPos * worldMatrix`
- ▶ Use HLSL “reflect” function
  - Reflects a vector based on a normal
  - Use vertex’s normal
  - Use output of function as cube map sample vector

# Environment mapping problems

- ▶ Environment maps are usually static
  - Doesn't always match up with real environment



# Environment mapping problems

- ▶ Static cube maps aren't always perfect



# Real-time cube maps

- ▶ You can render to a cube map!
  - Basically rendering 6 times
  - Each time to a different face
- ▶ Use resulting cube map as the reflection map
  - Too expensive to render for EACH object
  - Usually done on “main character”
- ▶ Can also render a cube shadow map
  - Used for point light shadow maps

# Real-time reflection example

- ▶ Watch the very top of the car



# Creating a Cubemap From Six Textures

Sample Code

# Signature

```
// -----
// Loads six individual textures (the six faces of a cube map), then
// creates a blank cube map and copies each of the six textures to
// another face. Afterwards, creates a shader resource view for
// the cube map and cleans up all of the temporary resources.
// -----
Microsoft::WRL::ComPtr<ID3D11ShaderResourceView> Game::CreateCubemap(
    const wchar_t* right,
    const wchar_t* left,
    const wchar_t* up,
    const wchar_t* down,
    const wchar_t* front,
    const wchar_t* back)
{
```

# Load the textures & get info

```
// Load the 6 textures into an array.  
// - We need references to the TEXTURES, not the SHADER RESOURCE VIEWS!  
// - Specifically NOT generating mipmaps, as we don't need them for the sky!  
// - Order matters here! +X, -X, +Y, -Y, +Z, -Z  
  
ID3D11Texture2D* textures[6] = {};  
CreateWICTextureFromFile(device.Get(), right, (ID3D11Resource**)&textures[0], 0);  
CreateWICTextureFromFile(device.Get(), left, (ID3D11Resource**)&textures[1], 0);  
CreateWICTextureFromFile(device.Get(), up, (ID3D11Resource**)&textures[2], 0);  
CreateWICTextureFromFile(device.Get(), down, (ID3D11Resource**)&textures[3], 0);  
CreateWICTextureFromFile(device.Get(), front, (ID3D11Resource**)&textures[4], 0);  
CreateWICTextureFromFile(device.Get(), back, (ID3D11Resource**)&textures[5], 0);  
  
// We'll assume all of the textures are the same color format and resolution,  
// so get the description of the first shader resource view  
D3D11_TEXTURE2D_DESC faceDesc = {};  
textures[0]->GetDesc(&faceDesc);
```

# Create the cube texture

```
// Describe the resource for the cube map, which is simply
// a "texture 2d array". This is a special GPU resource format,
// NOT just a C++ array of textures!!!
D3D11_TEXTURE2D_DESC cubeDesc = {};
cubeDesc.ArraySize      = 6; // Cube map!
cubeDesc.BindFlags      = D3D11_BIND_SHADER_RESOURCE; // We'll be using as a texture in a shader
cubeDesc.CPUAccessFlags = 0; // No read back
cubeDesc.Format         = faceDesc.Format; // Match the loaded texture's color format
cubeDesc.Width          = faceDesc.Width; // Match the size
cubeDesc.Height         = faceDesc.Height; // Match the size
cubeDesc.MipLevels      = 1; // Only need 1
cubeDesc.MiscFlags      = D3D11_RESOURCE_MISC_TEXTURECUBE; // This is a CUBE, not 6 separate textures
cubeDesc.Usage           = D3D11_USAGE_DEFAULT; // Standard usage
cubeDesc.SampleDesc.Count = 1;
cubeDesc.SampleDesc.Quality = 0;

// Create the actual texture resource
ID3D11Texture2D* cubeMapTexture = 0;
device->CreateTexture2D(&cubeDesc, 0, &cubeMapTexture);
```

# Loop and copy textures to cube

```
// Loop through the individual face textures and copy them,
// one at a time, to the cube map texture
for (int i = 0; i < 6; i++)
{
    // Calculate the subresource position to copy into
    unsigned int subresource =
        0, // Which mip (zero, since there's only one)
        i, // Which array element?
        1); // How many mip levels are in the texture?

    // Copy from one resource (texture) to another
    context->CopySubresourceRegion(
        cubeMapTexture, // Destination resource
        subresource, // Dest subresource index (one of the array elements)
        0, 0, 0, // XYZ location of copy
        textures[i], // Source resource
        0, // Source subresource index (we're assuming there's only one)
        0); // Source subresource "box" of data to copy (zero = whole thing)
}
```

# Create final SRV, clean up, return

```
// At this point, all of the faces have been copied into the
// cube map texture, so we can describe a shader resource view for it
D3D11_SHADER_RESOURCE_VIEW_DESC srvDesc = {};
srvDesc.Format = cubeDesc.Format;           // Same format as texture
srvDesc.ViewDimension = D3D11_SRV_DIMENSION_TEXTURECUBE; // Treat this as a cube!
srvDesc.TextureCube.MipLevels = 1;          // Only need access to 1 mip
srvDesc.TextureCube.MostDetailedMip = 0;    // Index of the first mip we want to see

// Make the SRV
Microsoft::WRL::ComPtr<ID3D11ShaderResourceView> cubeSRV;
device->CreateShaderResourceView(cubeMapTexture, &srvDesc, cubeSRV.GetAddressOf());

// Now that we're done, clean up the stuff we don't need anymore
cubeMapTexture->Release(); // Done with this particular reference (the SRV has another)
for (int i = 0; i < 6; i++)
    textures[i]->Release();

// Send back the SRV, which is what we need for our shaders
return cubeSRV;
}
```